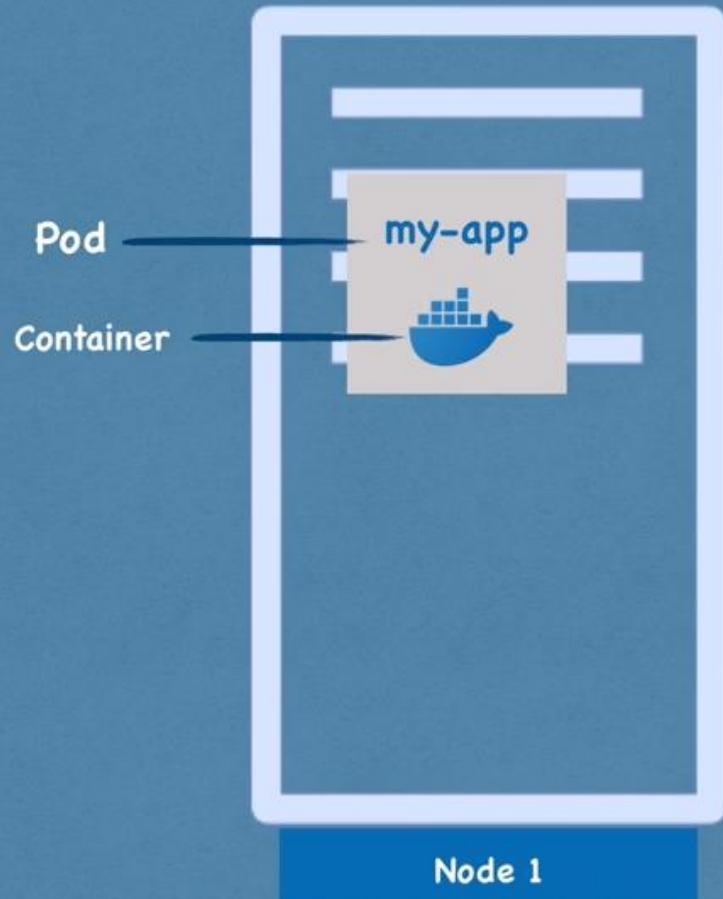




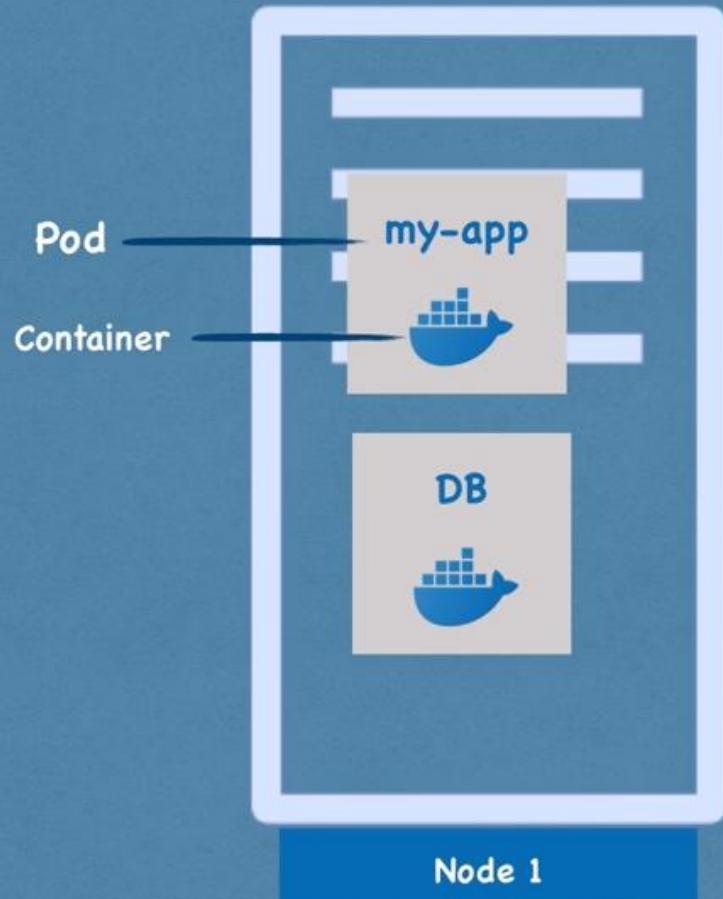
# Node and Pod





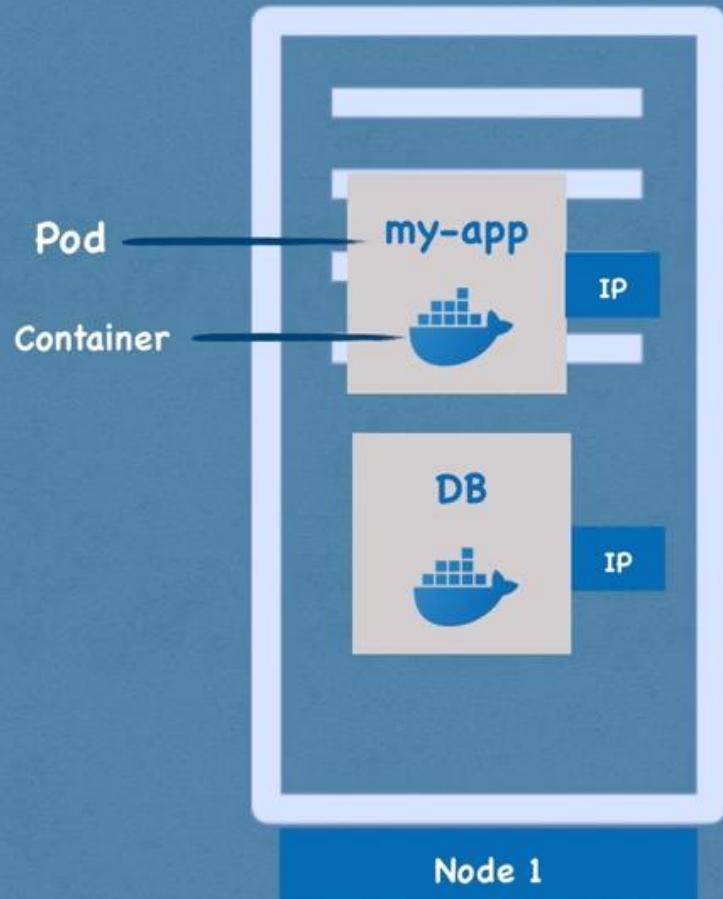
## Pod:

- Smallest unit of K8s
- Abstraction over container



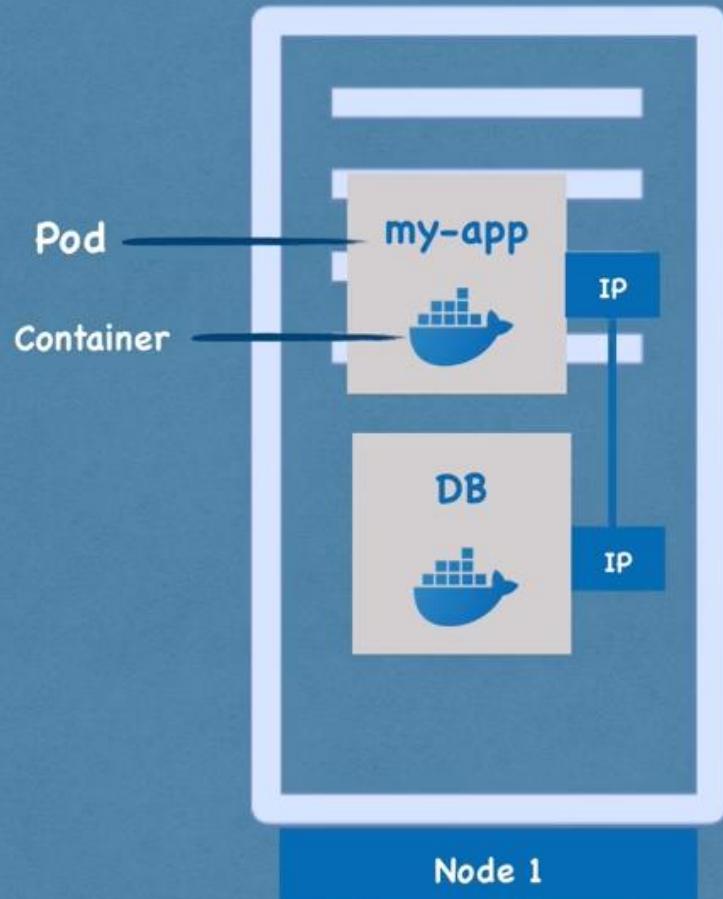
## Pod:

- Smallest unit of K8s
- Abstraction over container
- Usually 1 application per Pod



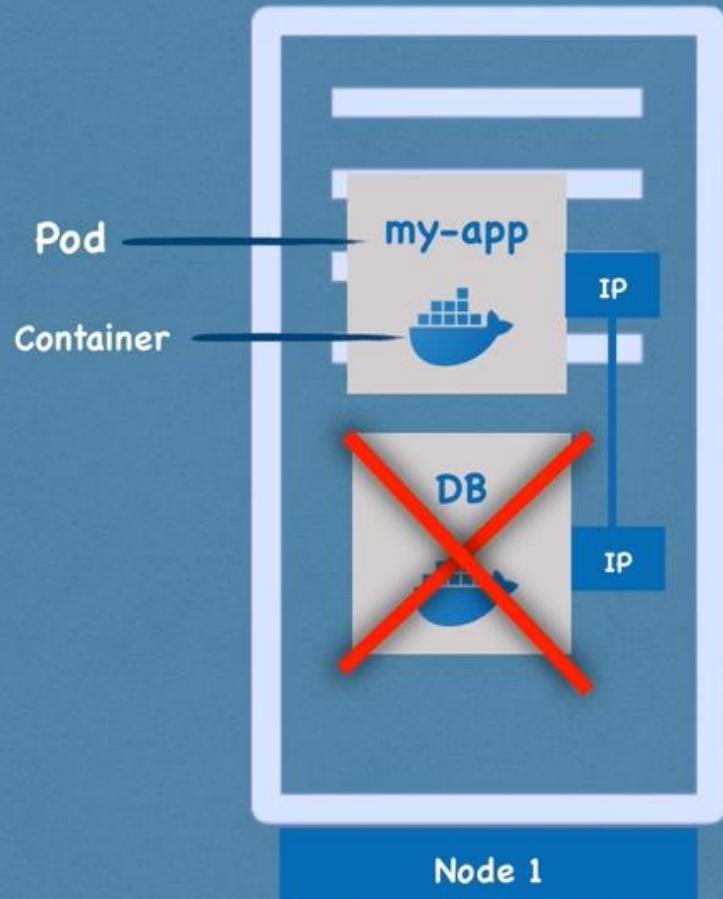
## Pod:

- Smallest unit of K8s
- Abstraction over container
- Usually 1 application per Pod
- Each Pod gets its own IP address



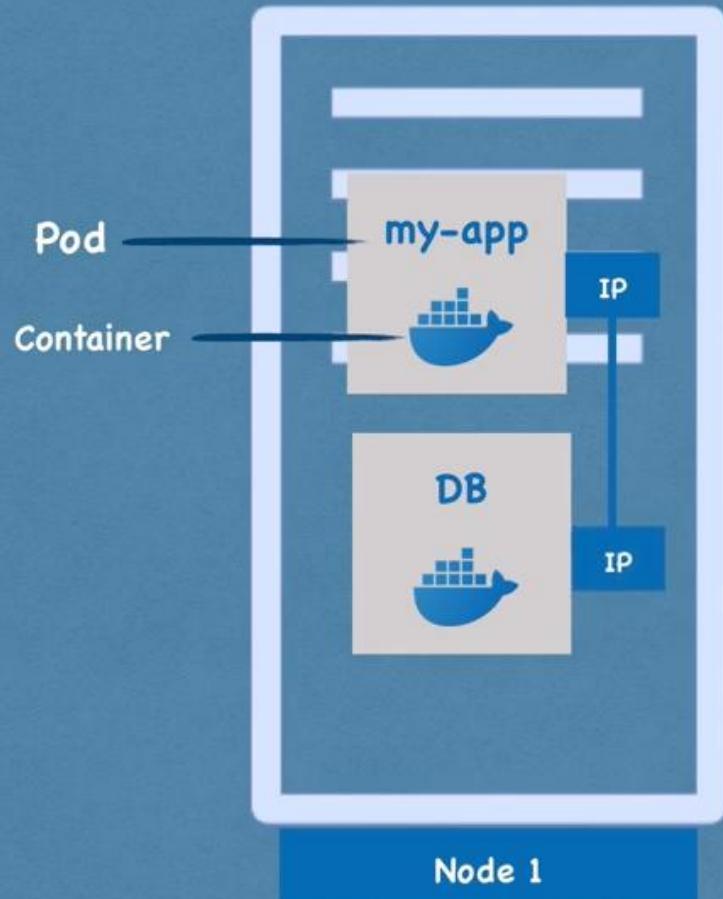
## Pod:

- Smallest unit of K8s
- Abstraction over container
- Usually 1 application per Pod
- Each Pod gets its own IP address



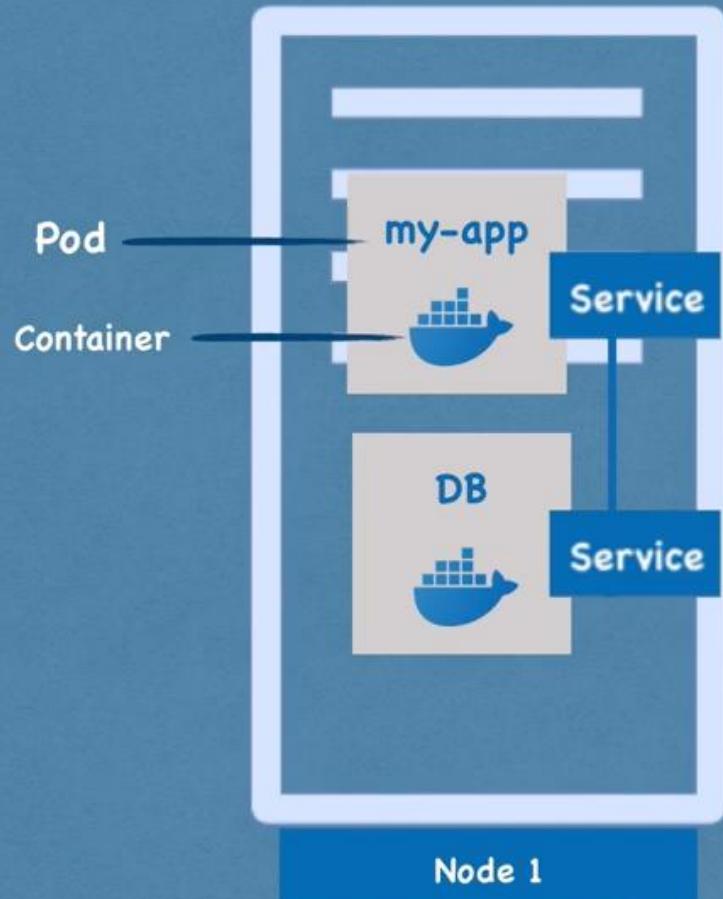
## Pod:

- Smallest unit of K8s
- Abstraction over container
- Usually 1 application per Pod
- Each Pod gets its own IP address



## Pod:

- Smallest unit of K8s
- Abstraction over container
- Usually 1 application per Pod
- Each Pod gets its own IP address
- New IP address on re-creation



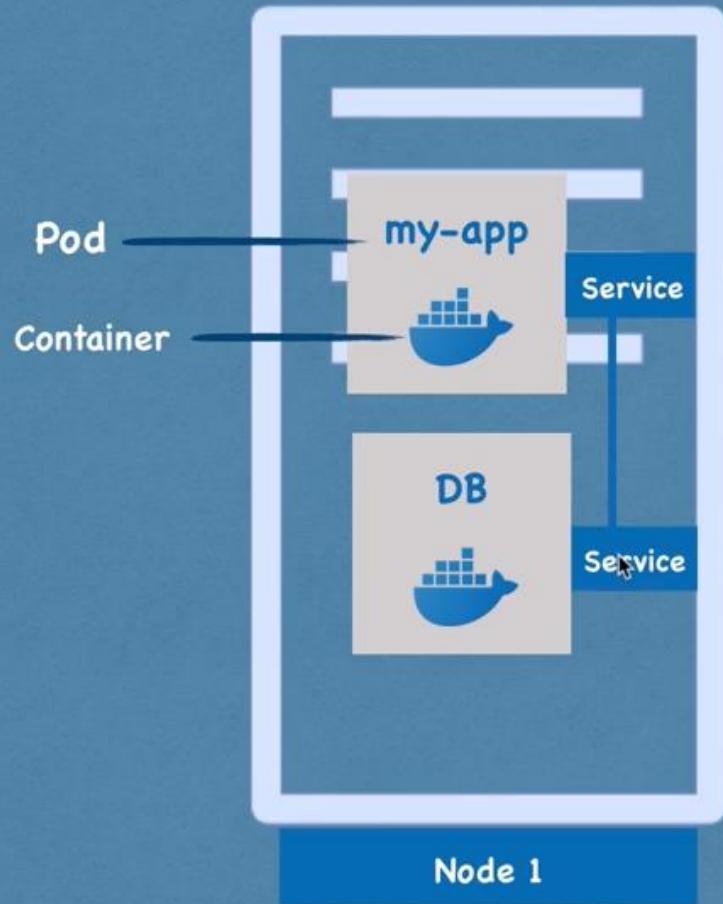
## Pod:

- Smallest unit of K8s
- Abstraction over container
- Usually 1 application per Pod
- Each Pod gets its own IP address
- New IP address on re-creation



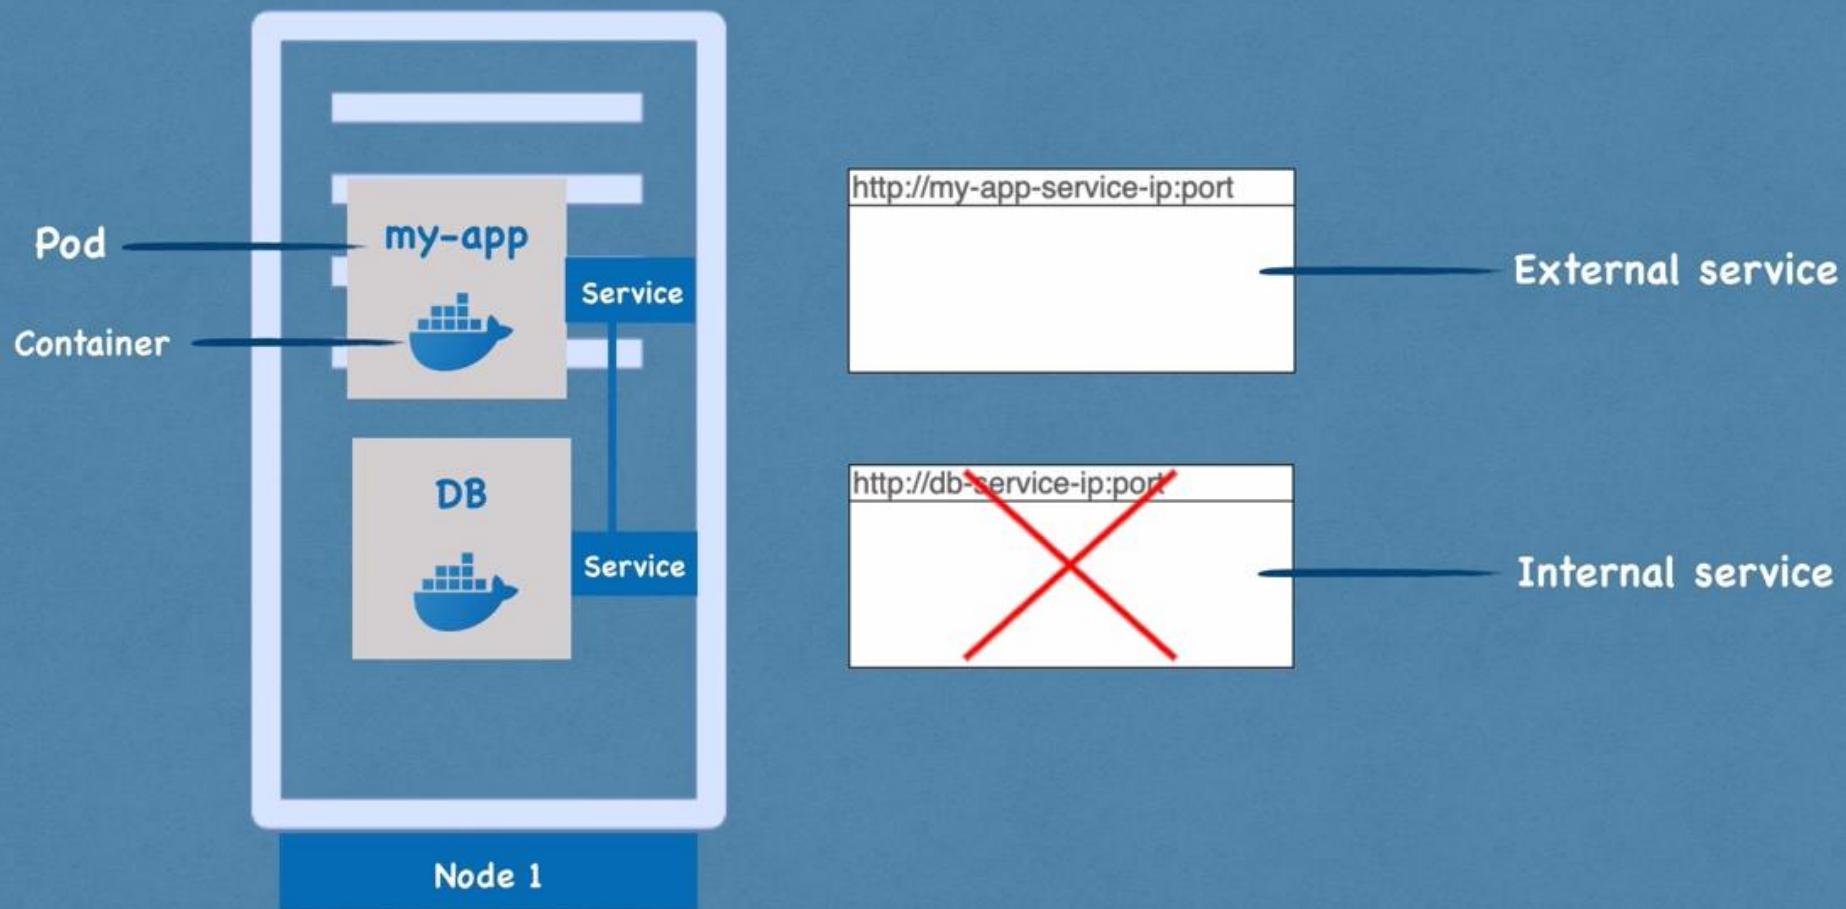
# Service and Ingress





## Service:

- permanent IP address

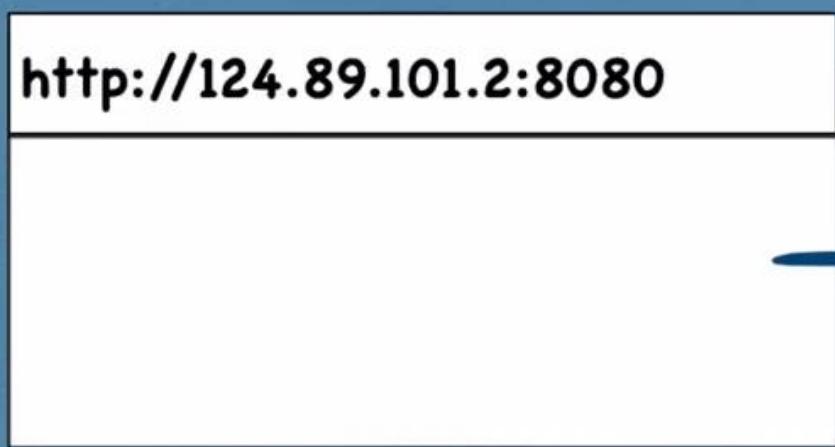




External service



Internal service



External service



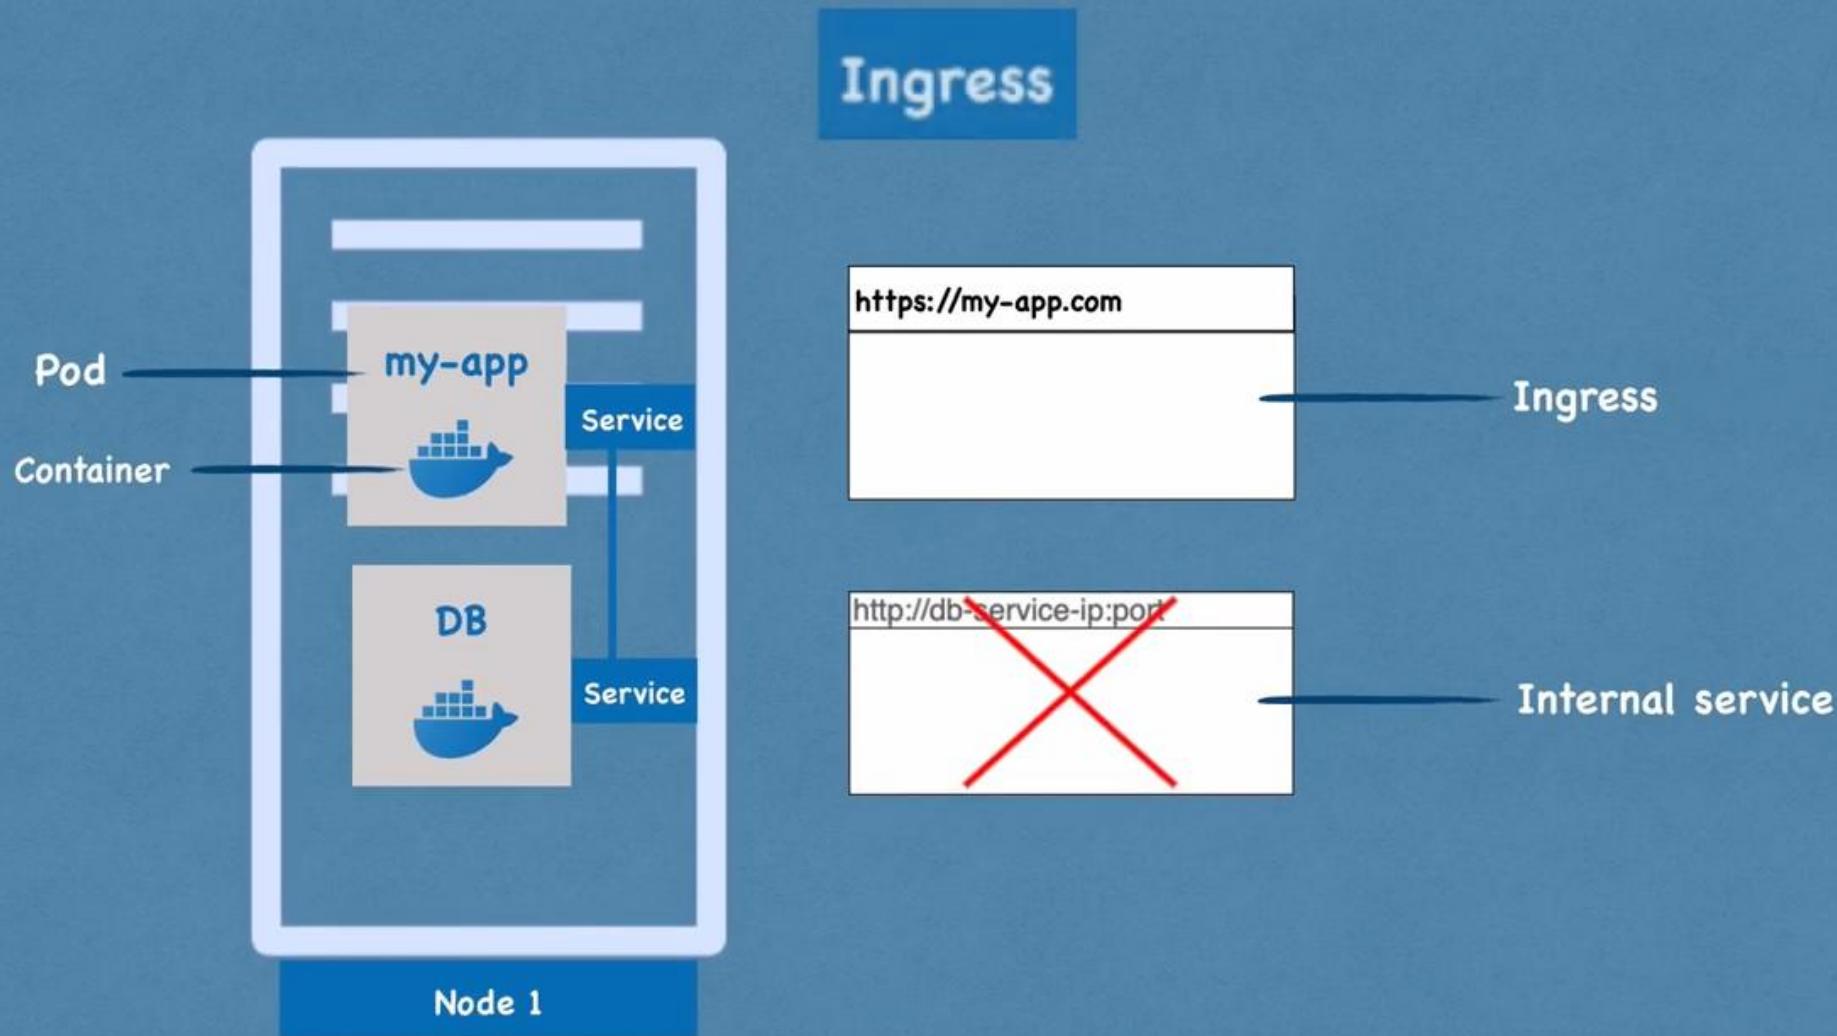
Internal service

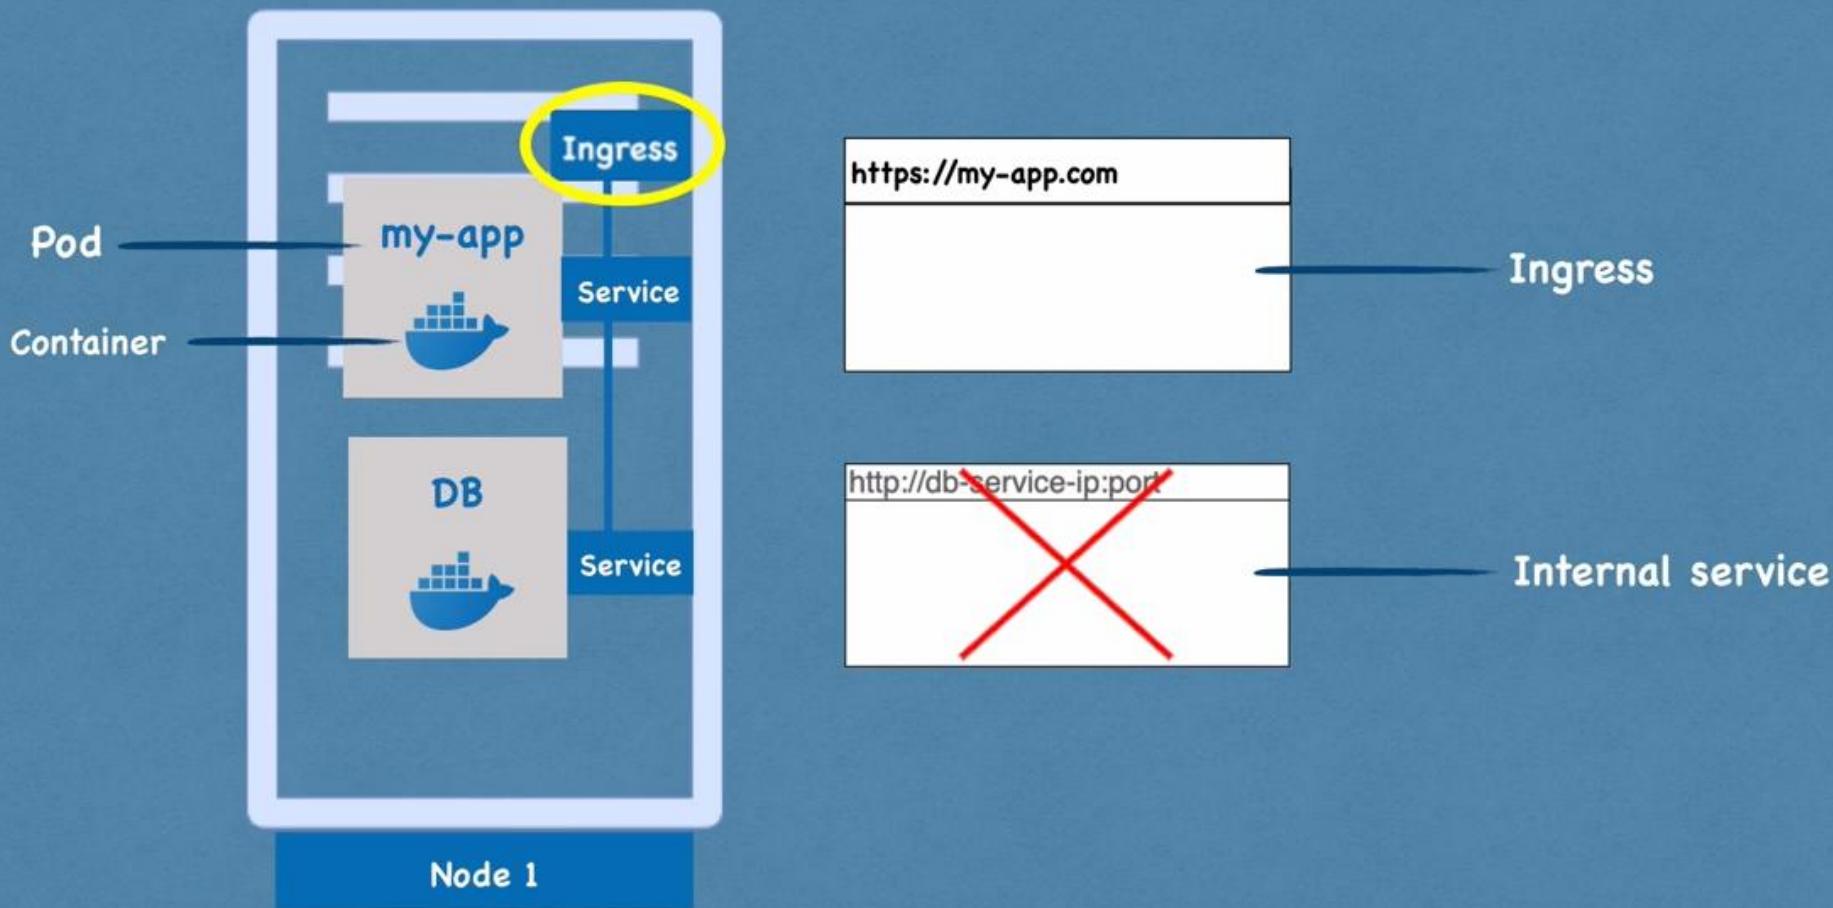


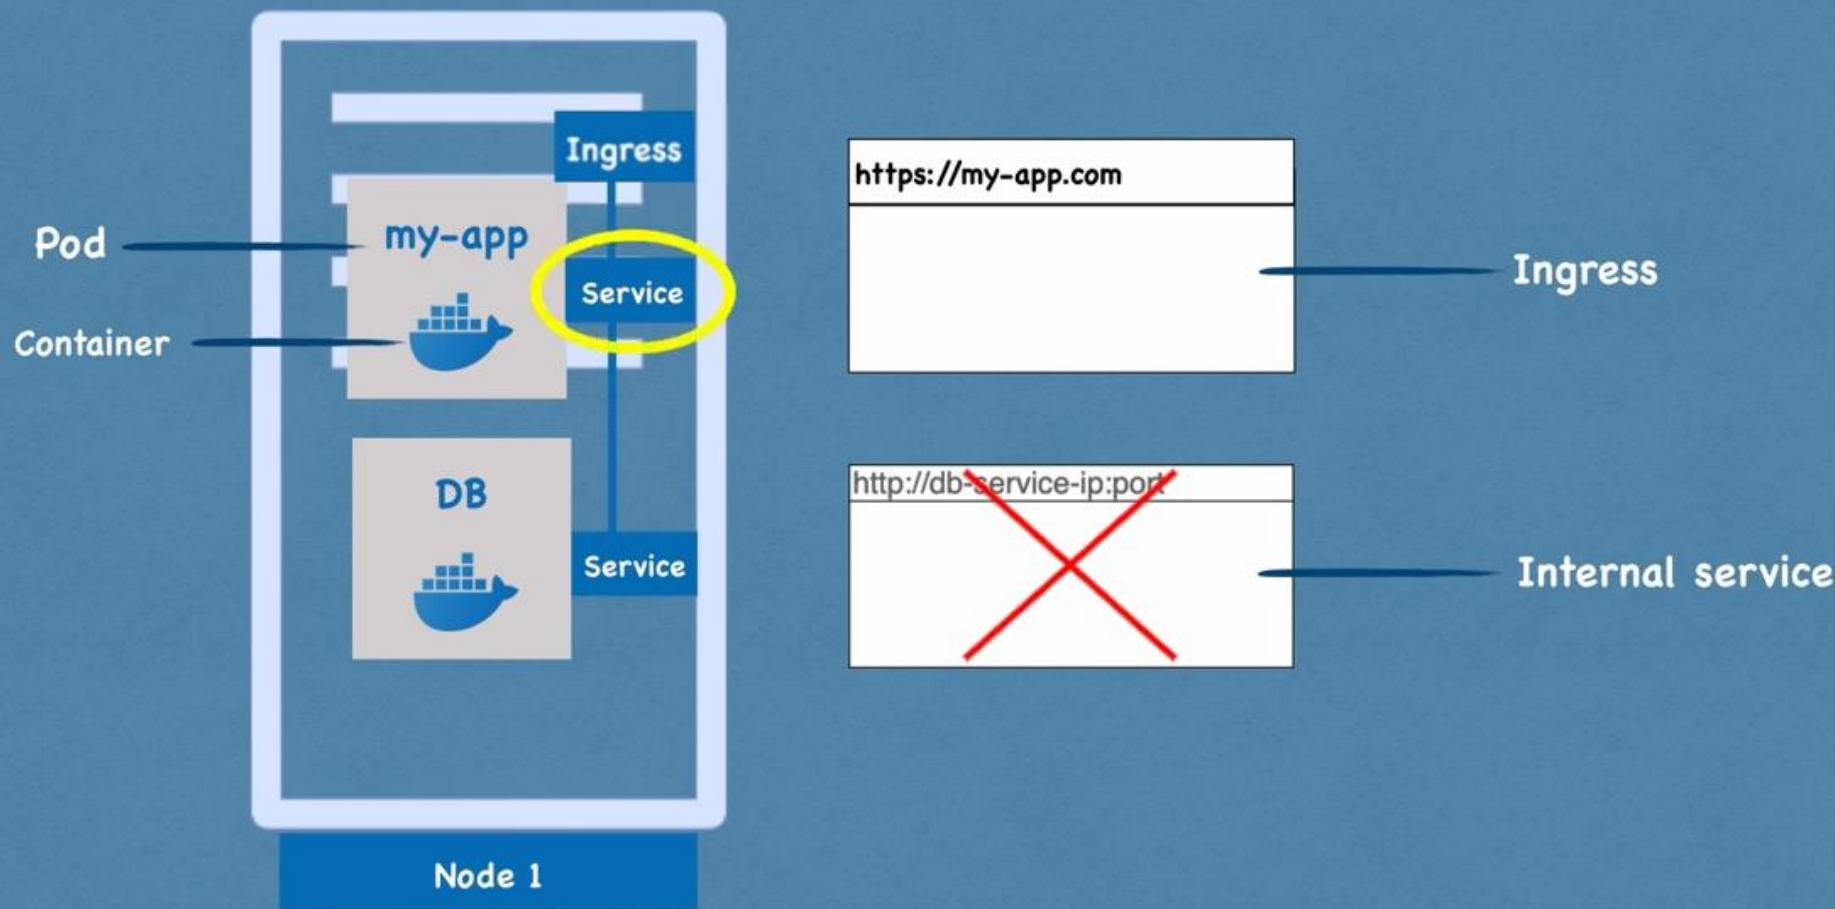
External service



Internal service





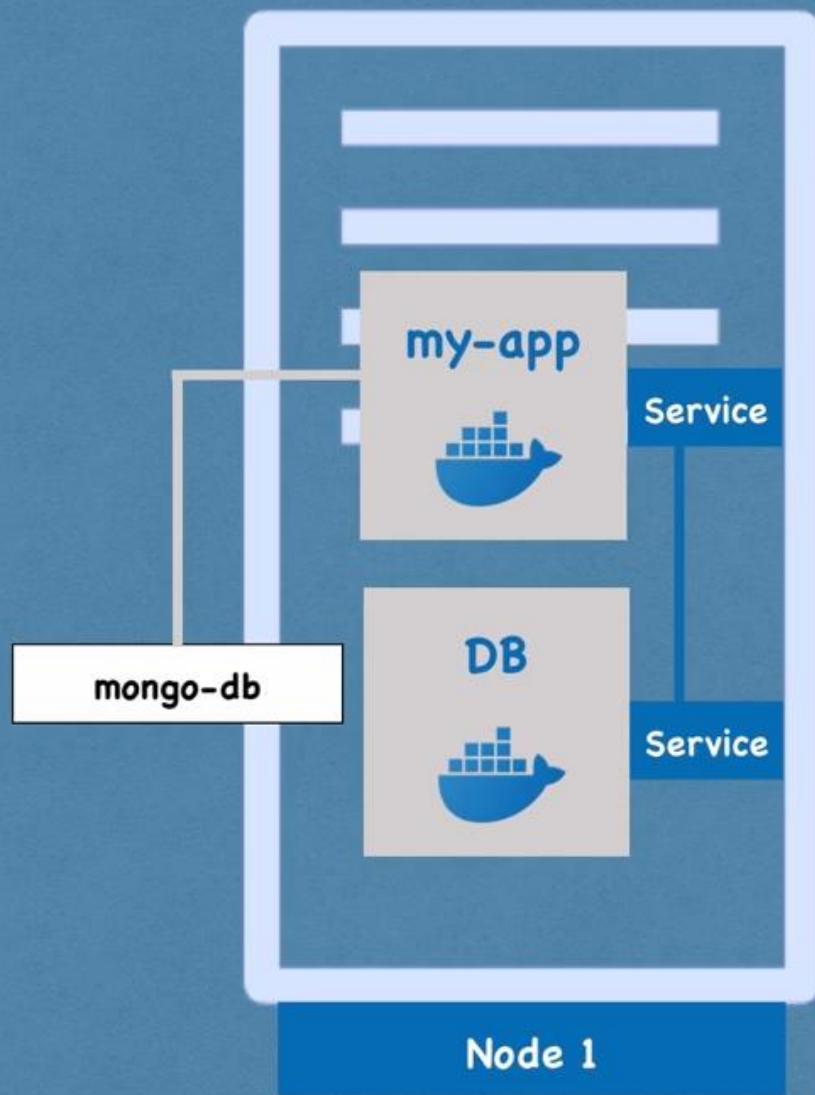


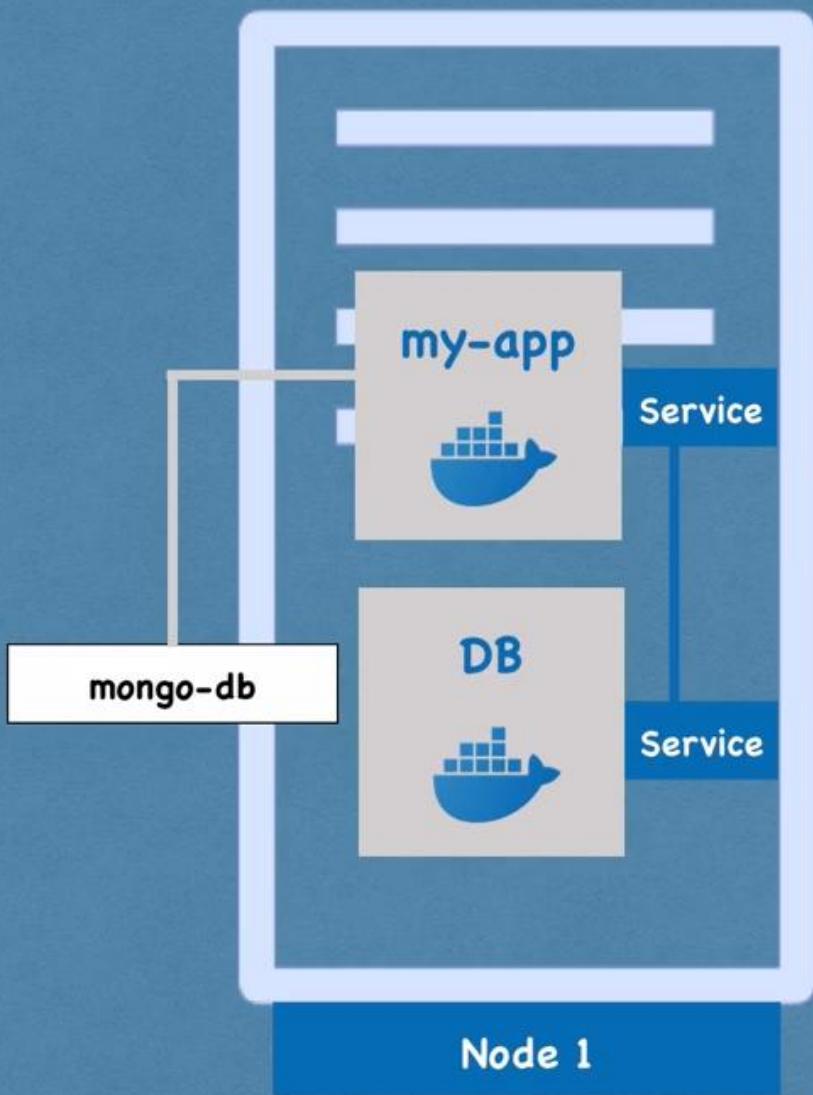


# ConfigMap and Secret



**Database Url usually in  
the built application!**





Database Url usually in  
the built application!

Repository

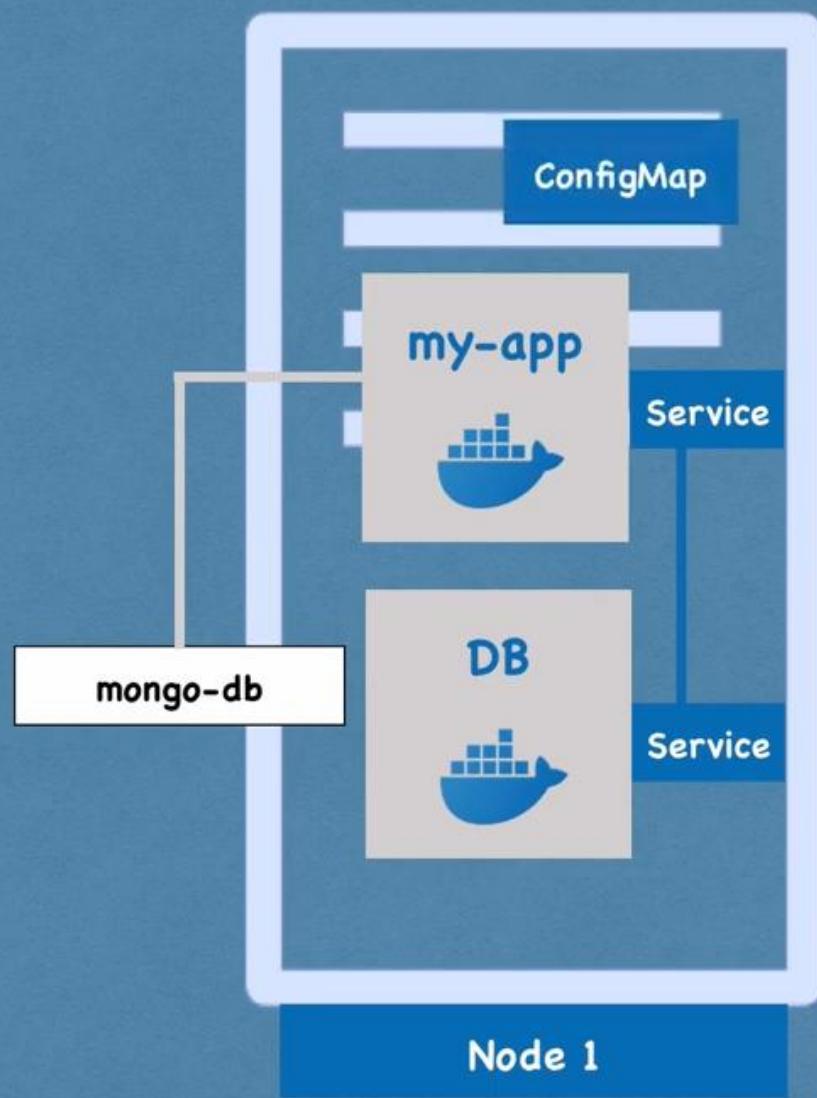


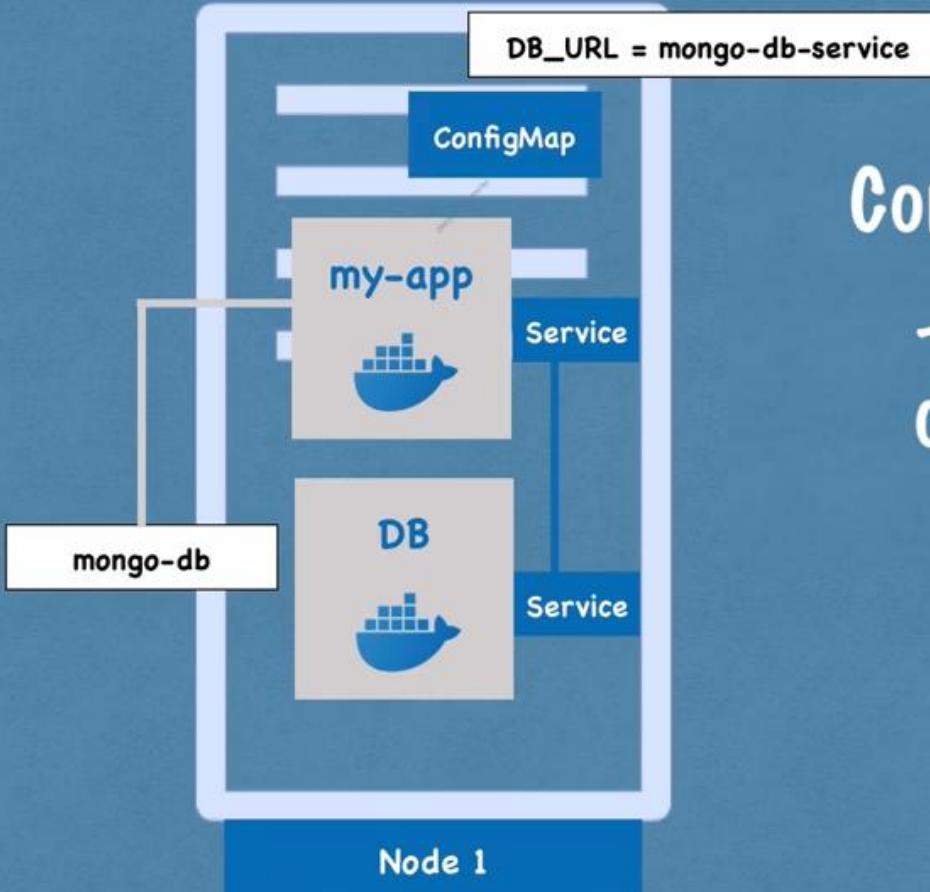
re-build



push it to repo

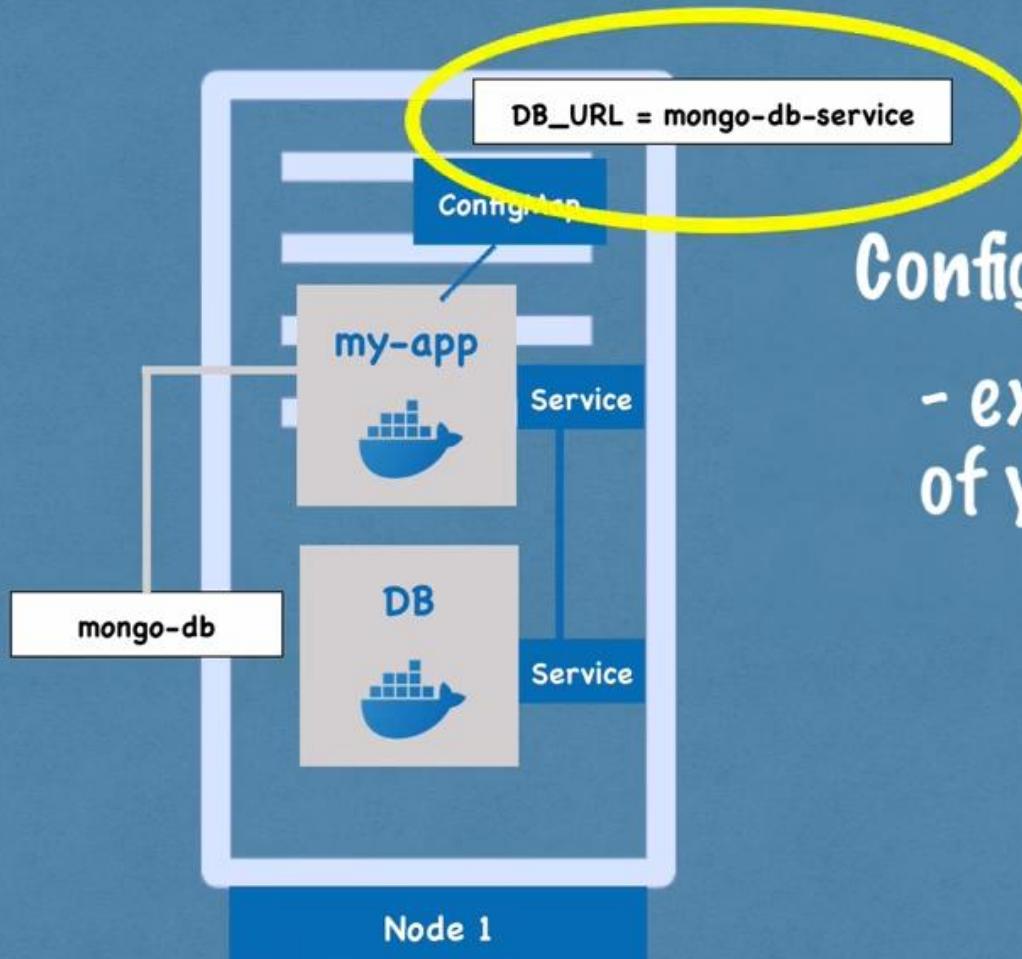
pull it in your pod





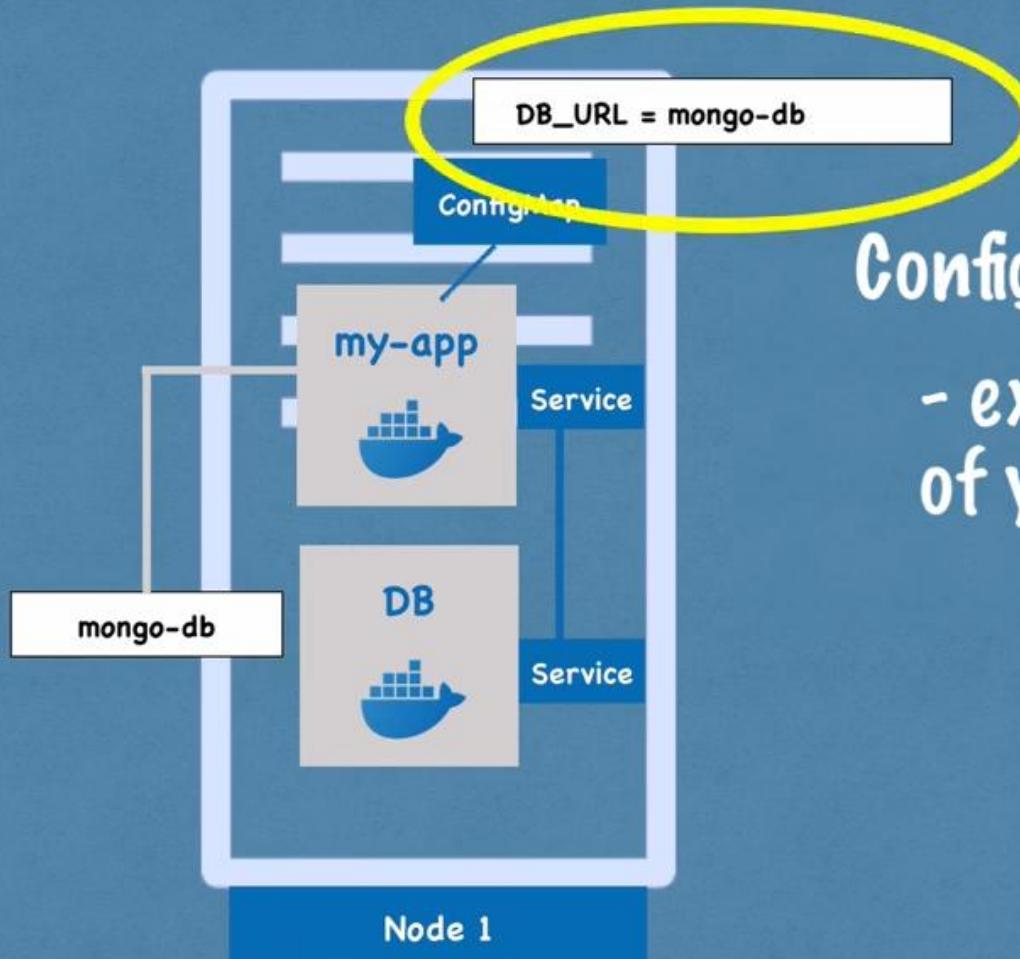
## ConfigMap:

*- external configuration  
of your application*



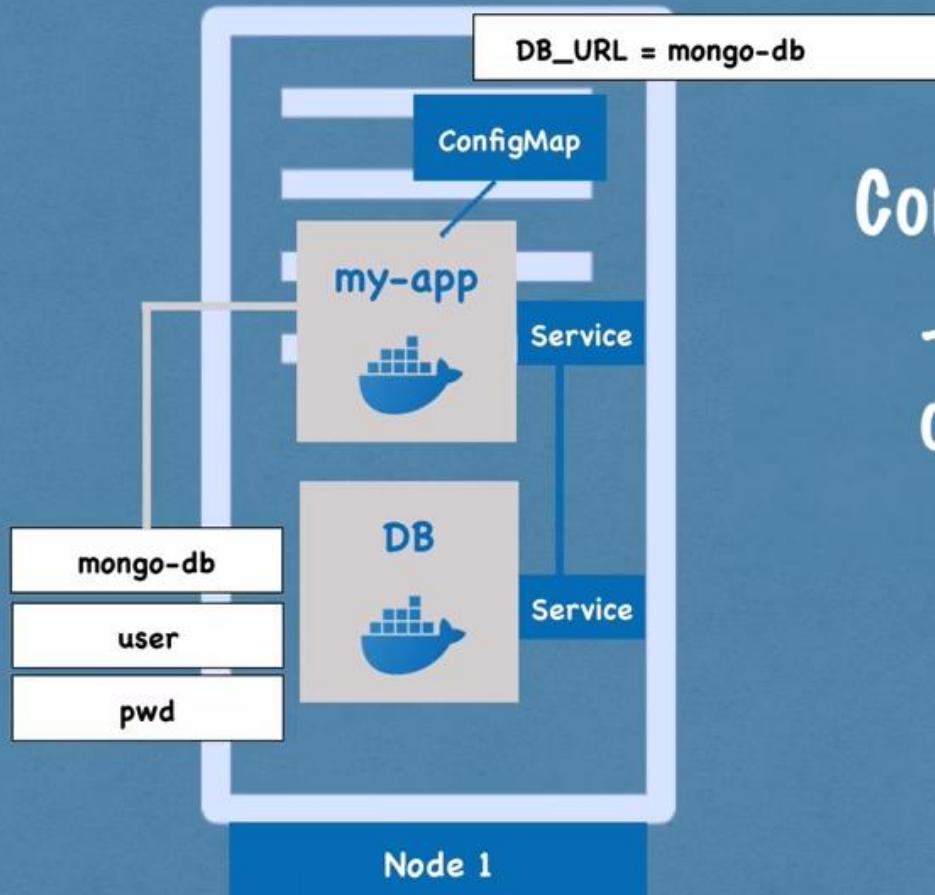
## ConfigMap:

- external configuration  
of your application



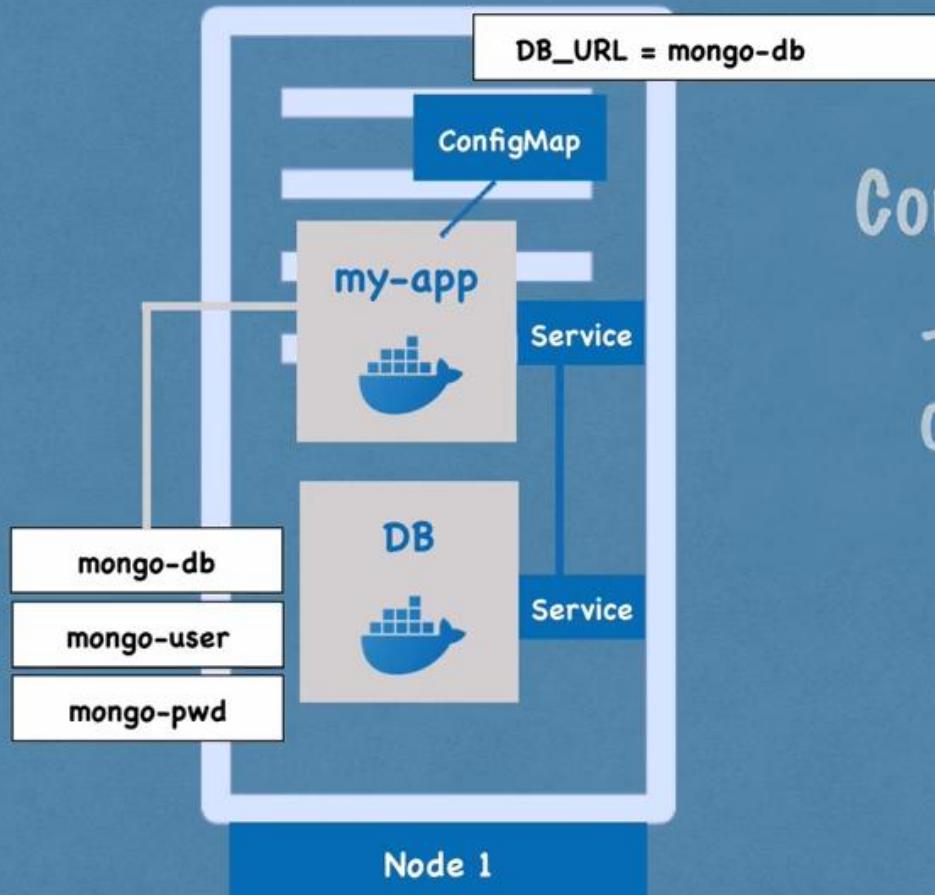
## ConfigMap:

- external configuration  
of your application



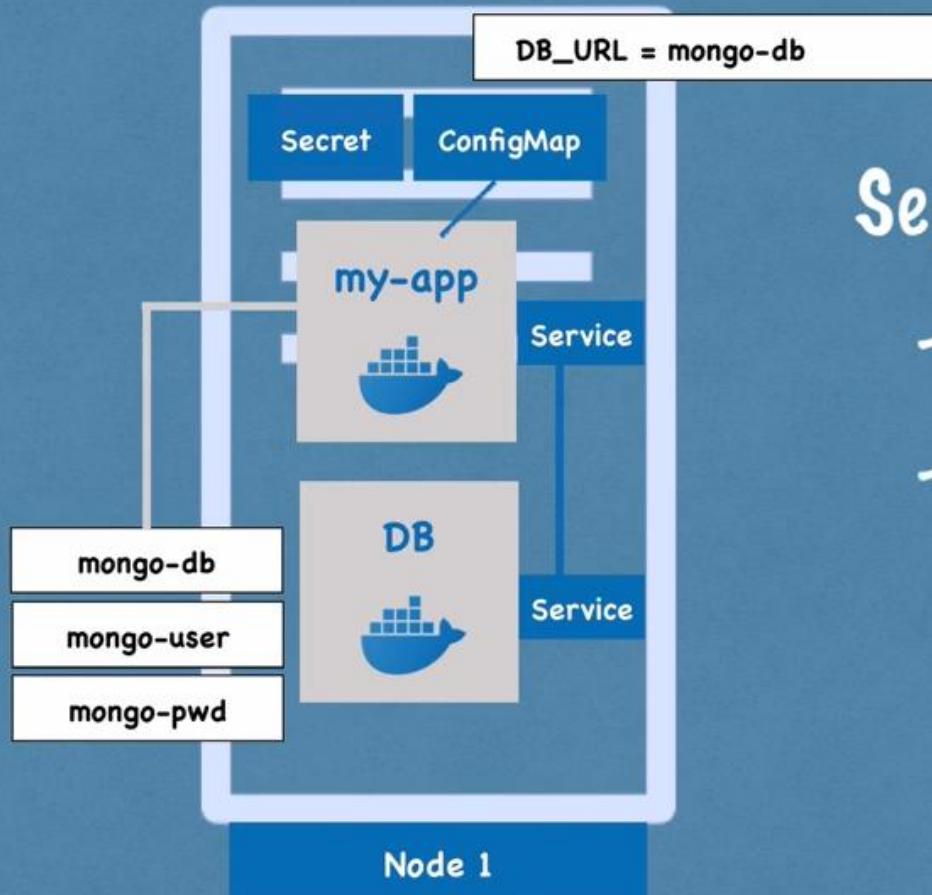
## ConfigMap:

- external configuration  
of your application



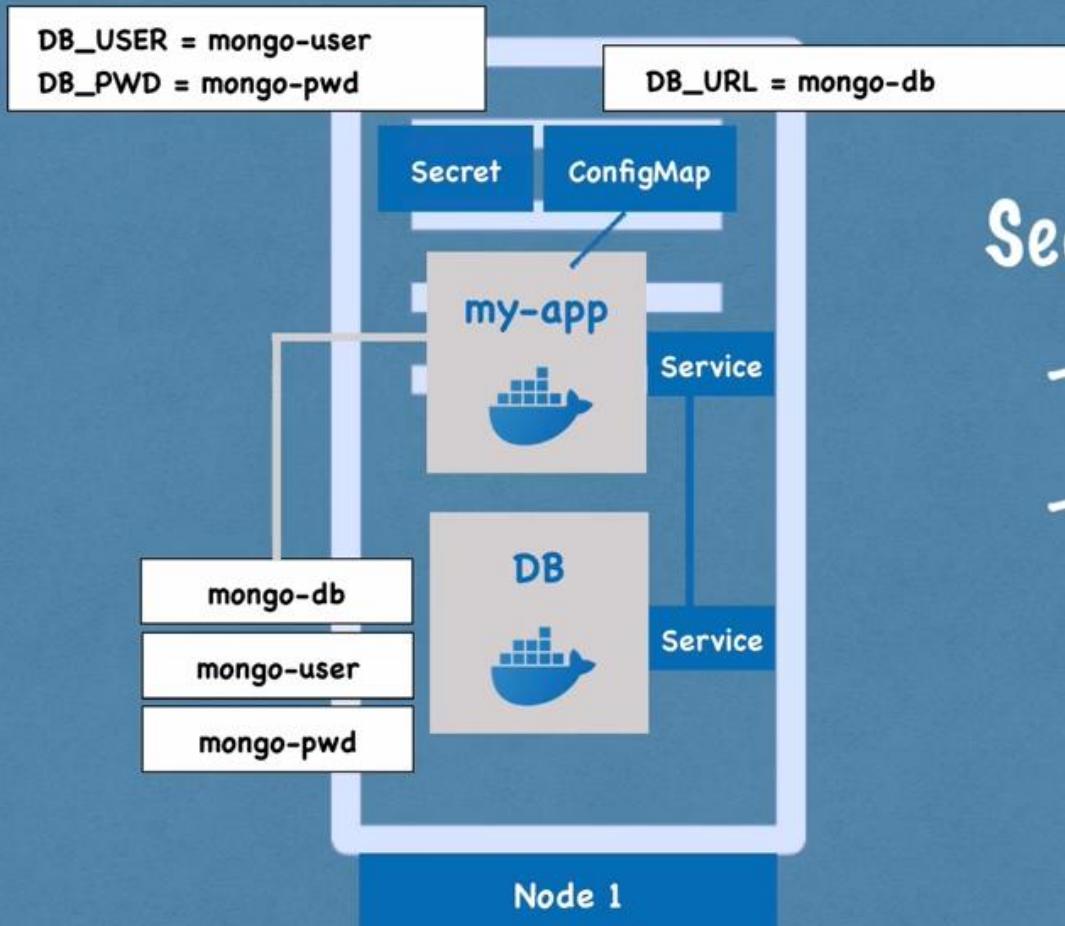
## ConfigMap:

*- external configuration  
of your application*



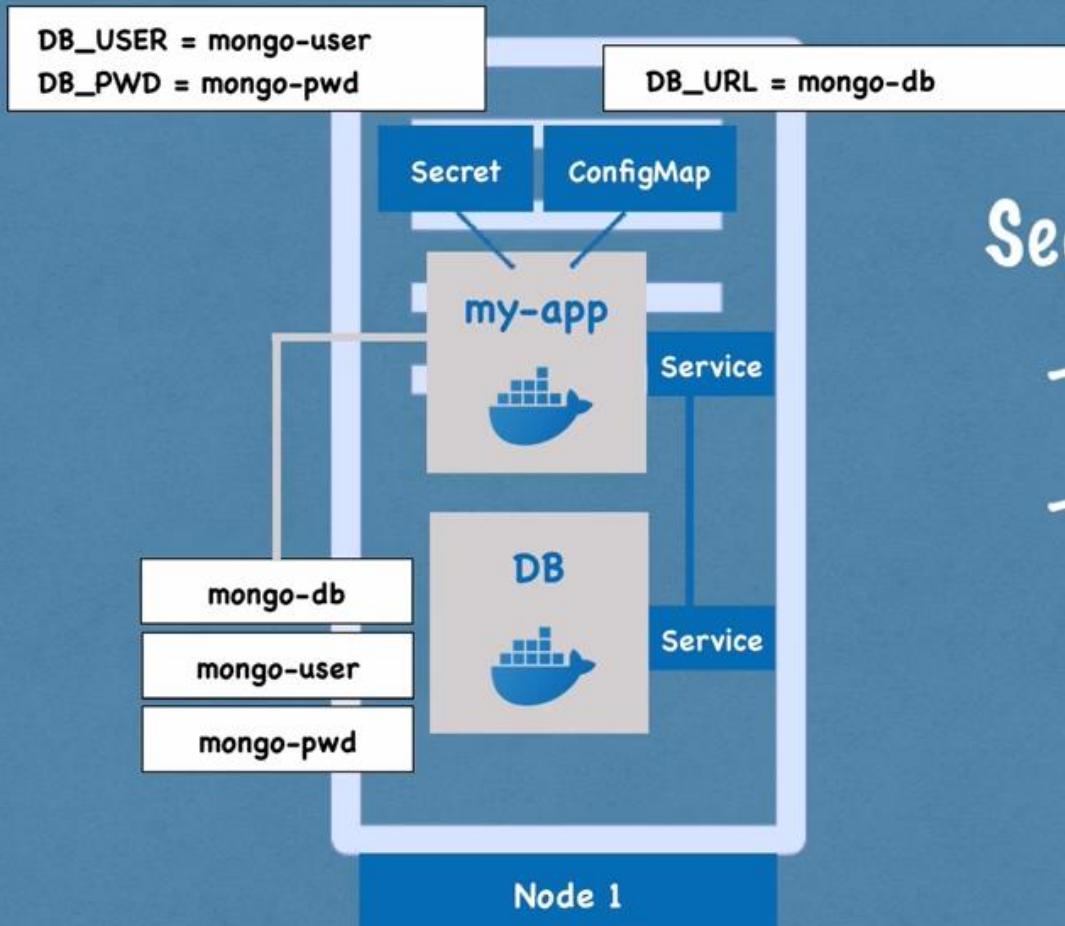
## Secret:

- used to store secret data
- base64 encoded



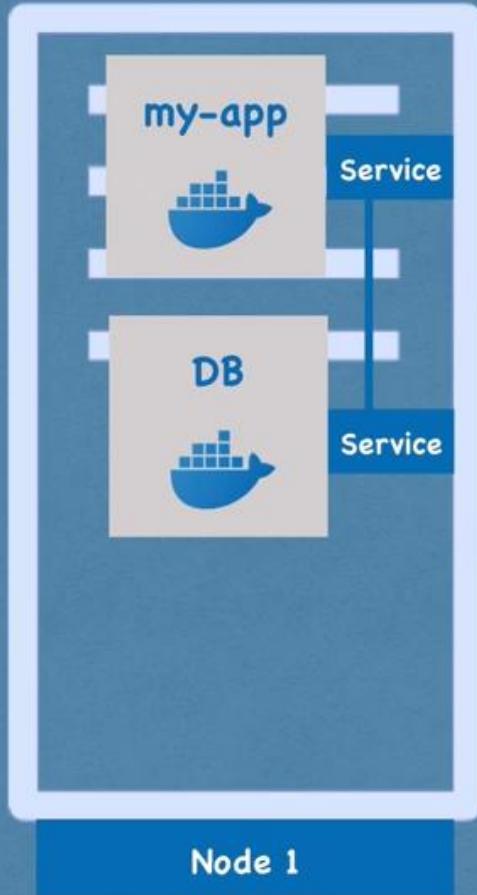
## Secret:

- used to store secret data
- base64 encoded



## Secret:

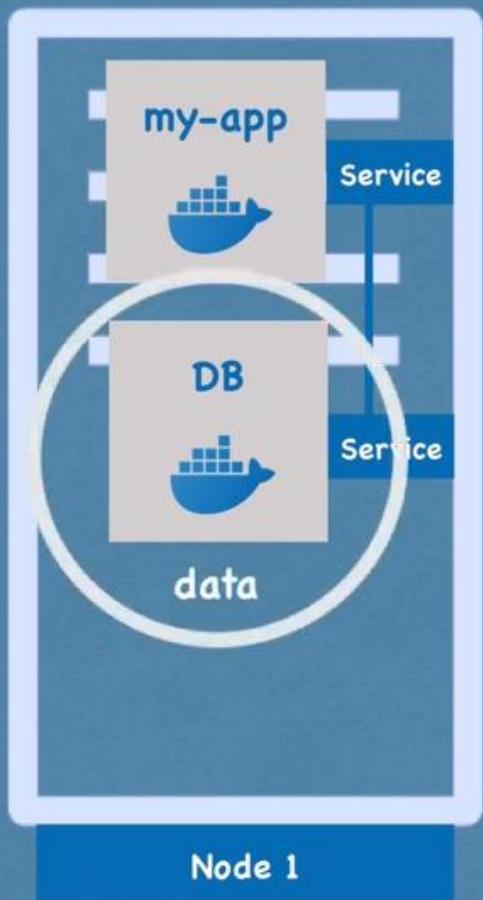
- used to store secret data
- base64 encoded

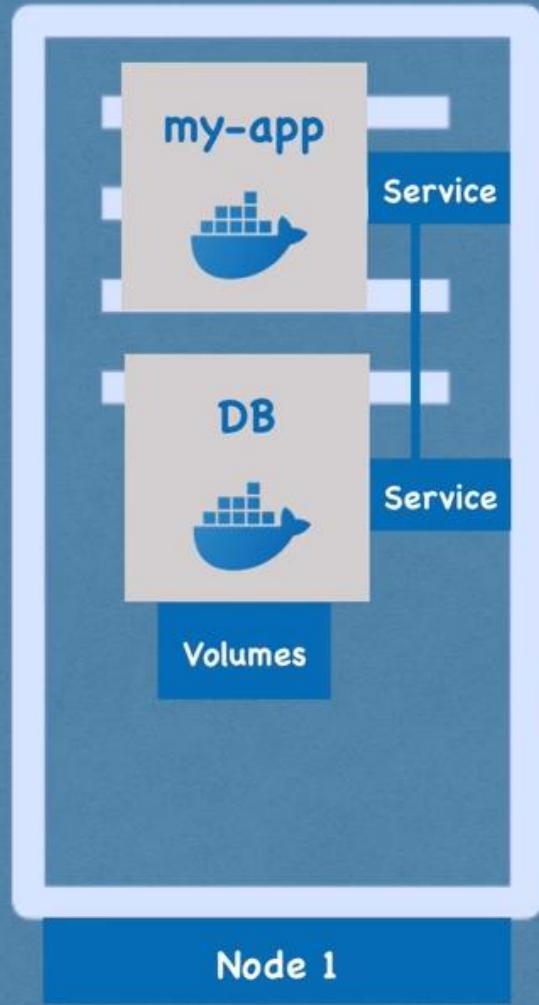


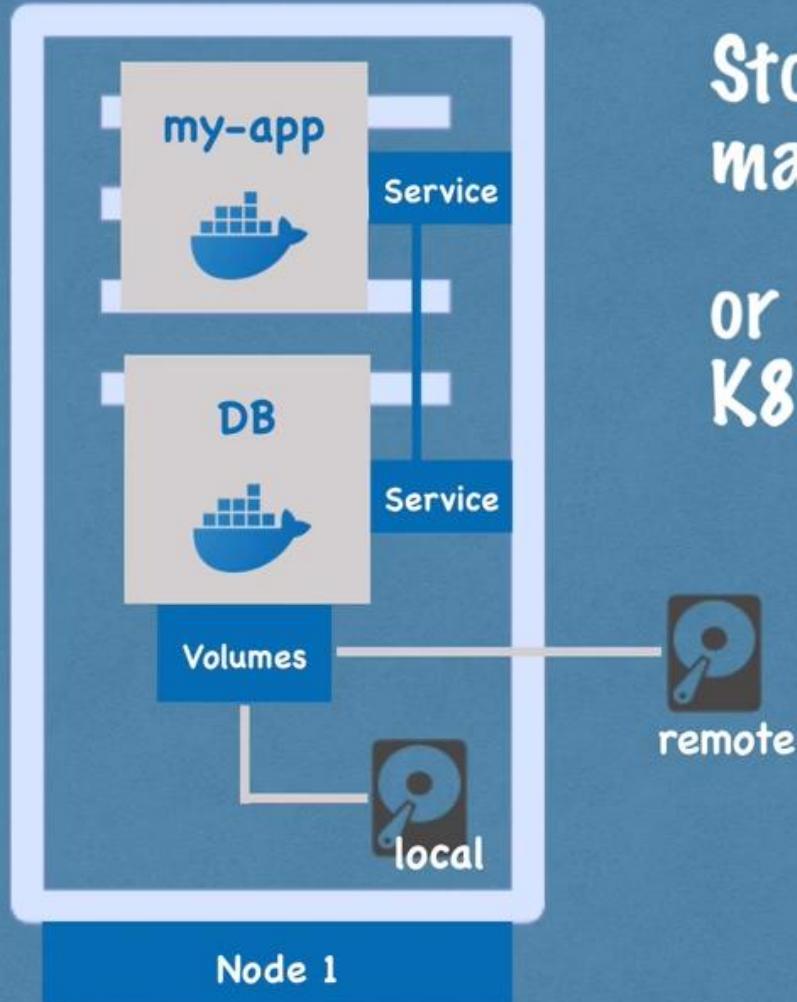
Pod  
Service      Ingress  
ConfigMap      Secrets



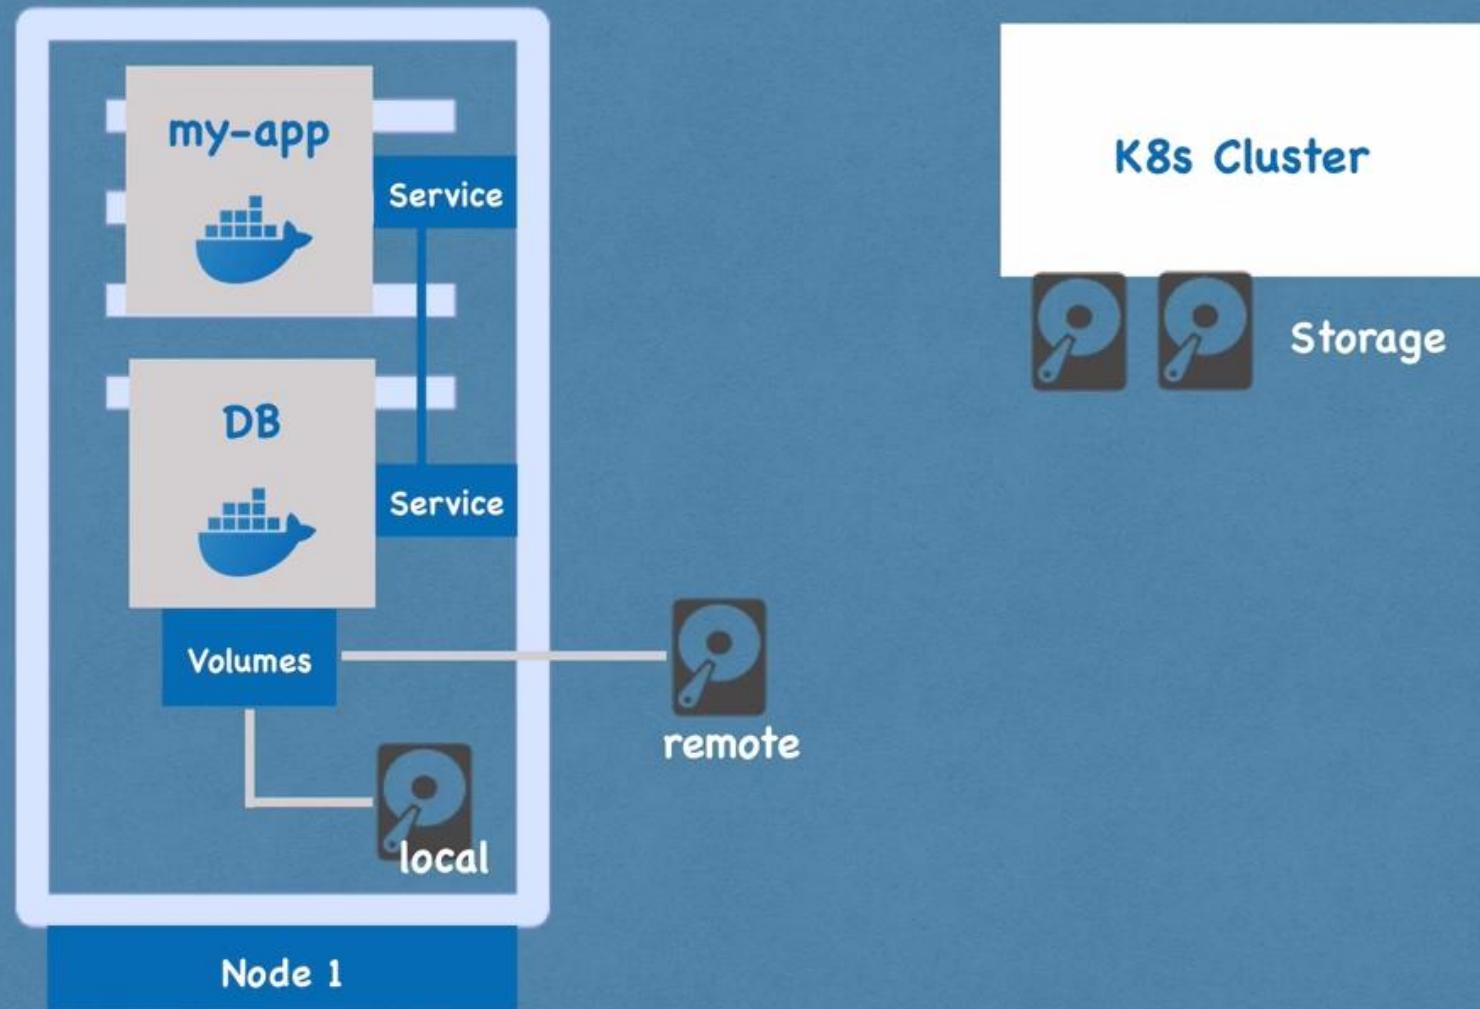
# Volumes

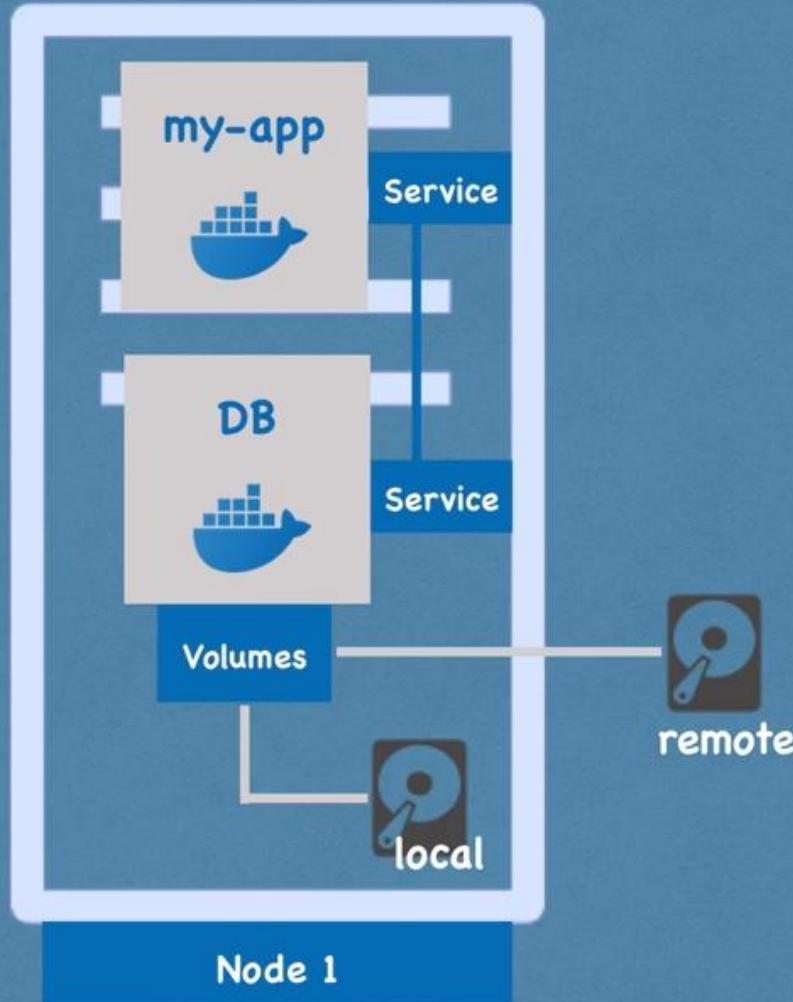






Storage can be on a local machine  
or remote, outside of the K8s cluster



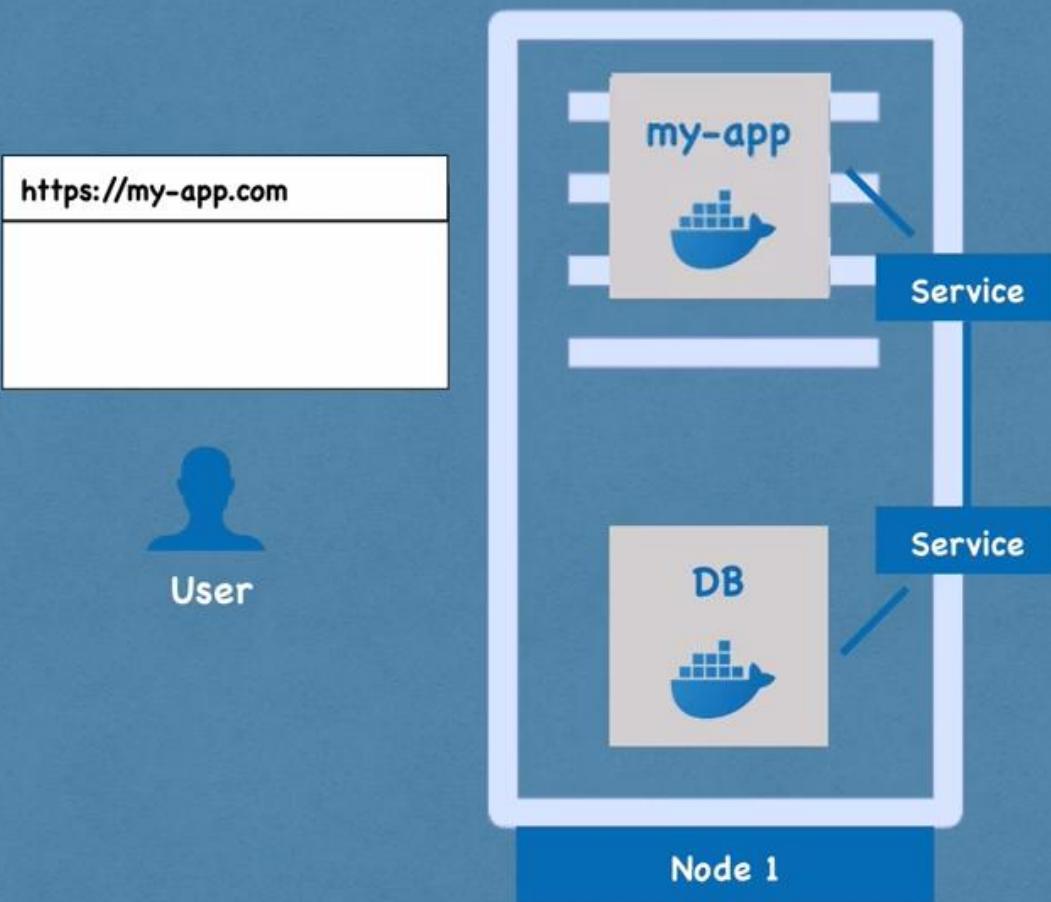


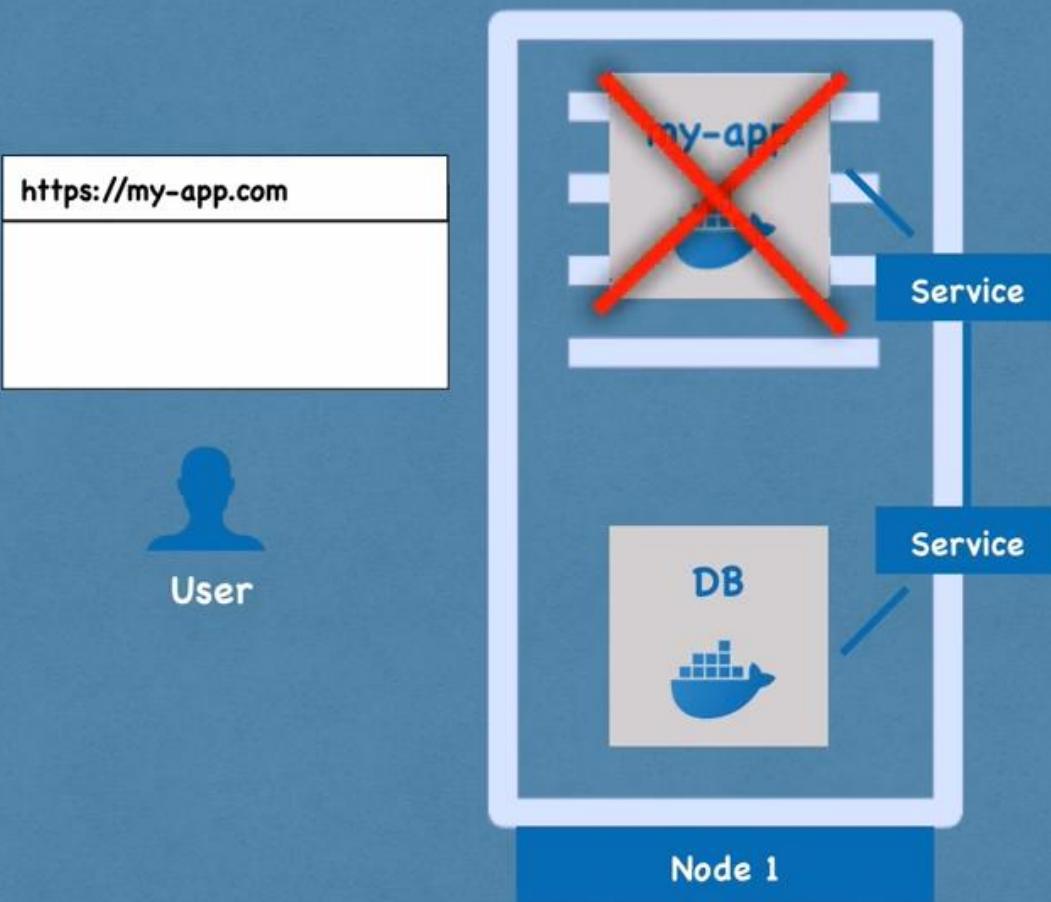
K8s doesn't manage  
data persistance!



# Deployment and Stateful Set

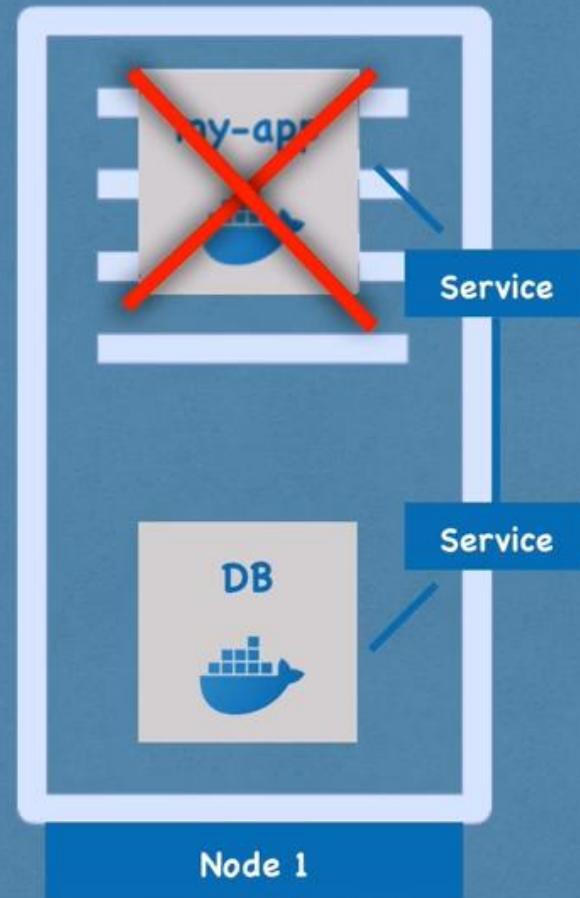
```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: my-app
  labels:
    app: my-app
spec:
  replicas: 2
  selector:
    matchLabels:
      app: my-app
  template:
    metadata:
      labels:
        app: my-app
    spec:
      containers:
        - name: my-app
          image: my-image
          env:
            - name: SOME_ENV
              value: $SOME_ENV
      ports:
        - containerPort: 8080
```





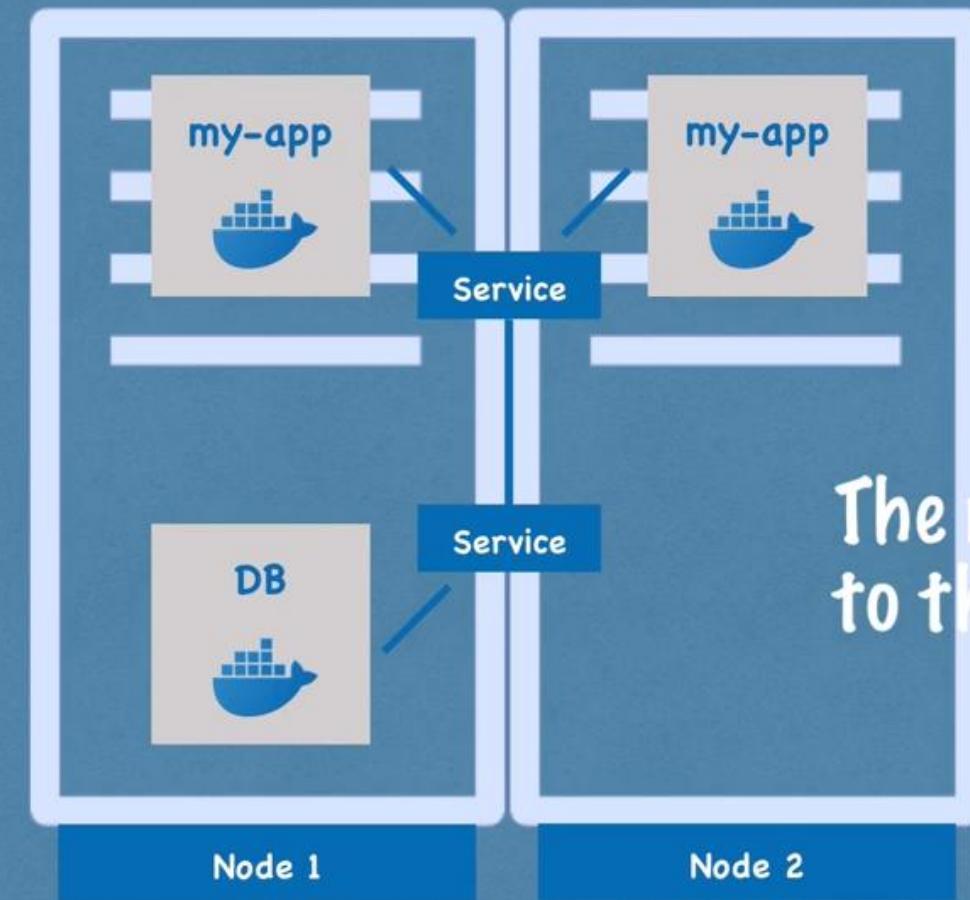


User





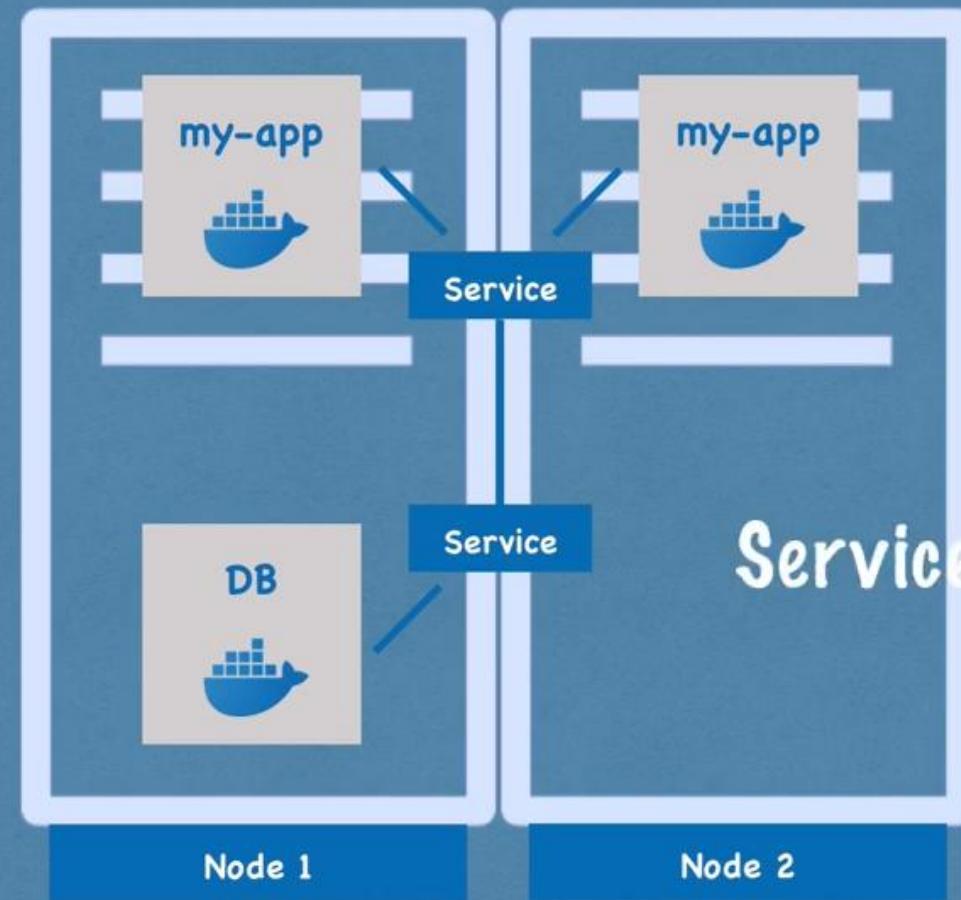
User



The replica is connected  
to the same Service



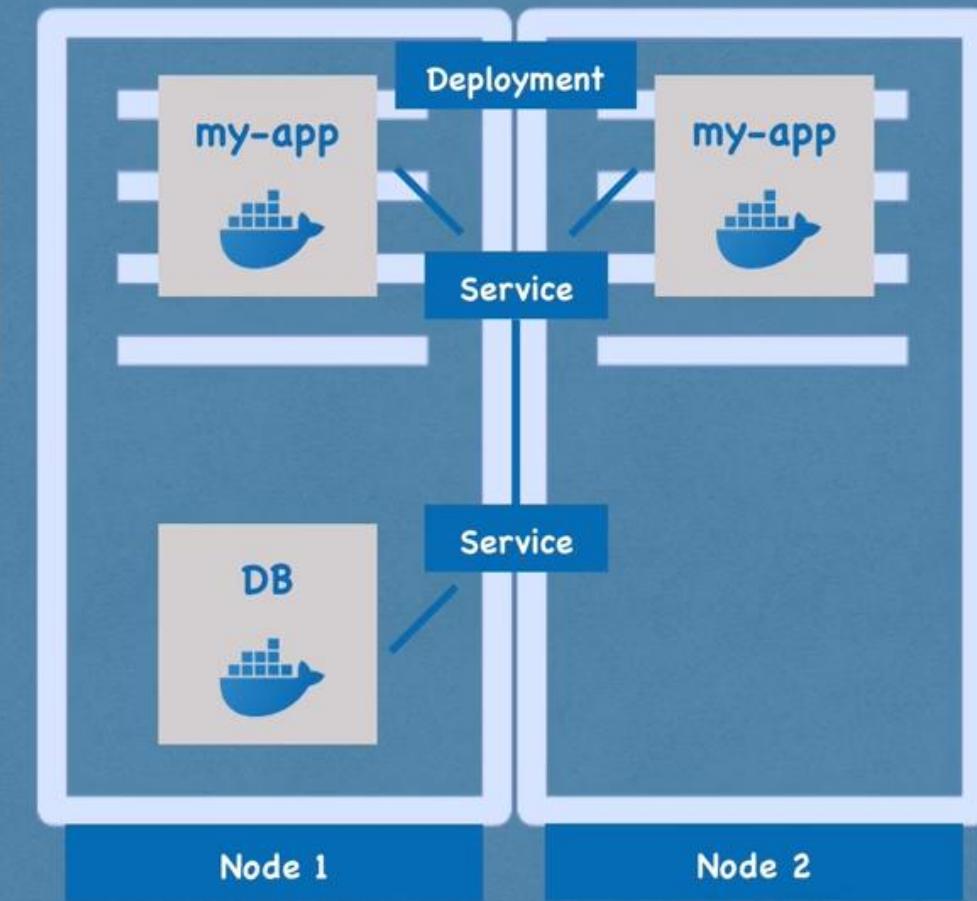
User



Service is a load balancer!

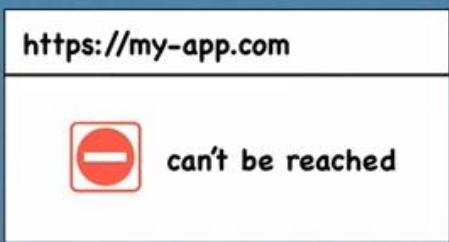


User

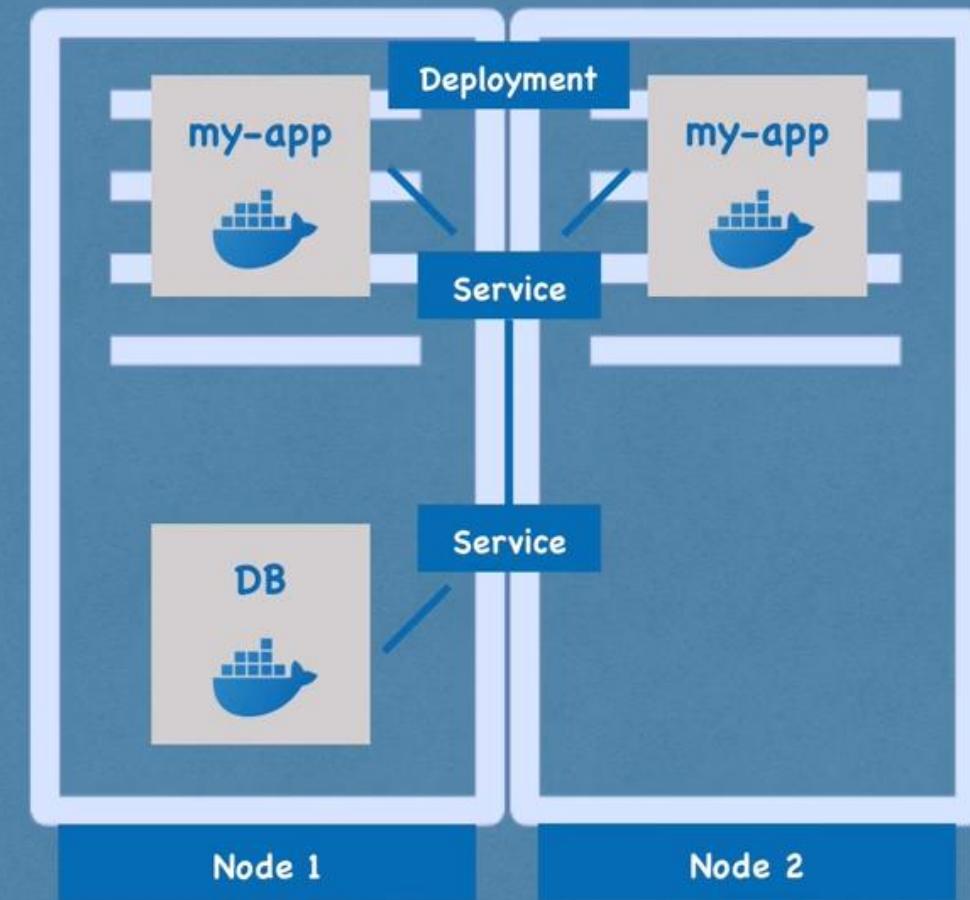


## Deployment:

- blueprint for  
my-app Pods

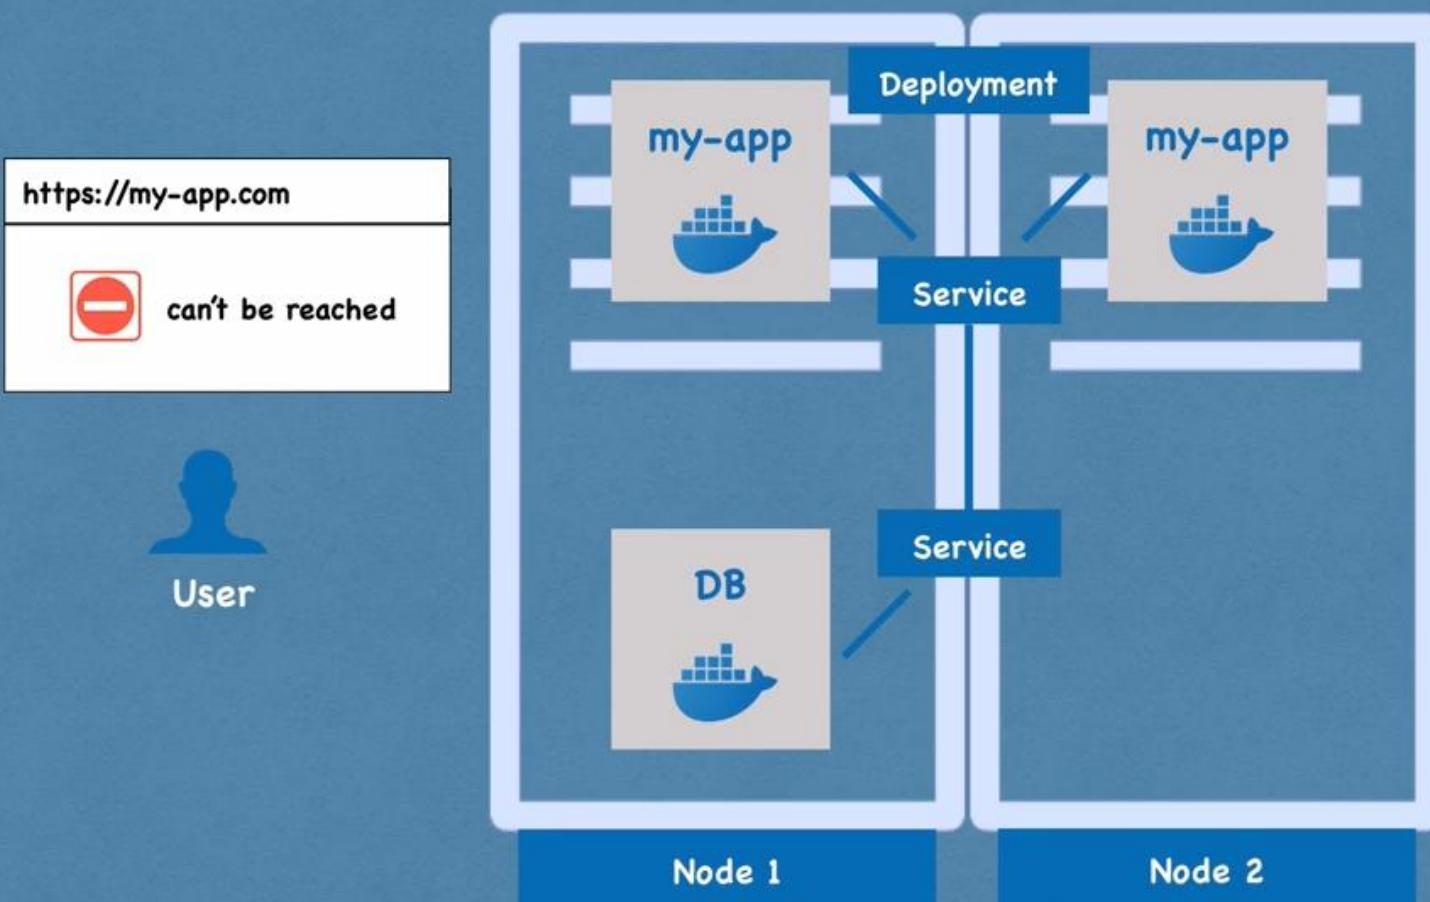


User



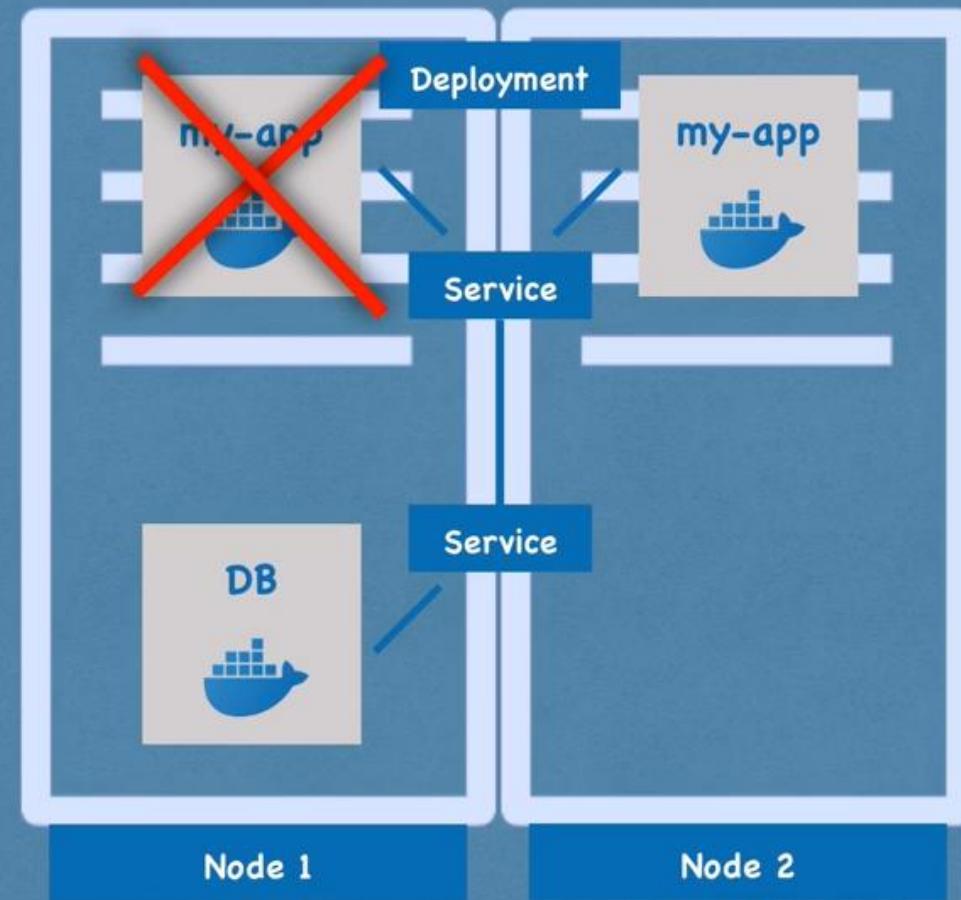
## Deployment:

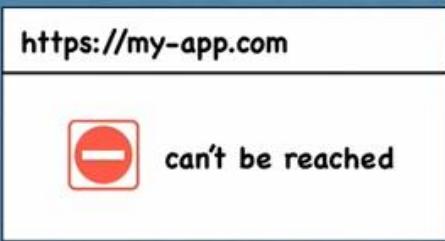
- blueprint for my-app Pods
- abstraction of Pods



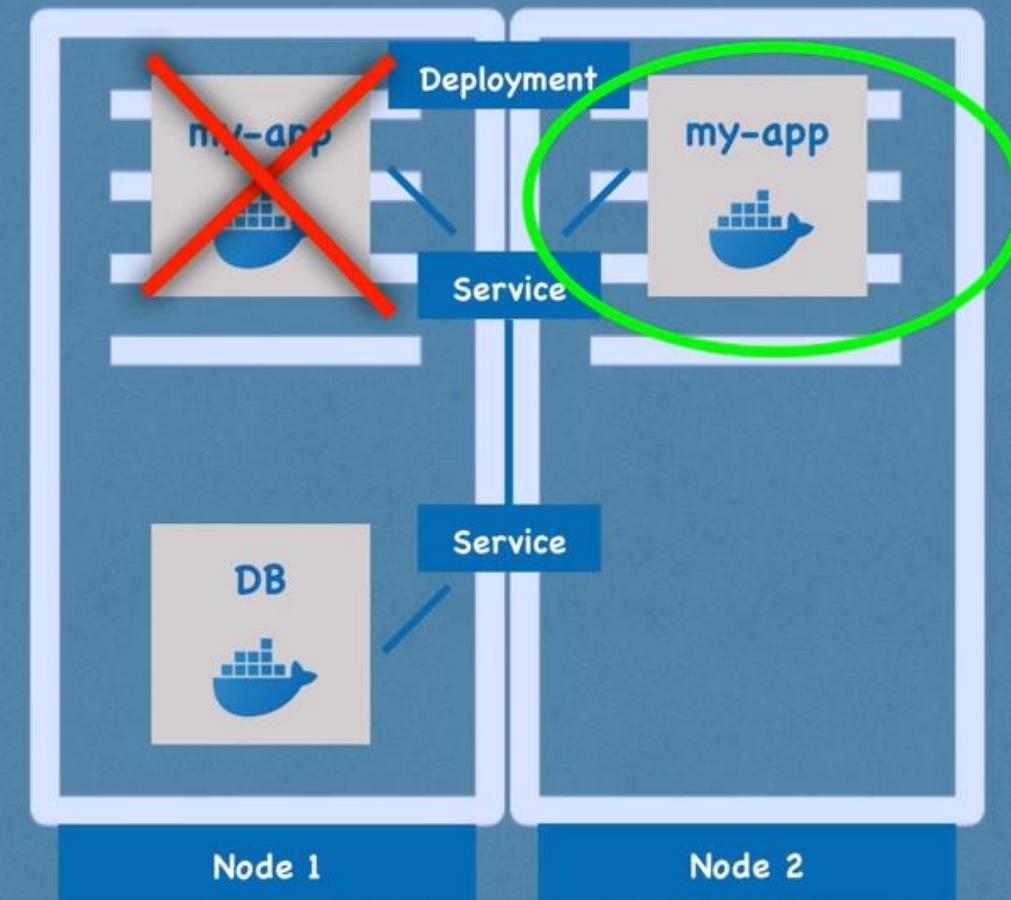


User



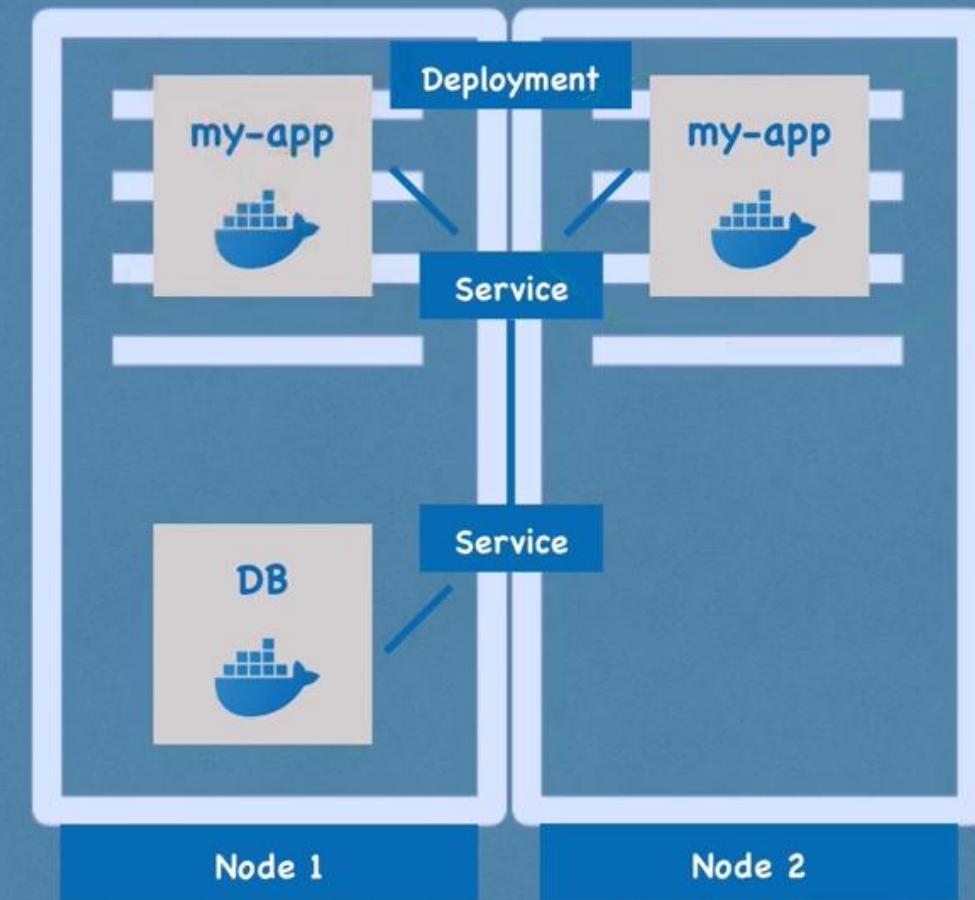


User





User

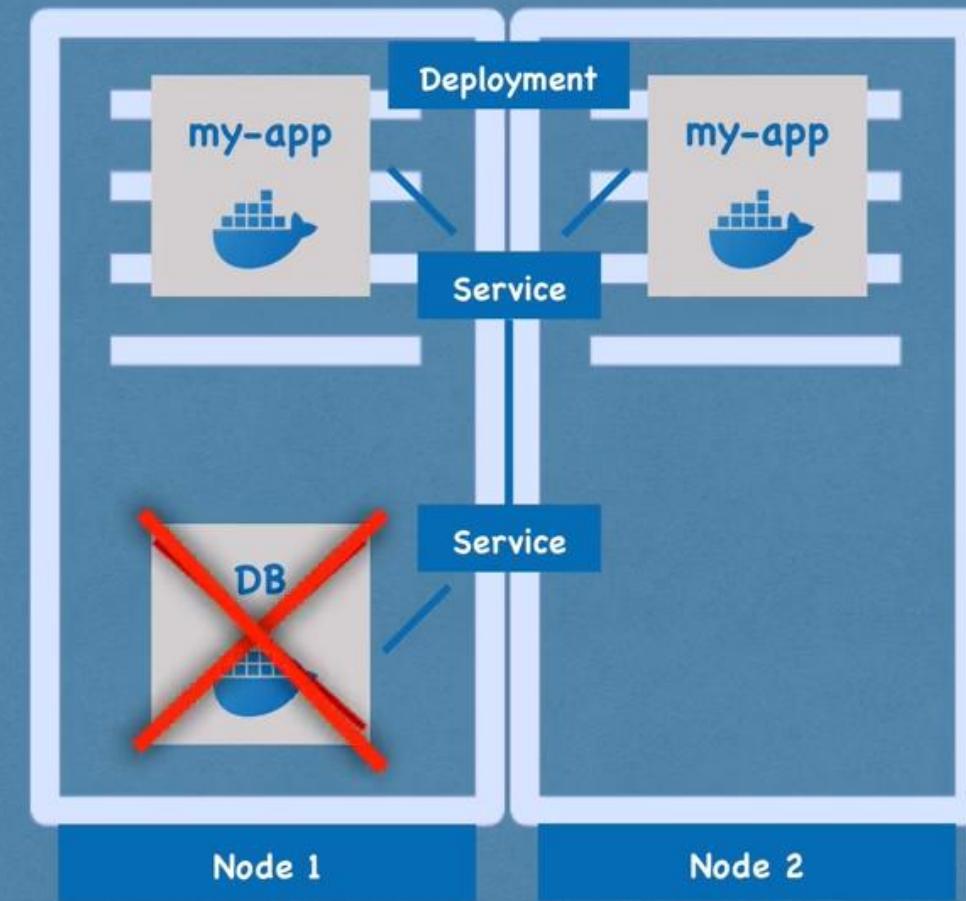


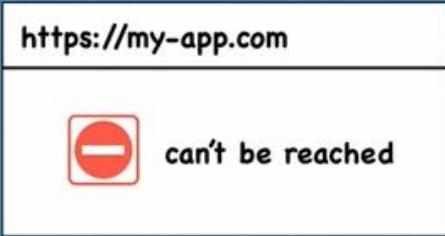
<https://my-app.com>

 can't be reached

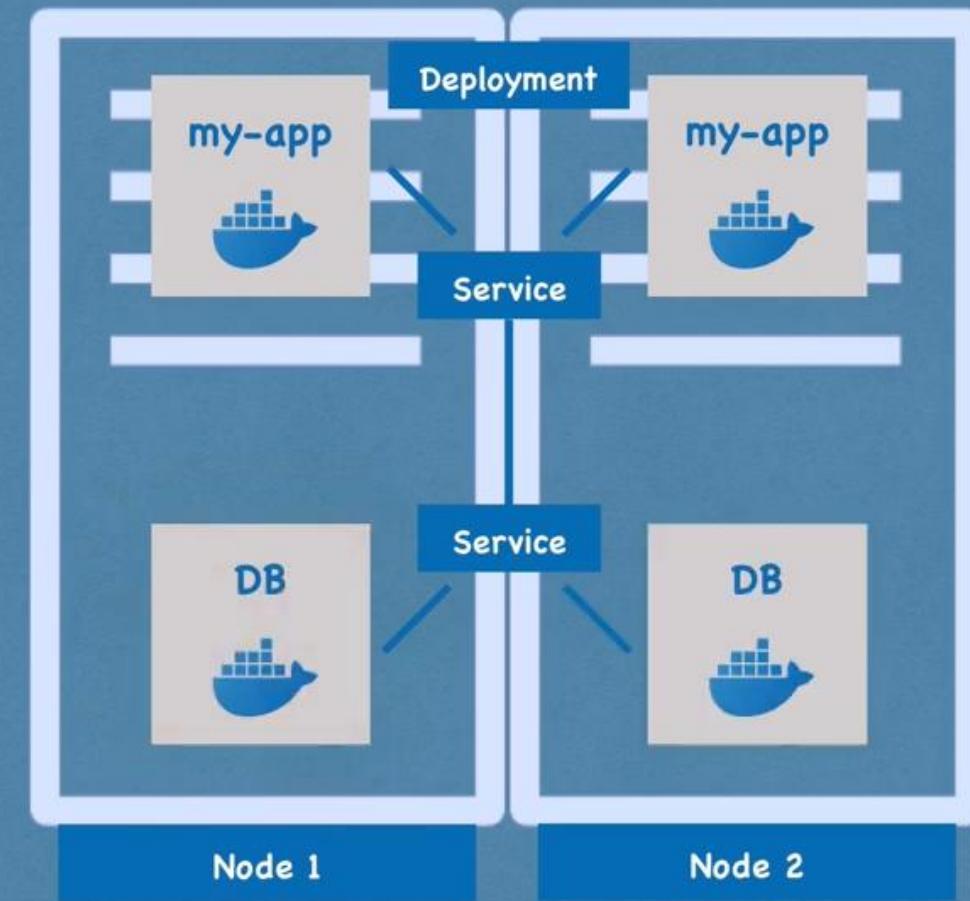


User

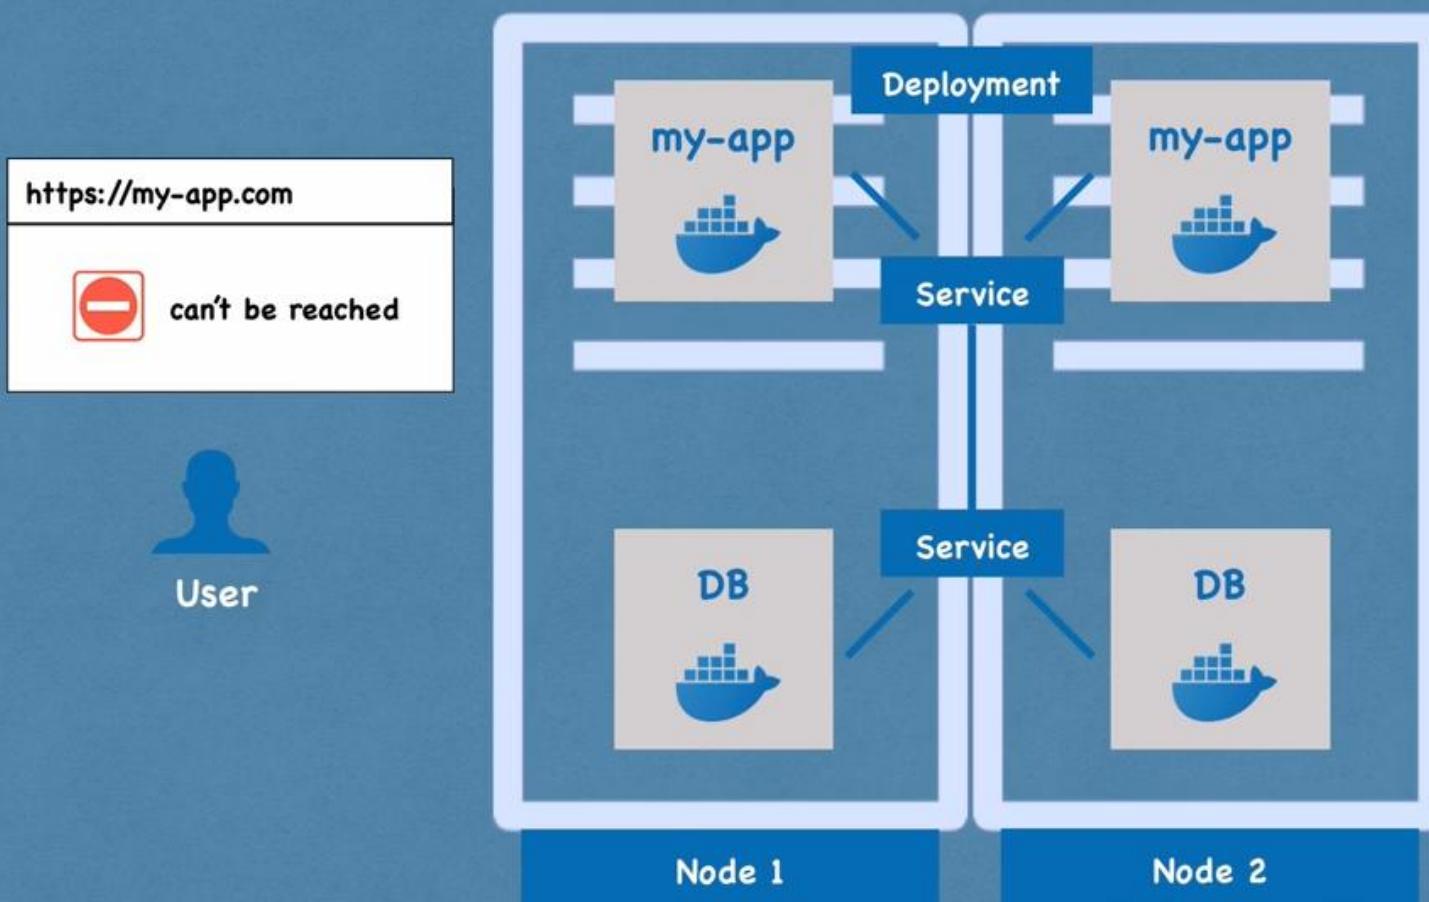




User



# DB can't be replicated via Deployment!

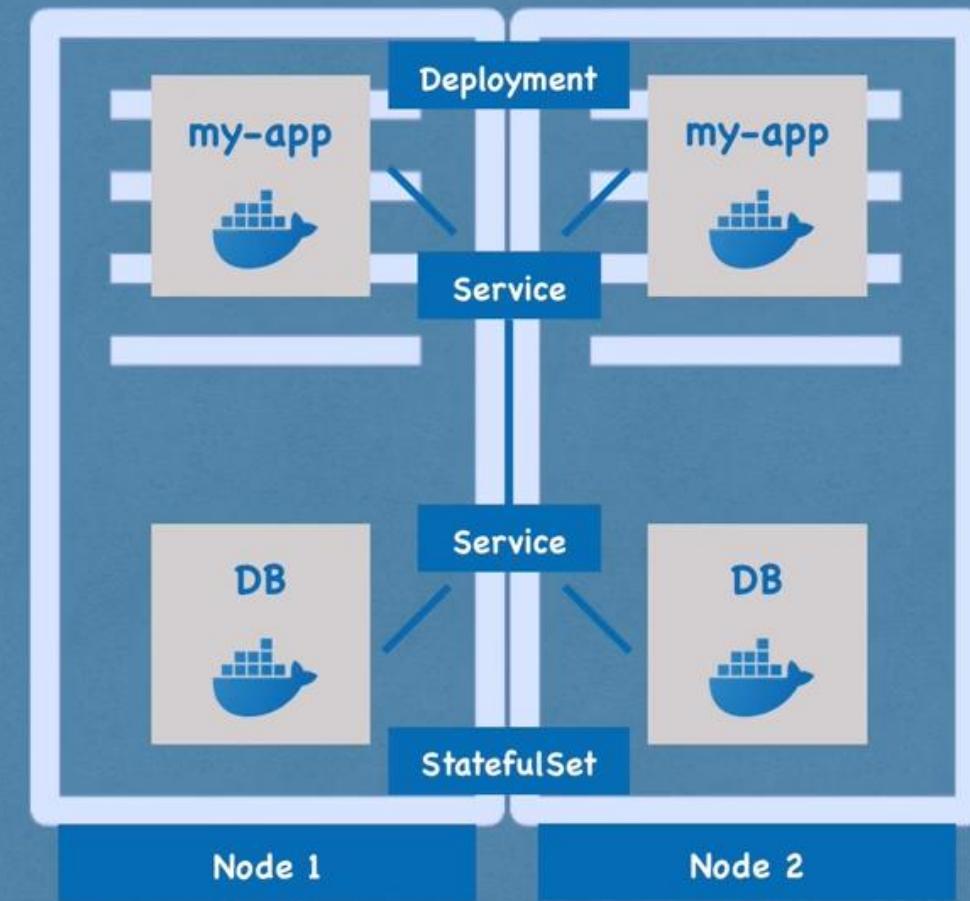


<https://my-app.com>

 can't be reached



User

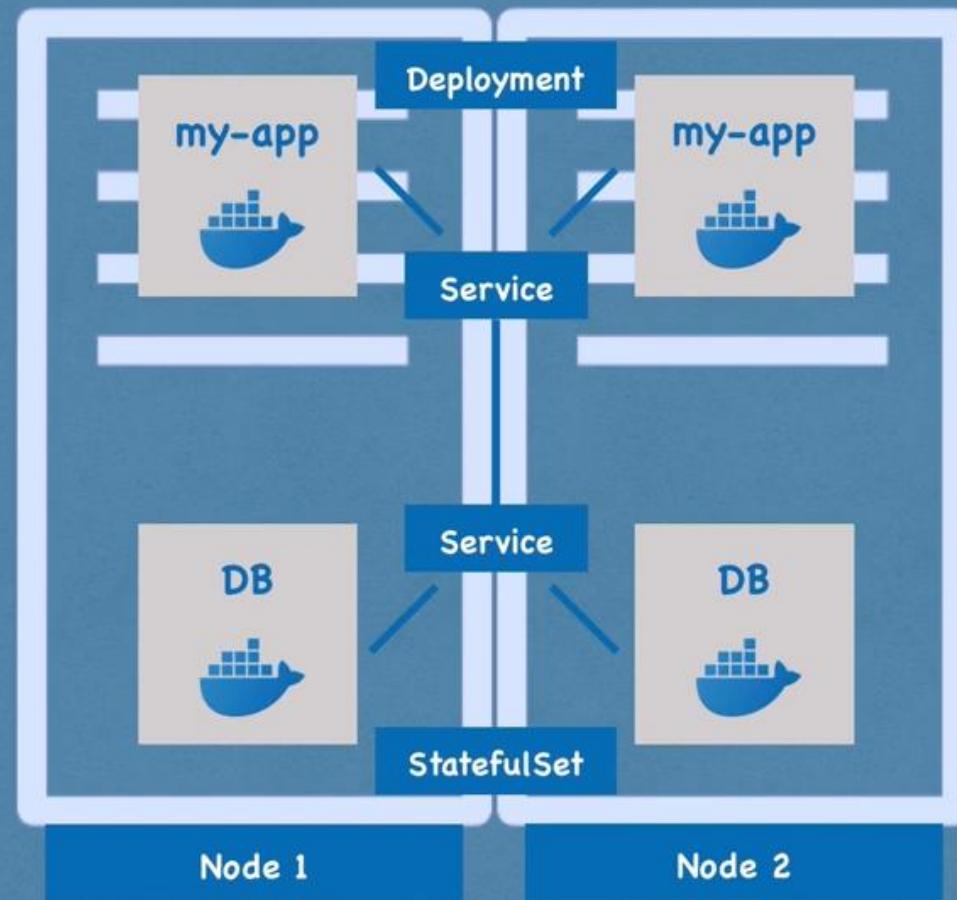


<https://my-app.com>

 can't be reached



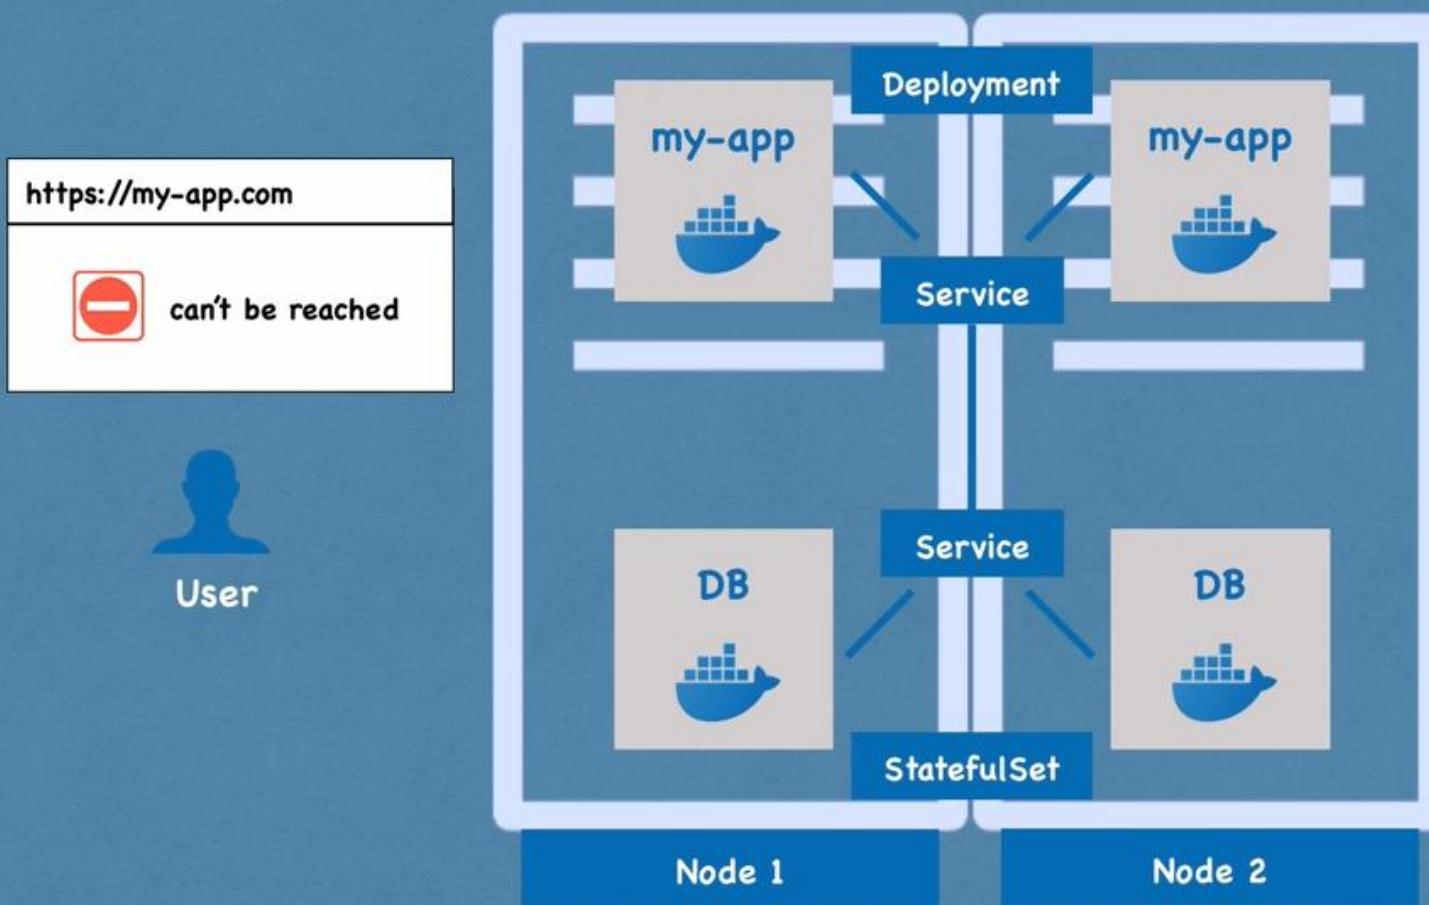
User

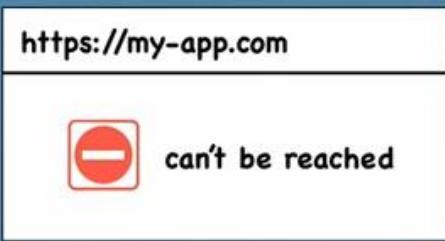


**Deployment for  
stateless Apps**

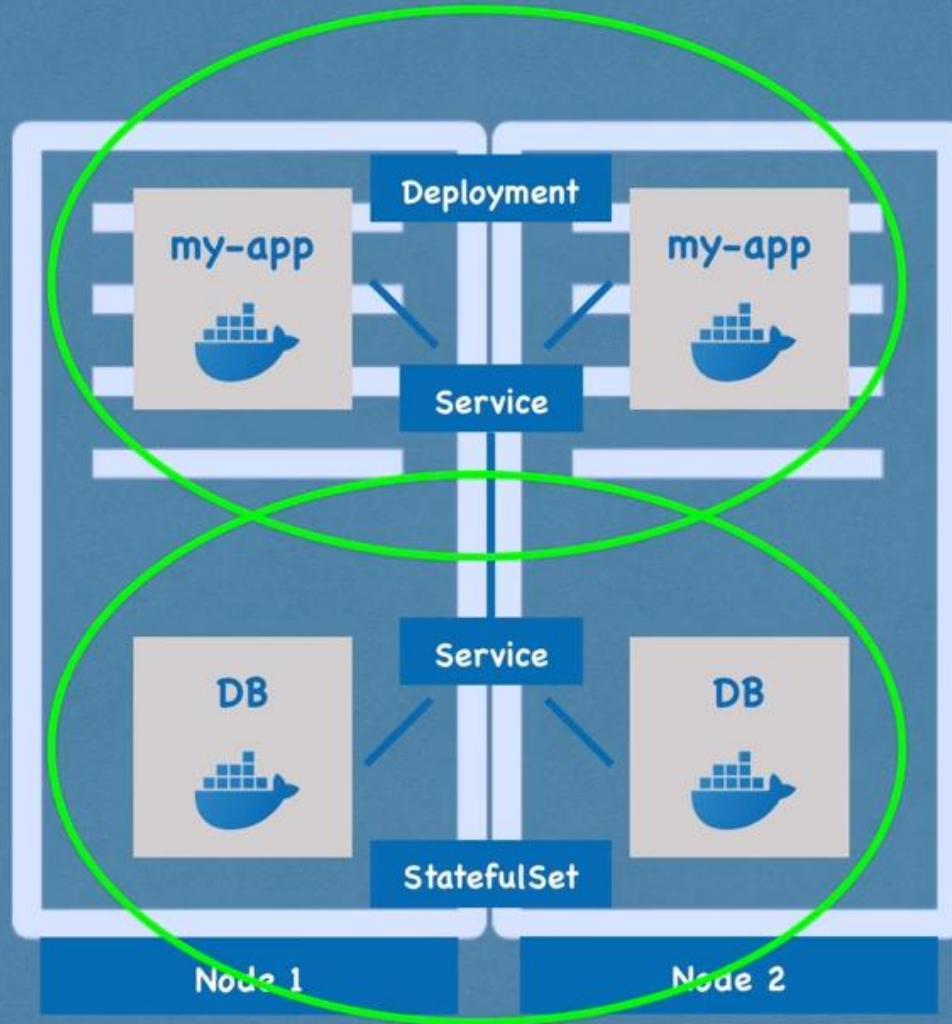
**StatefulSet for  
stateful Apps  
or Databases**

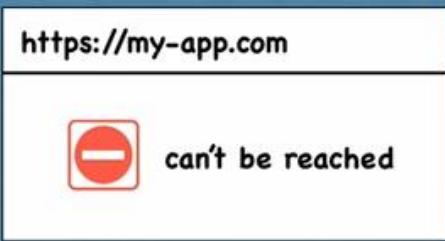
DB are often hosted outside of the K8s cluster



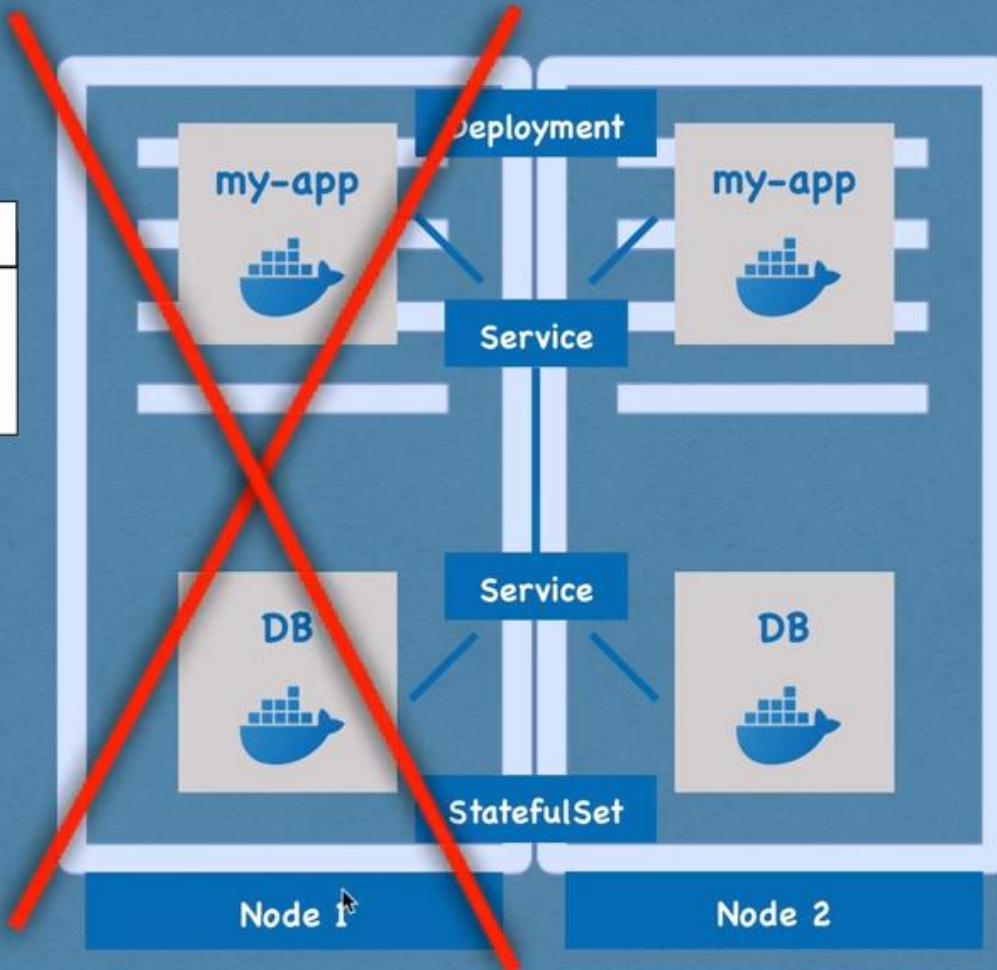


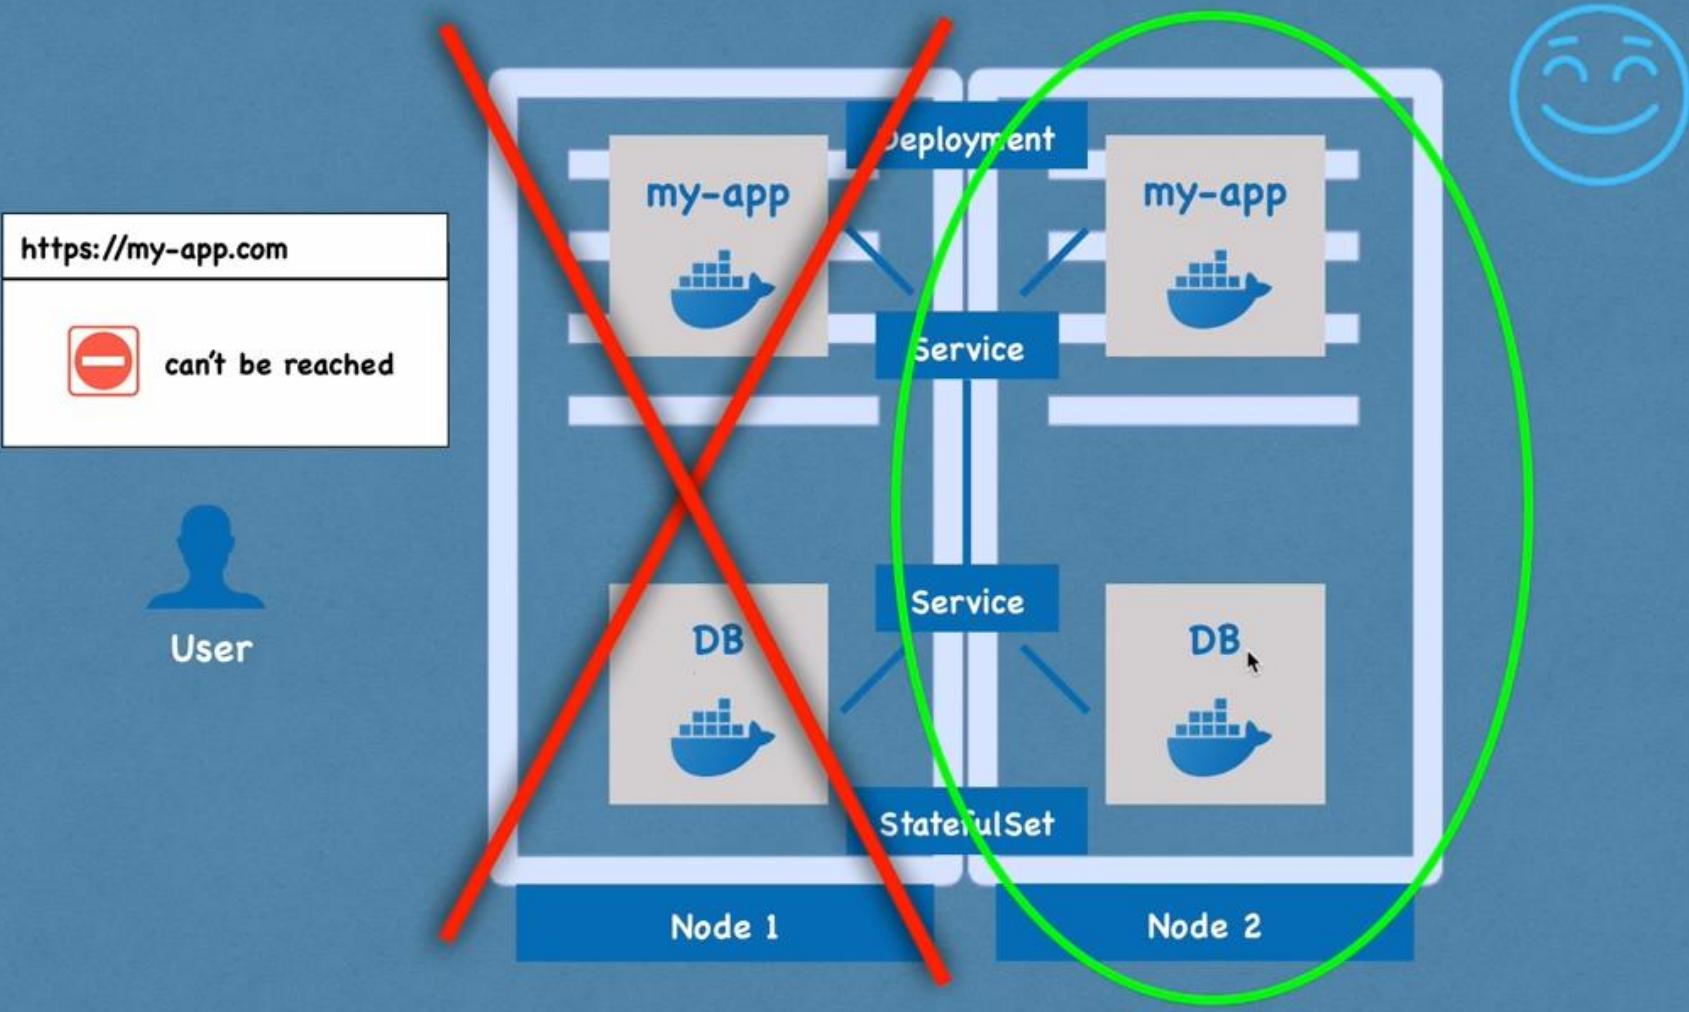
User





User







# Main Kubernetes Components summarized

Pod  
Service   Ingress

Volumes

ConfigMap

Secrets

Deployment

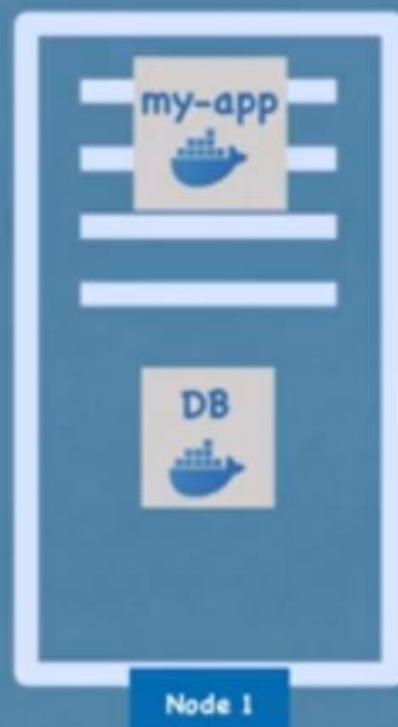
StatefulSet



# Node processes



# Worker machine in K8s cluster





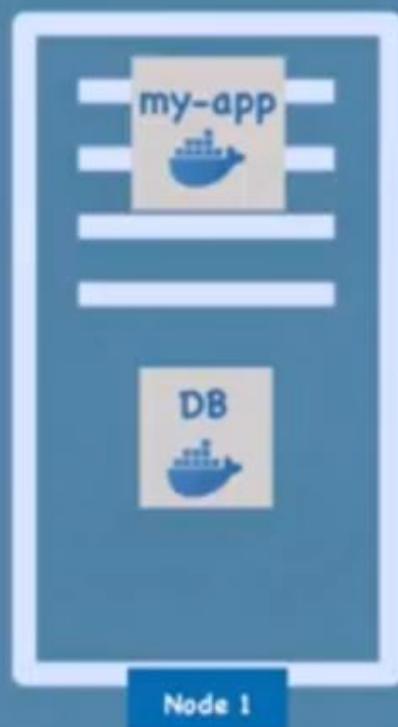
## Worker machine in K8s cluster



- each Node has multiple Pods on it



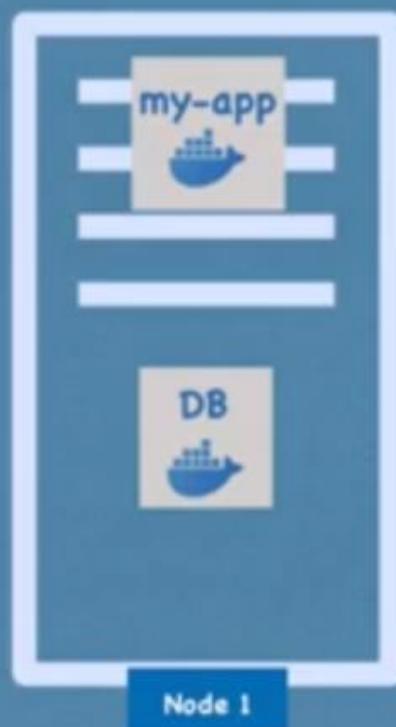
## Worker machine in K8s cluster



- each Node has multiple Pods on it
- 3 processes must be installed on every Node



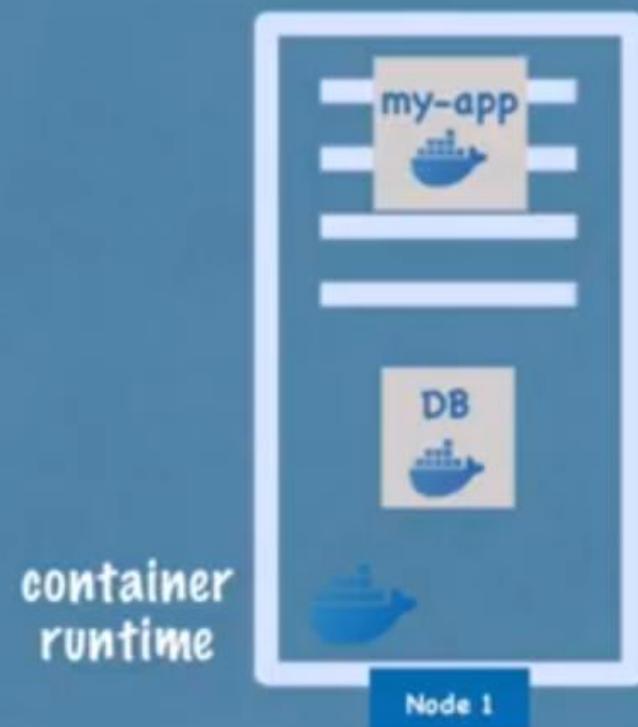
## Worker machine in K8s cluster



- each Node has multiple Pods on it
- 3 processes must be installed on every Node
- Worker Nodes do the actual work

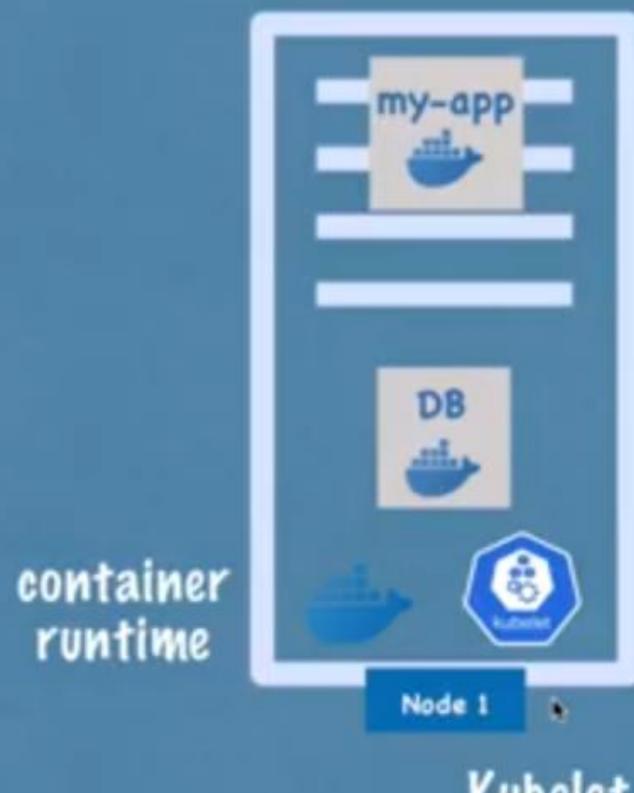


# Worker machine in K8s cluster





## Worker machine in K8s cluster



Kubelet interacts with both - the container and node

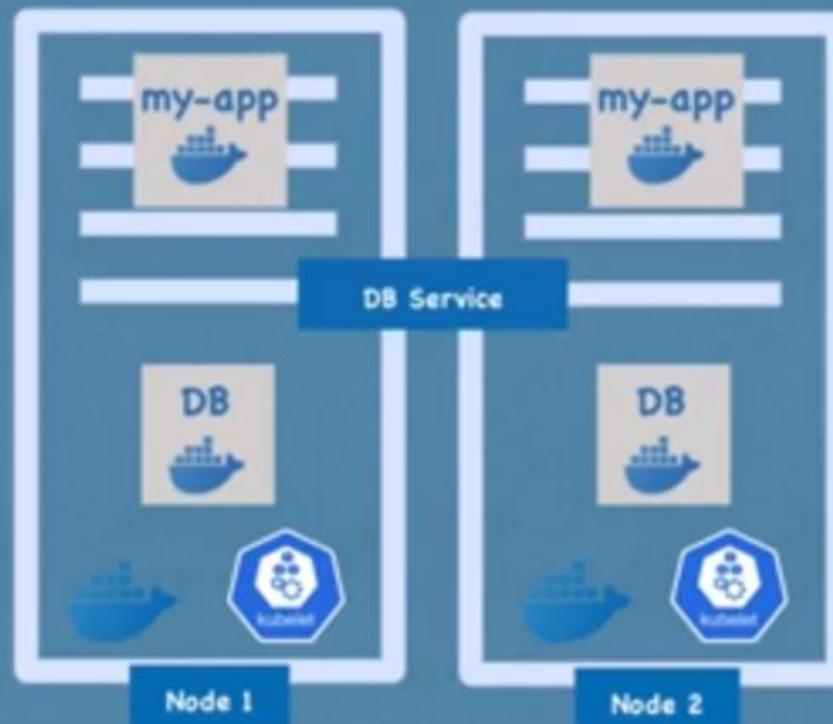
Kubelet starts the pod with a container inside



# Worker machine in K8s cluster

Communication  
via Services

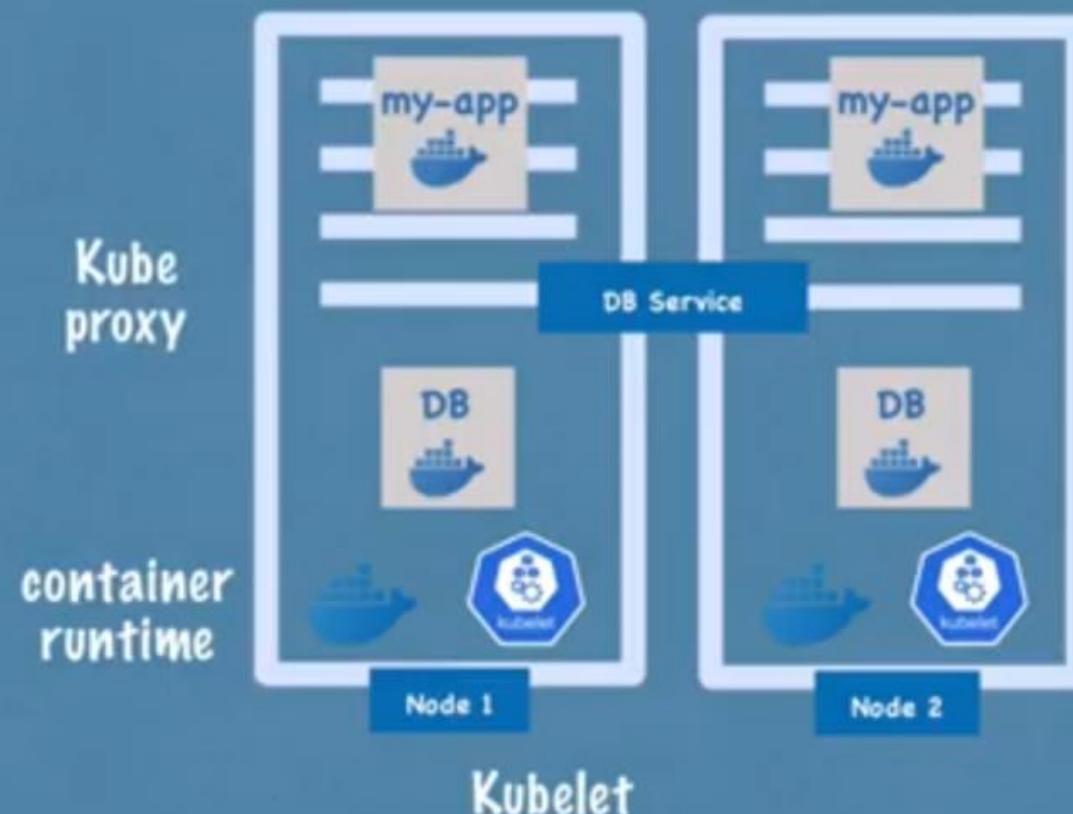
container  
runtime



Kubelet



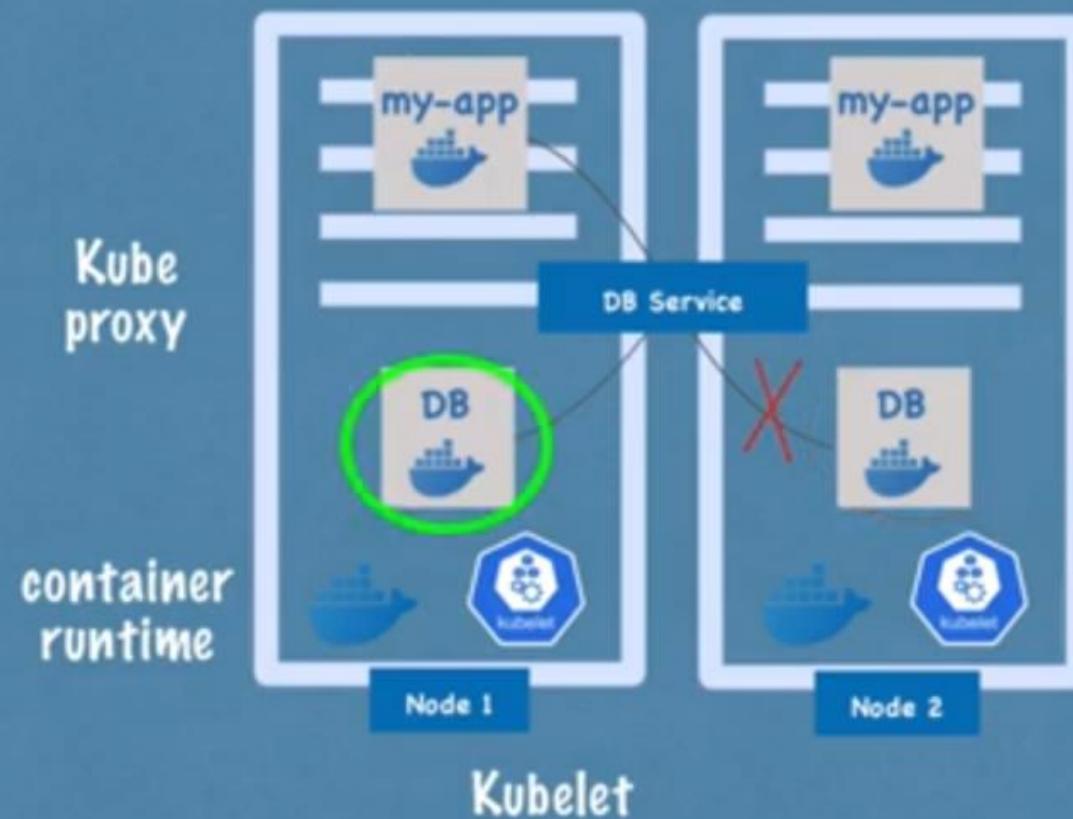
## Worker machine in K8s cluster



Kube Proxy  
forwards the  
requests

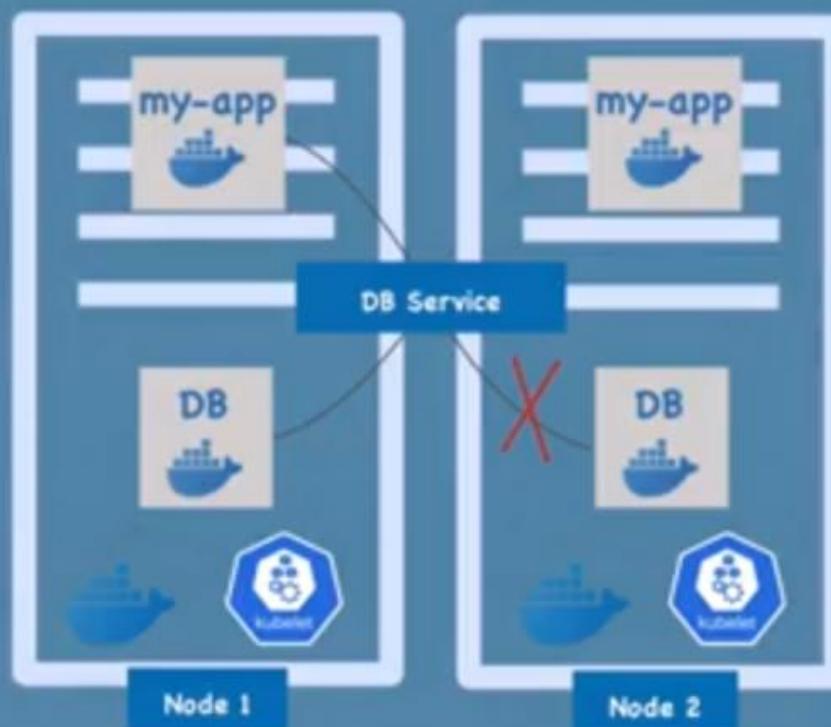


# Worker machine in K8s cluster





# Worker machine in K8s cluster



3 node processes:

1. Kubelet

2. Kube Proxy

3. Container runtime

So, how do you interact with this cluster?

How to:

- > schedule pod?
- > monitor?
- > re-schedule/re-start pod?
- > join a new node?



Managing processes are done by  
Master Nodes

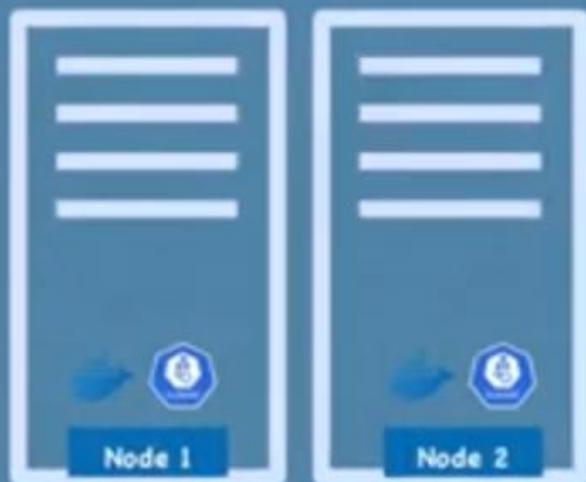




# Master processes

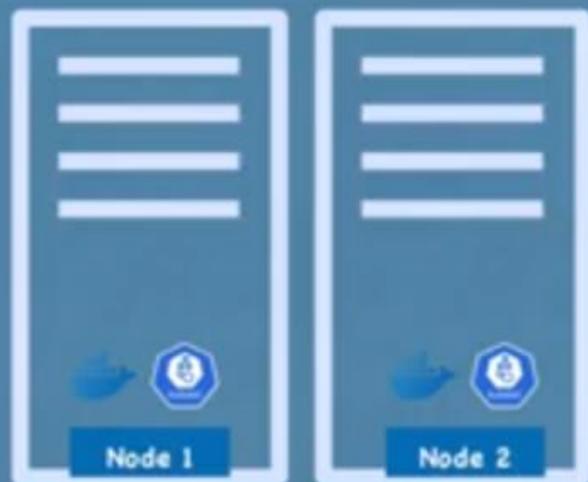


## Master processes





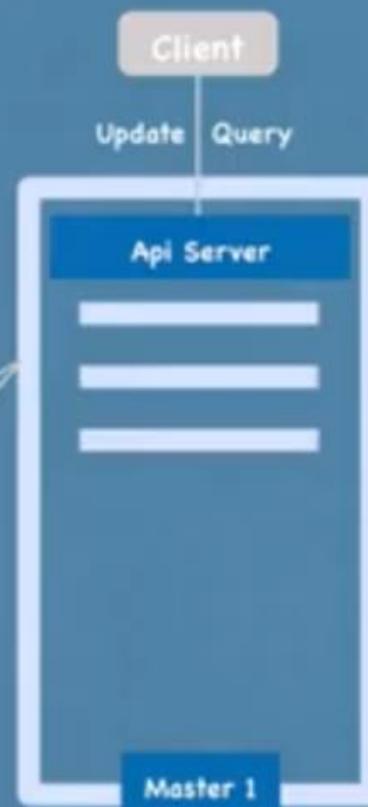
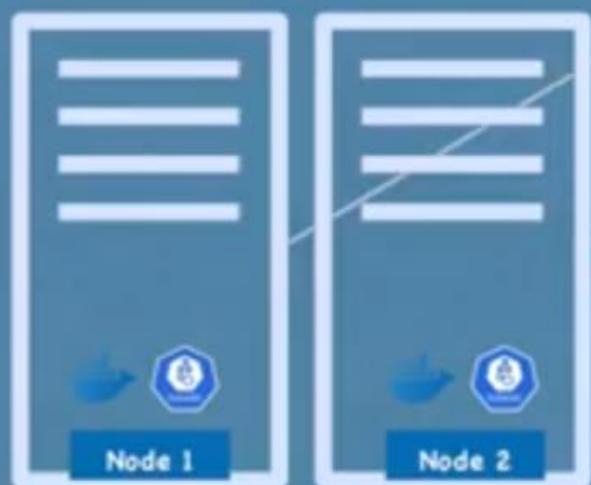
## Master processes



4 processes run on every  
master node!



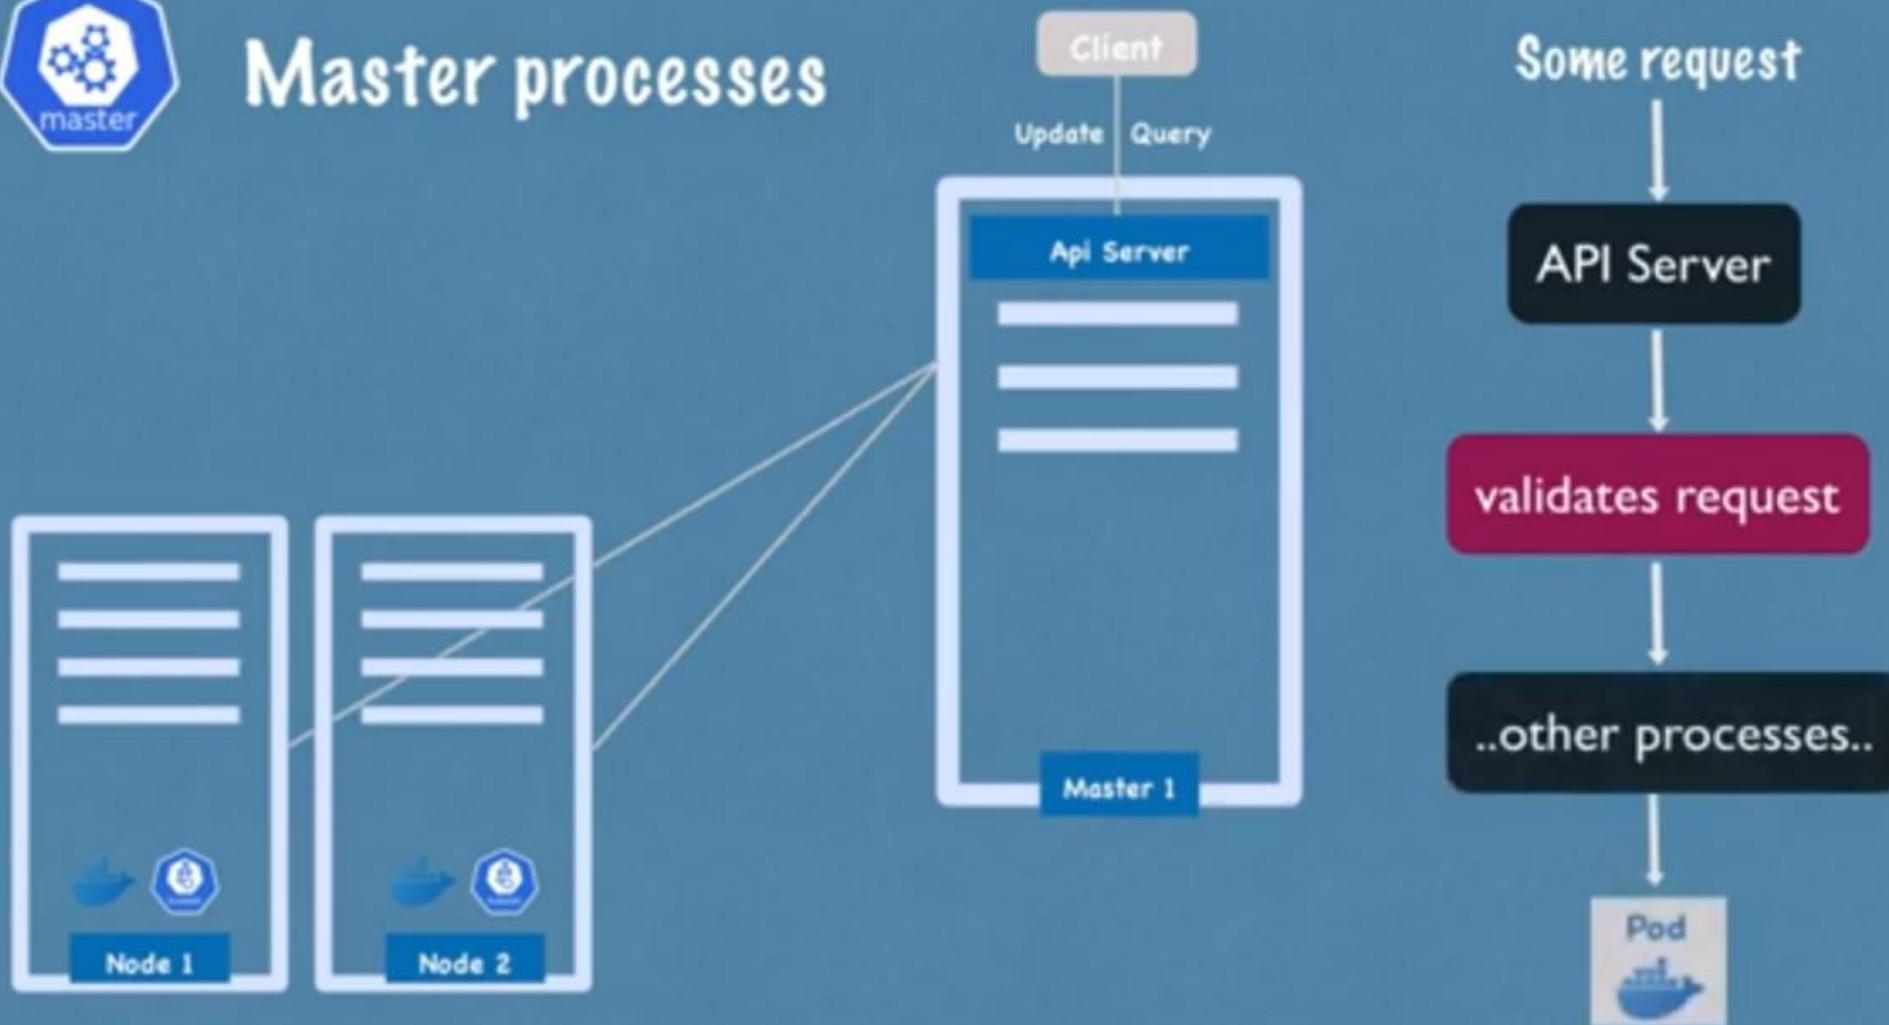
## Master processes



- cluster gateway
- acts as a gatekeeper for authentication!

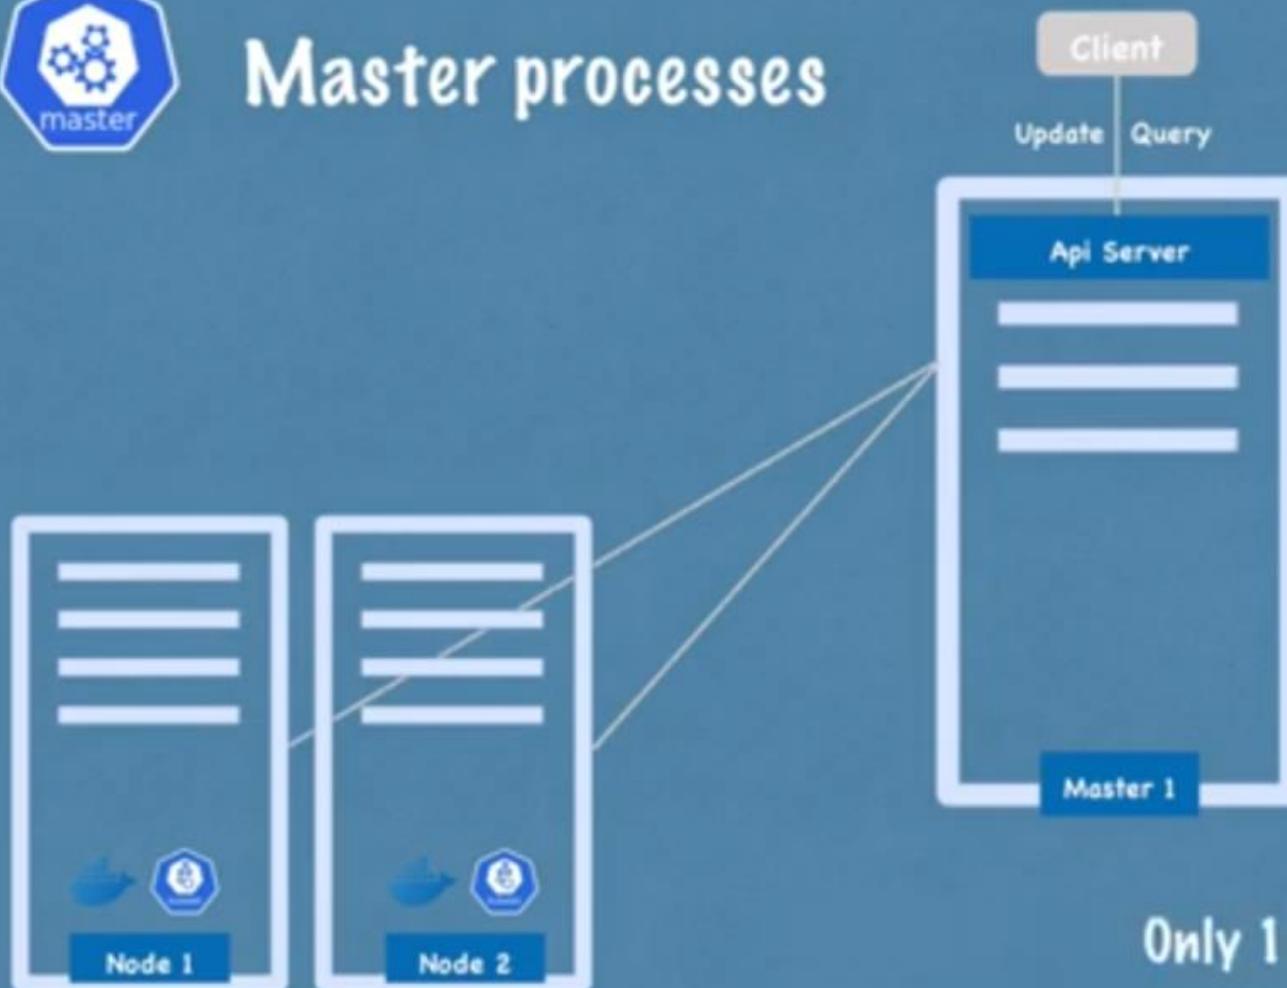


## Master processes





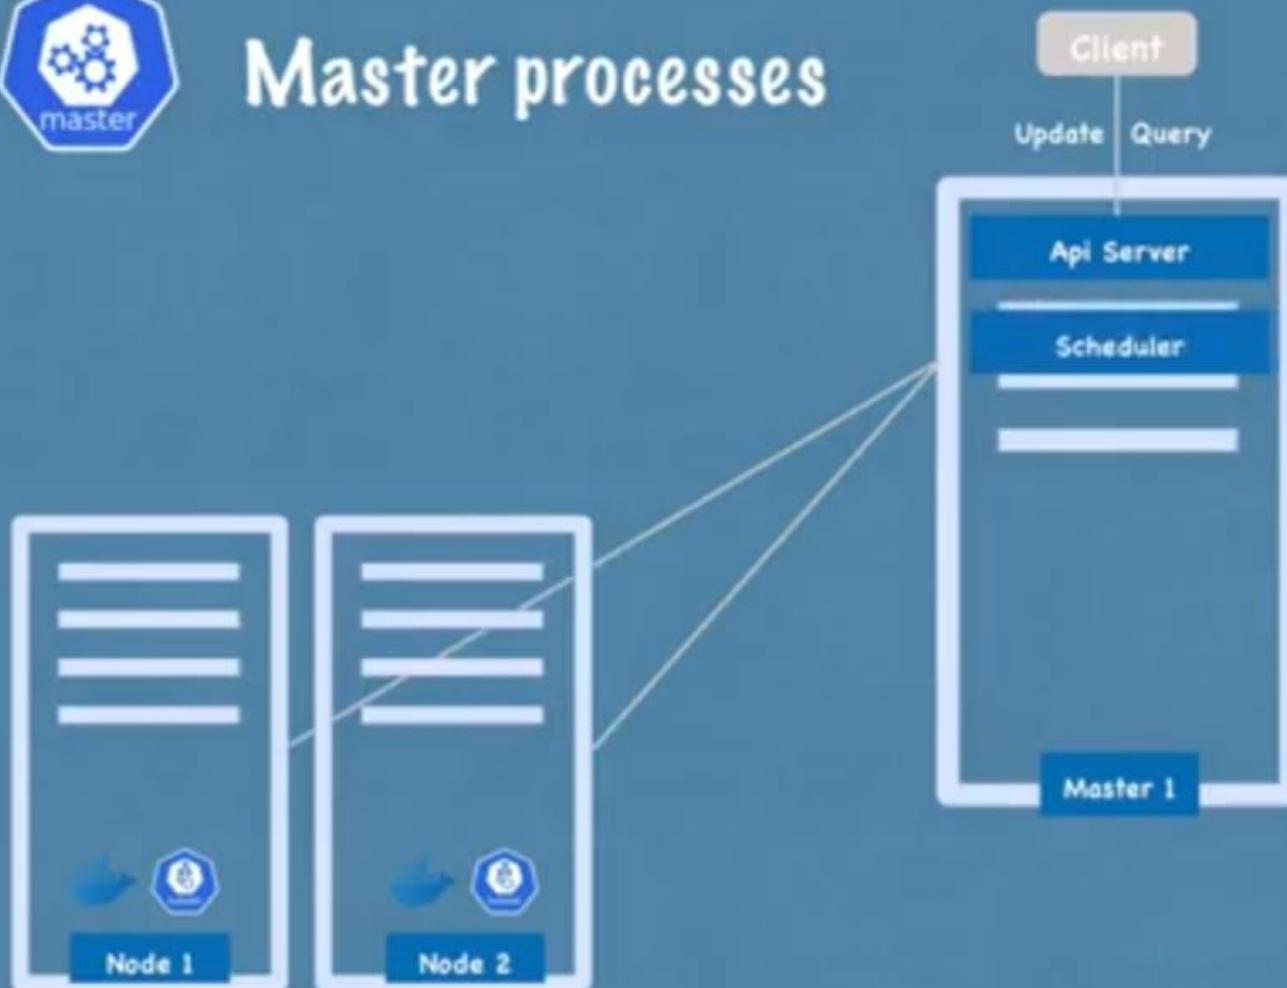
## Master processes



Only 1 entrypoint into the cluster

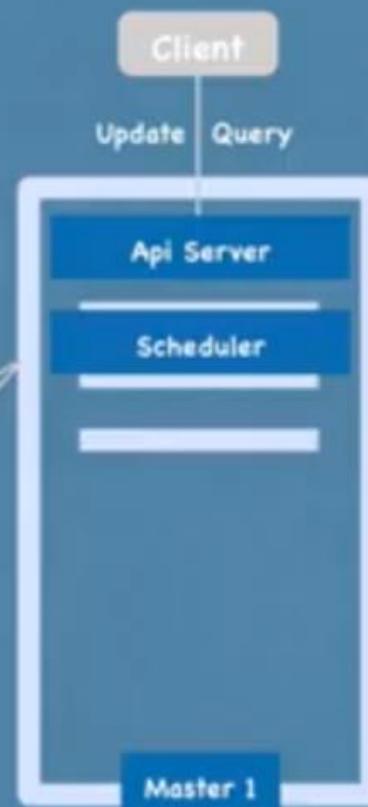
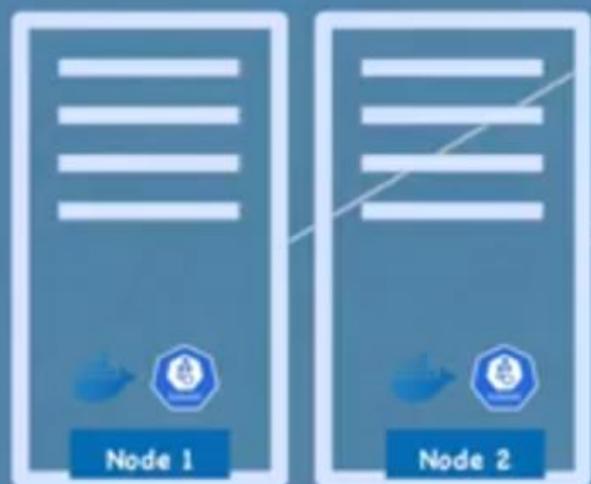


## Master processes





## Master processes



Schedule new Pod

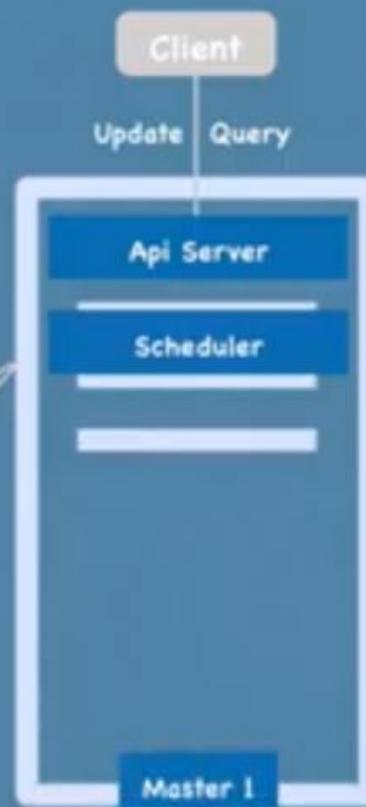
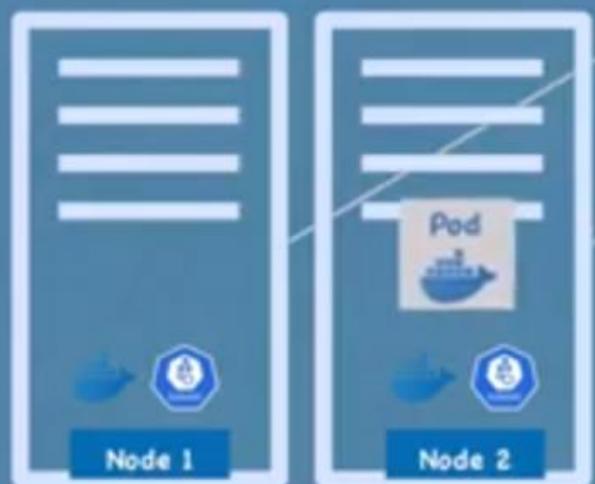
API Server

Scheduler





## Master processes



Schedule new Pod

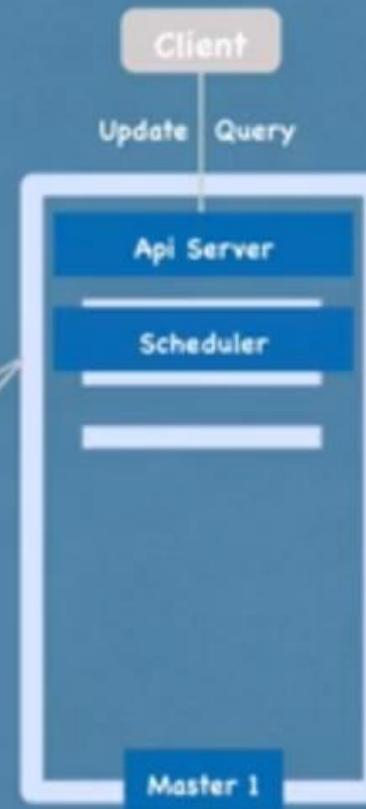
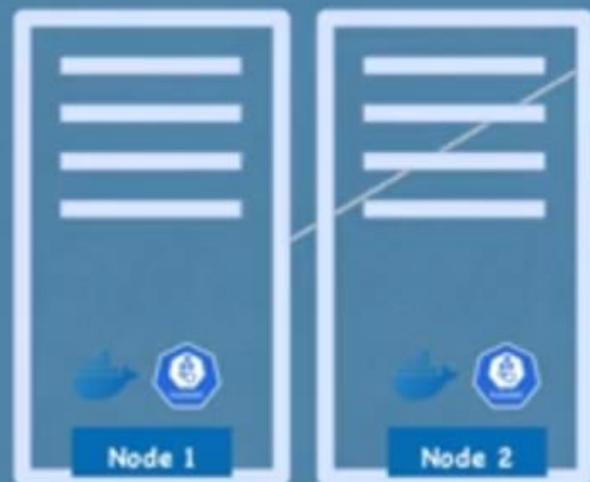
API Server

Scheduler

Master 1



## Master processes



Schedule new Pod

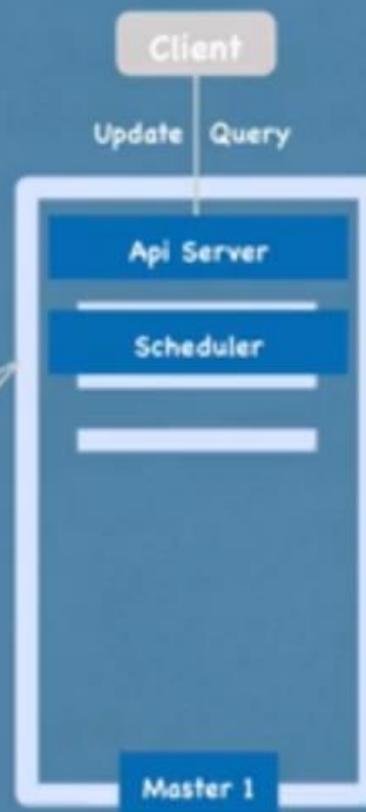
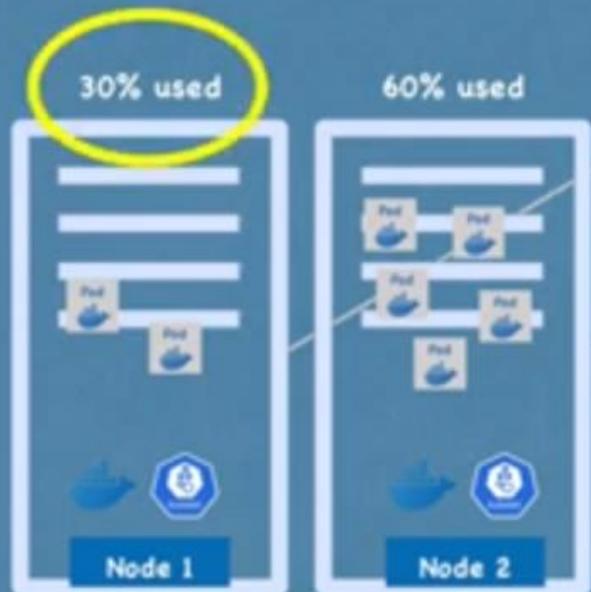
API Server

Scheduler

Where to put the Pod?



## Master processes



Schedule new Pod

API Server

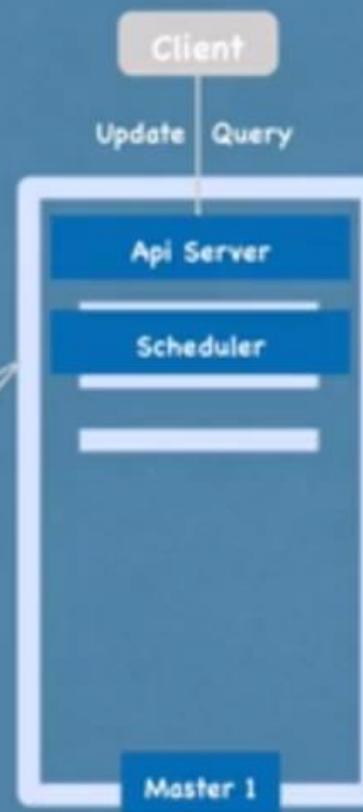
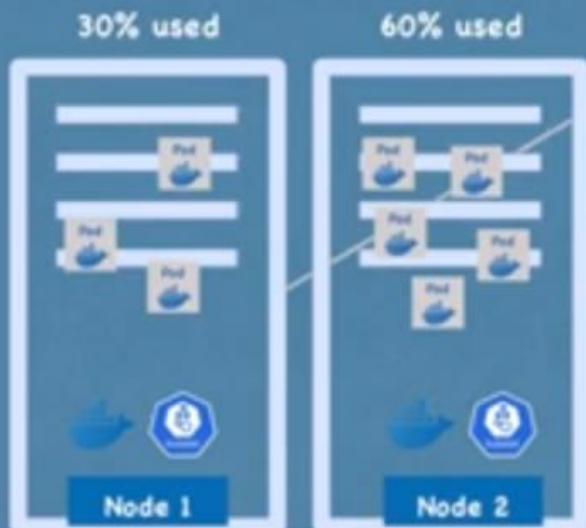
Scheduler

Where to put the Pod?



## Master processes

Scheduler just decides on which Node new Pod should be scheduled



Schedule new Pod

API Server

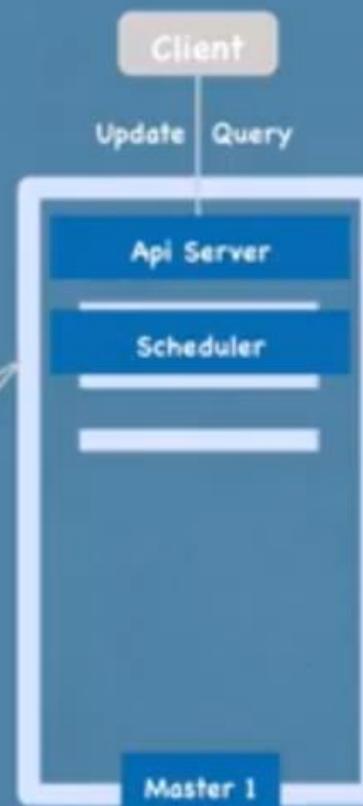
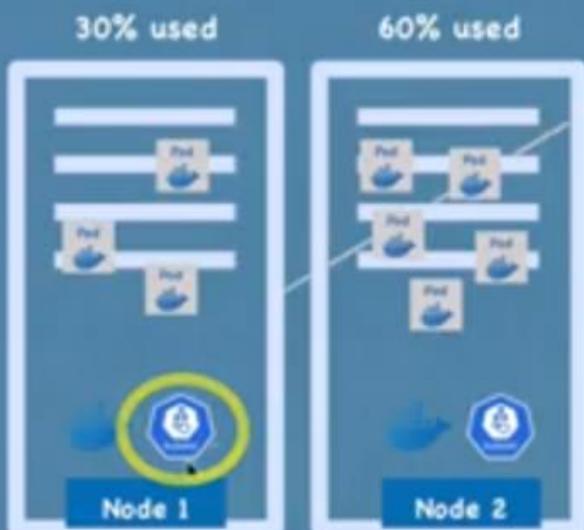
Scheduler

Where to put the Pod?



## Master processes

Scheduler just decides on which Node new Pod should be scheduled



Schedule new Pod

API Server

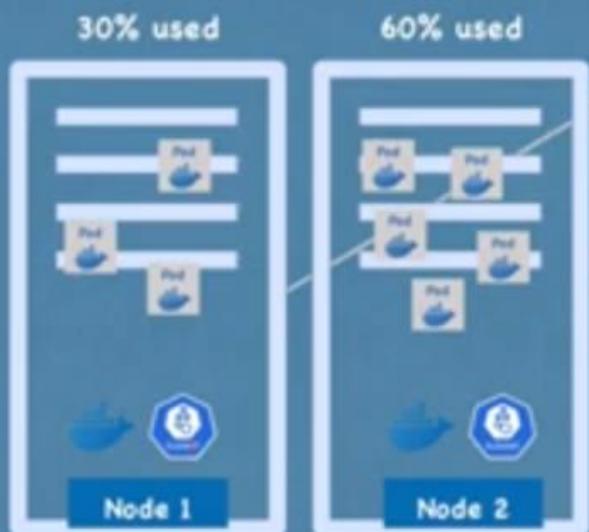
Scheduler

Where to put the Pod?

Kubelet

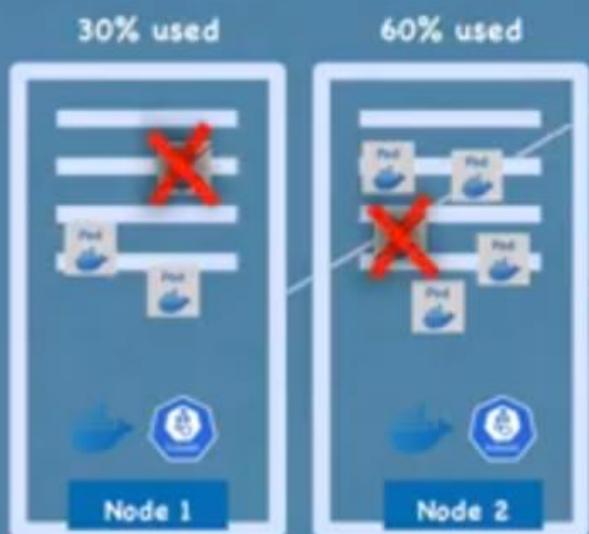


# Master processes



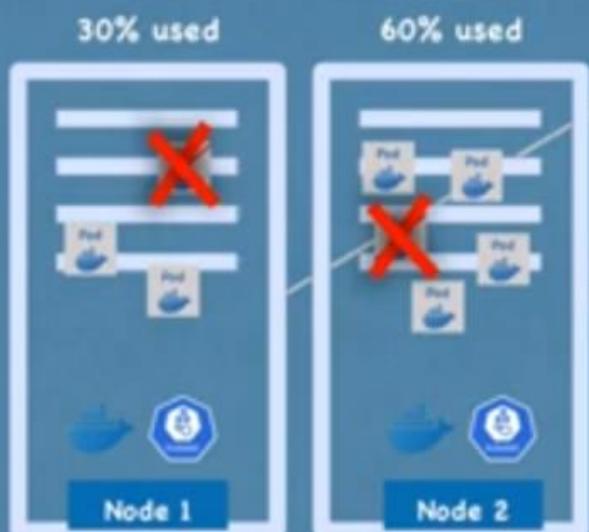


# Master processes





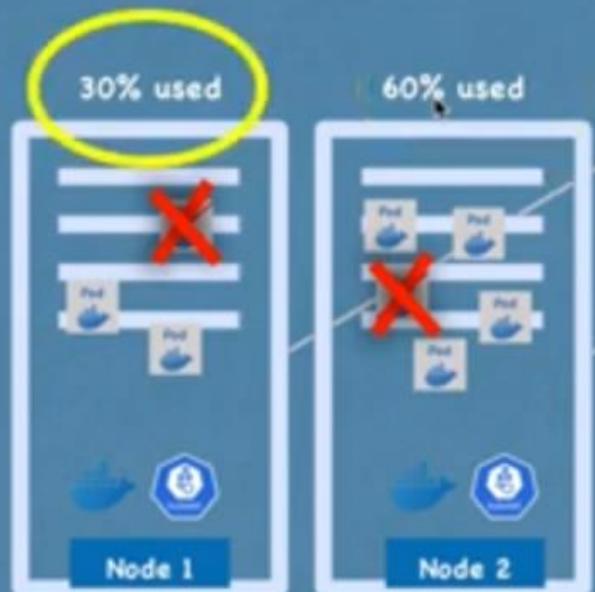
## Master processes



- detects cluster state changes



## Master processes

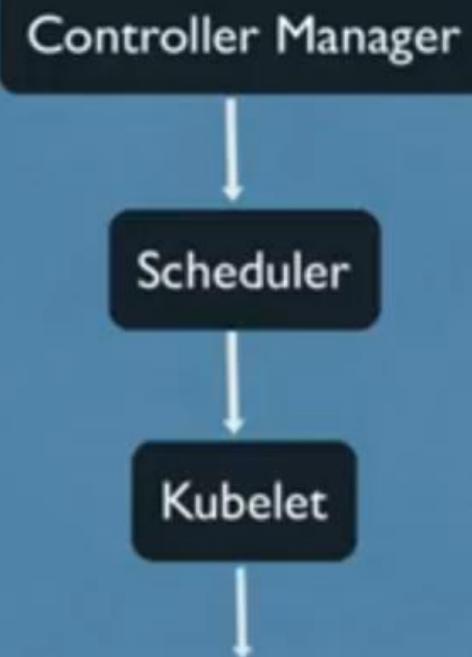
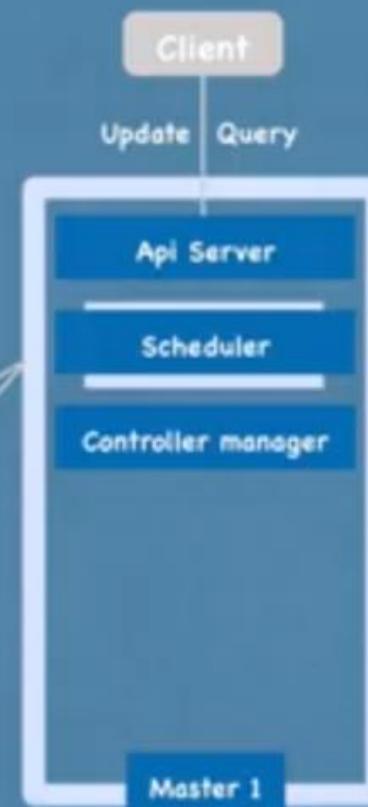
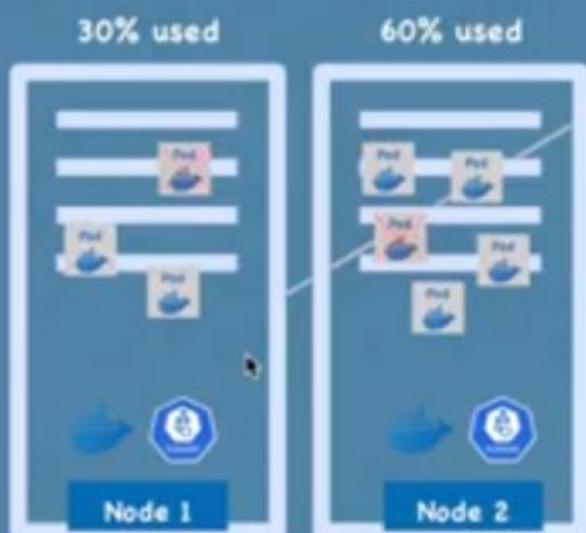


Controller Manager

Scheduler

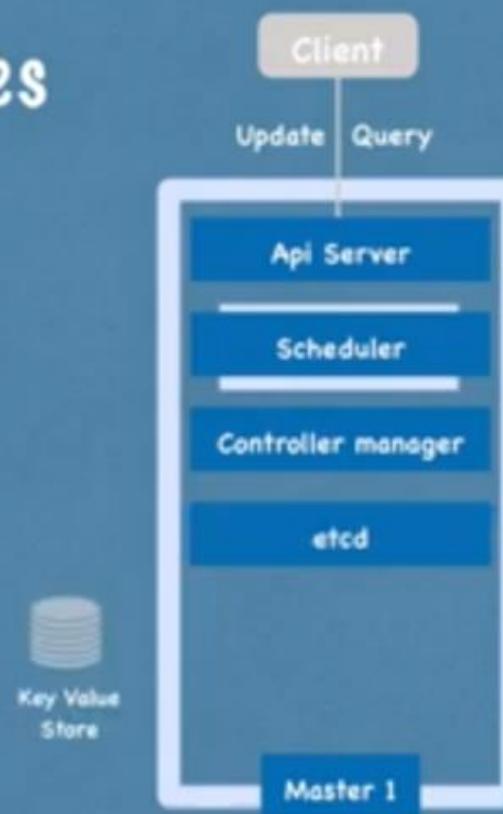
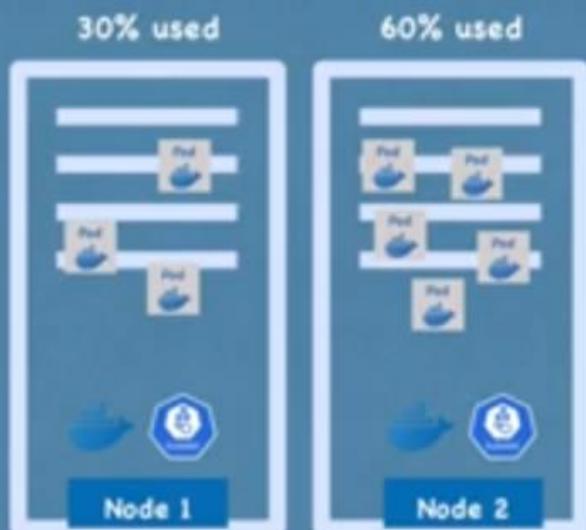


# Master processes



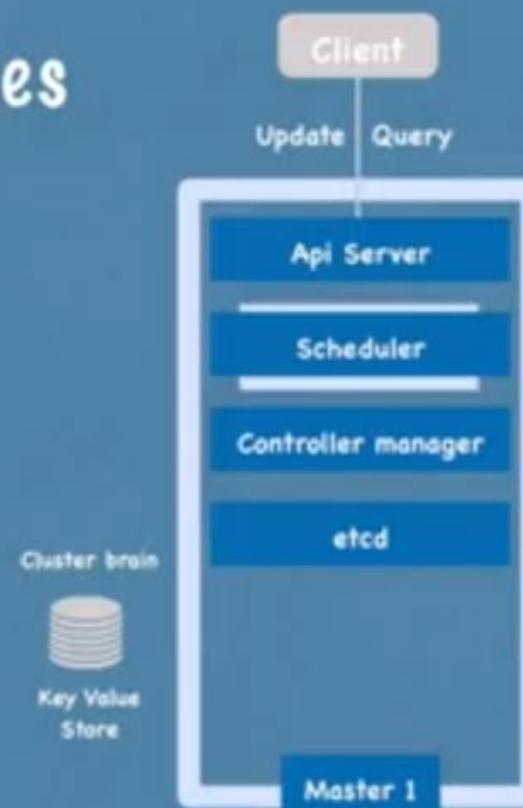
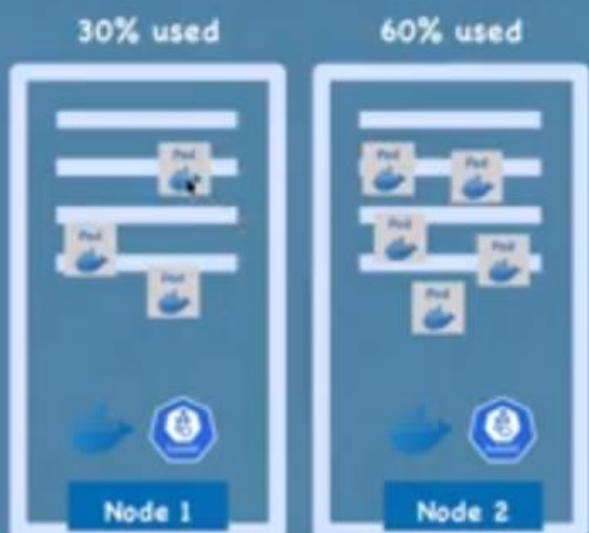


# Master processes





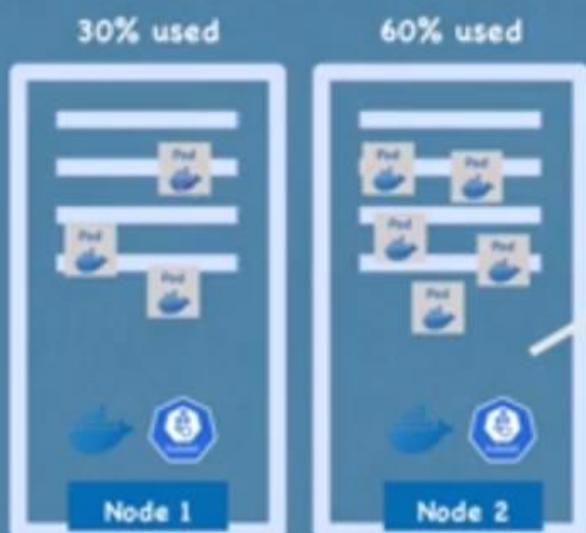
## Master processes



- etcd is the **cluster brain!**



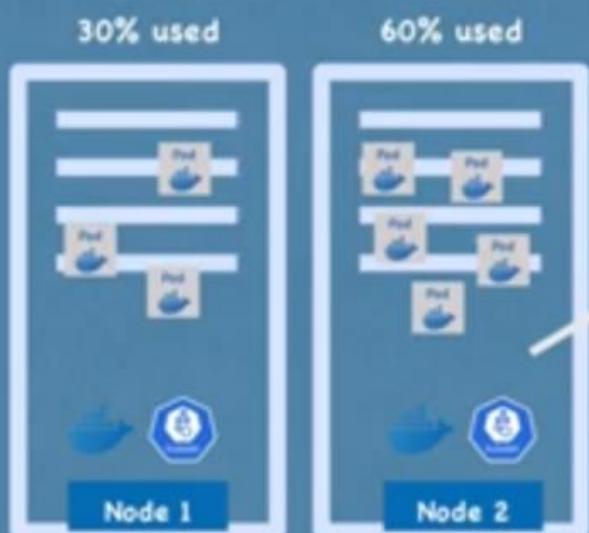
## Master processes



- etcd is the **cluster brain!**
- Cluster changes get stored in the key value store



# Master processes

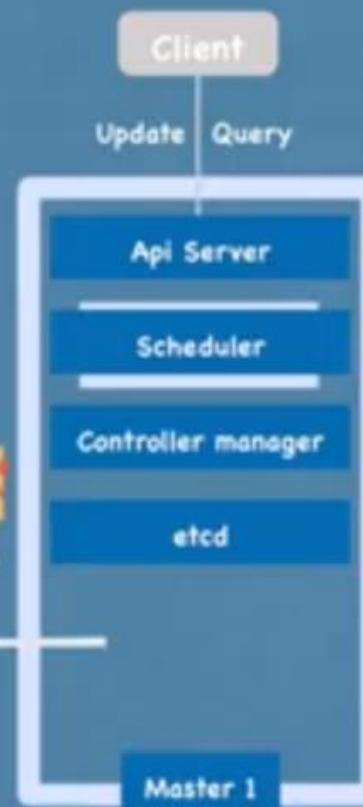
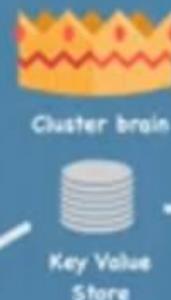
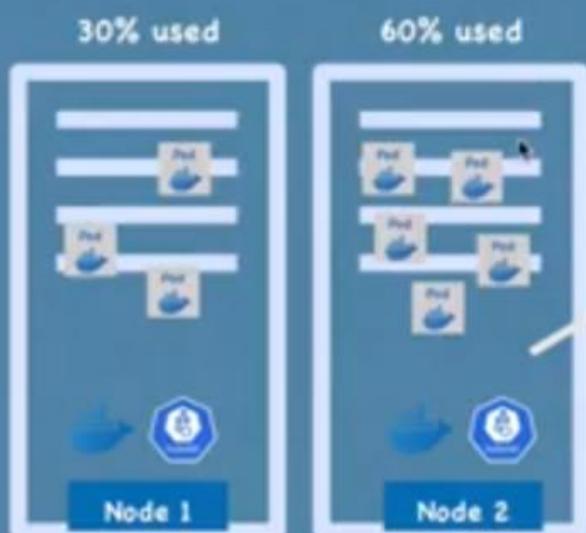


data





# Master processes

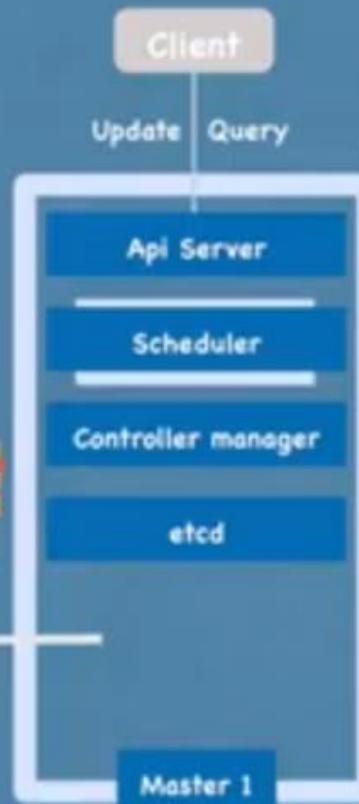
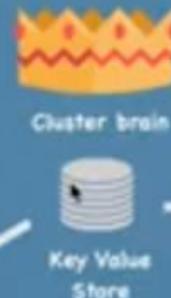
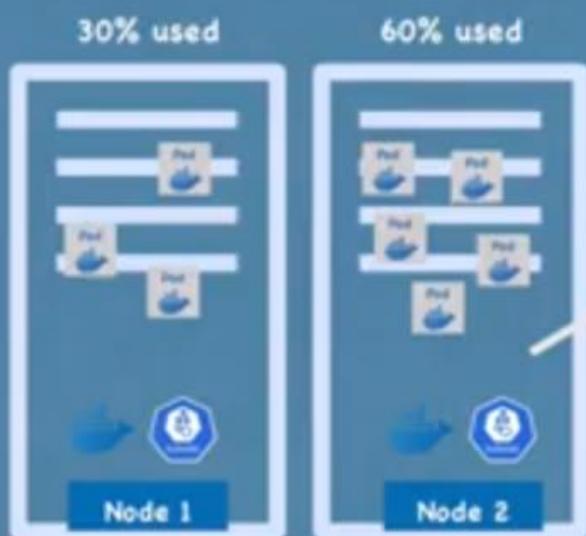


- > Is the cluster healthy?
- > What resources are available?
- > Did the cluster state change?





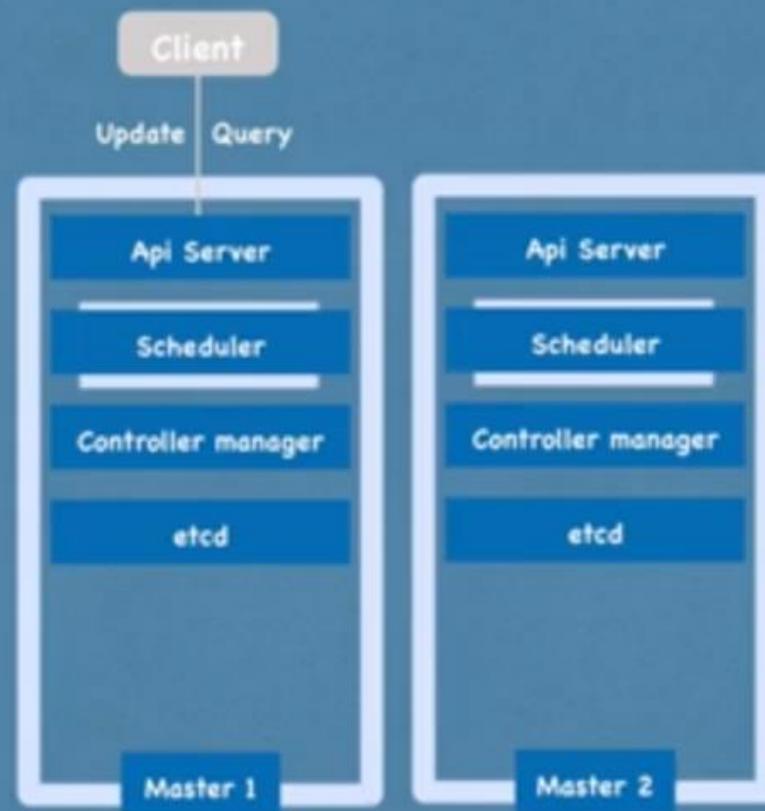
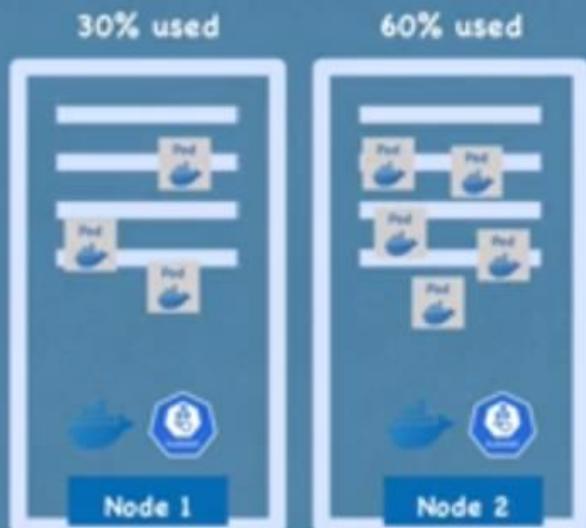
## Master processes



Application data is **NOT** stored in etcd!

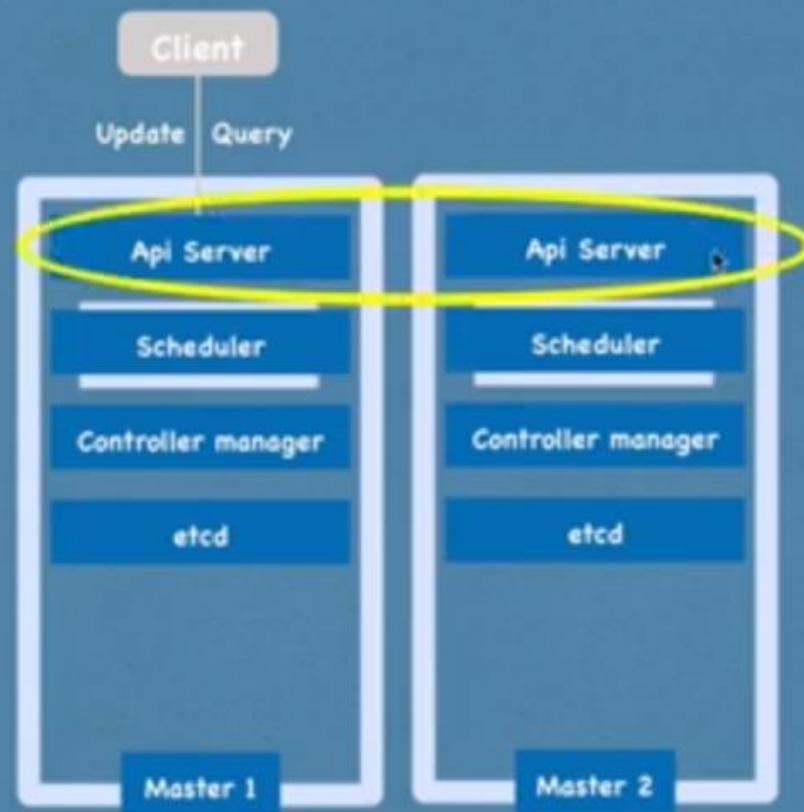
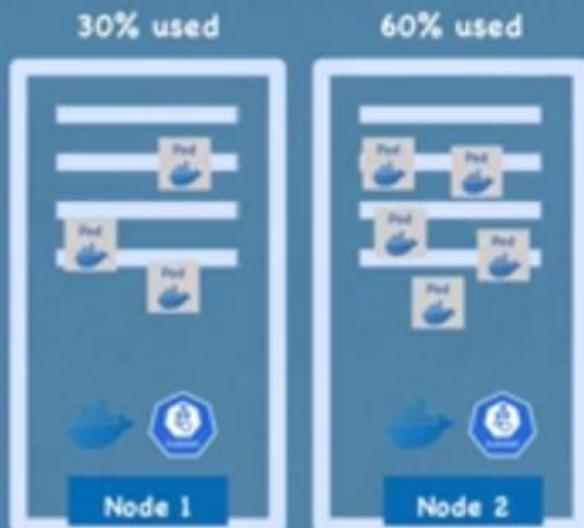


# Master processes





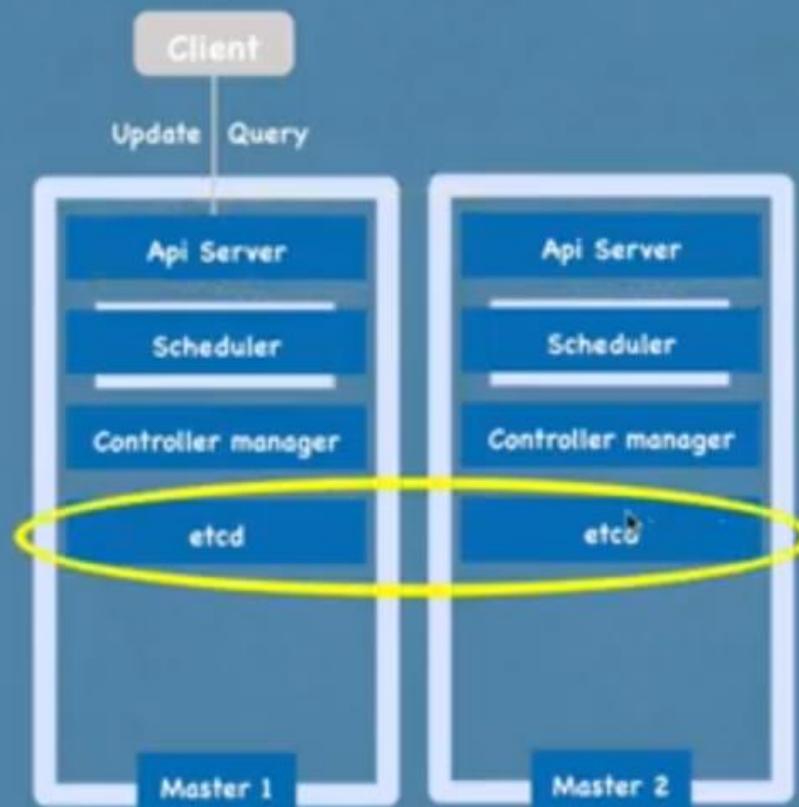
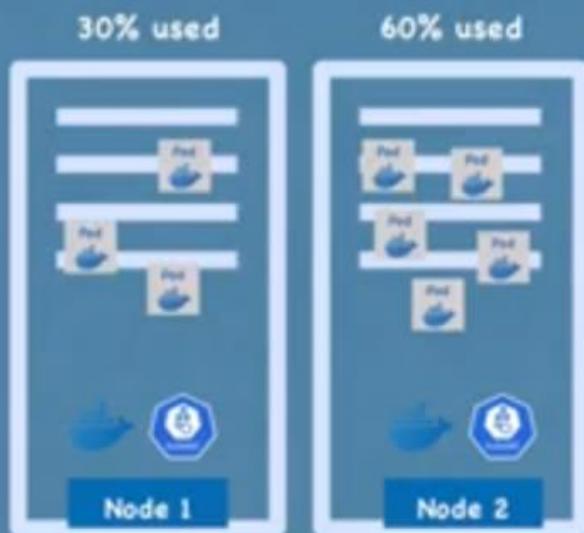
## Master processes



API server is load balanced.



## Master processes



Distributed storage across all master nodes

# Example cluster setup

## Example Cluster Set-Up

2 Master Nodes

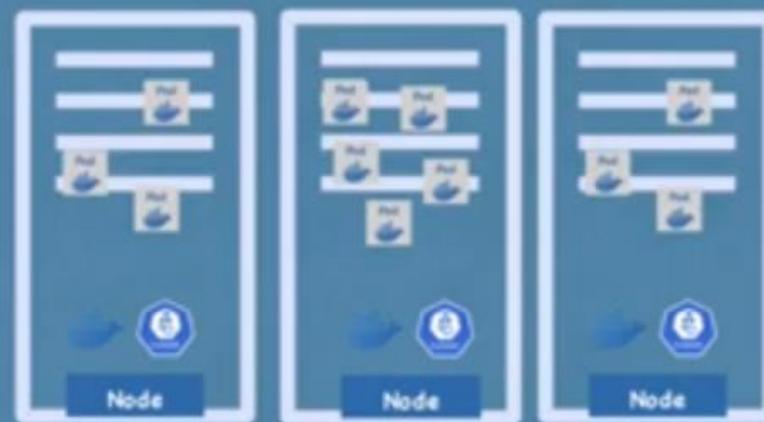
3 Worker Nodes



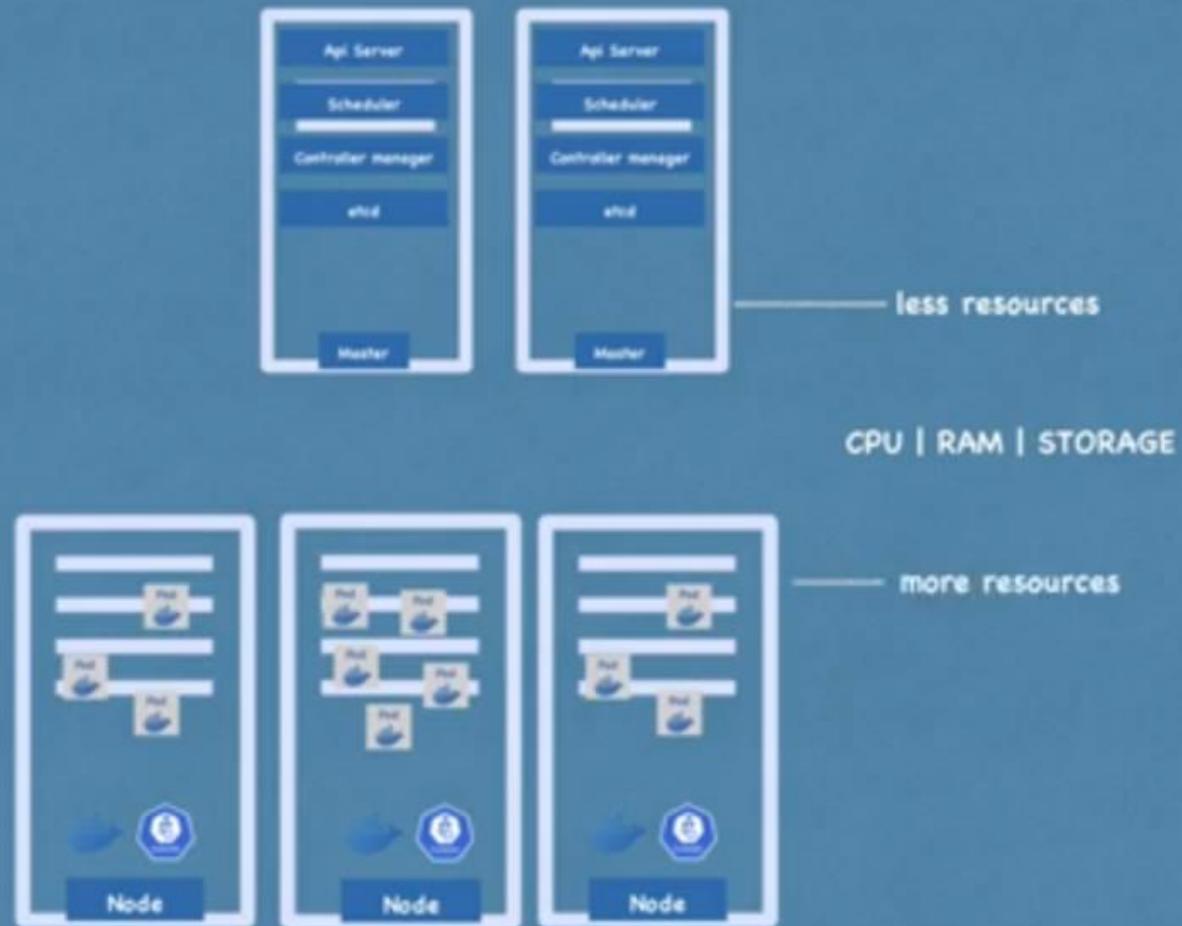
## Example Cluster Set-Up



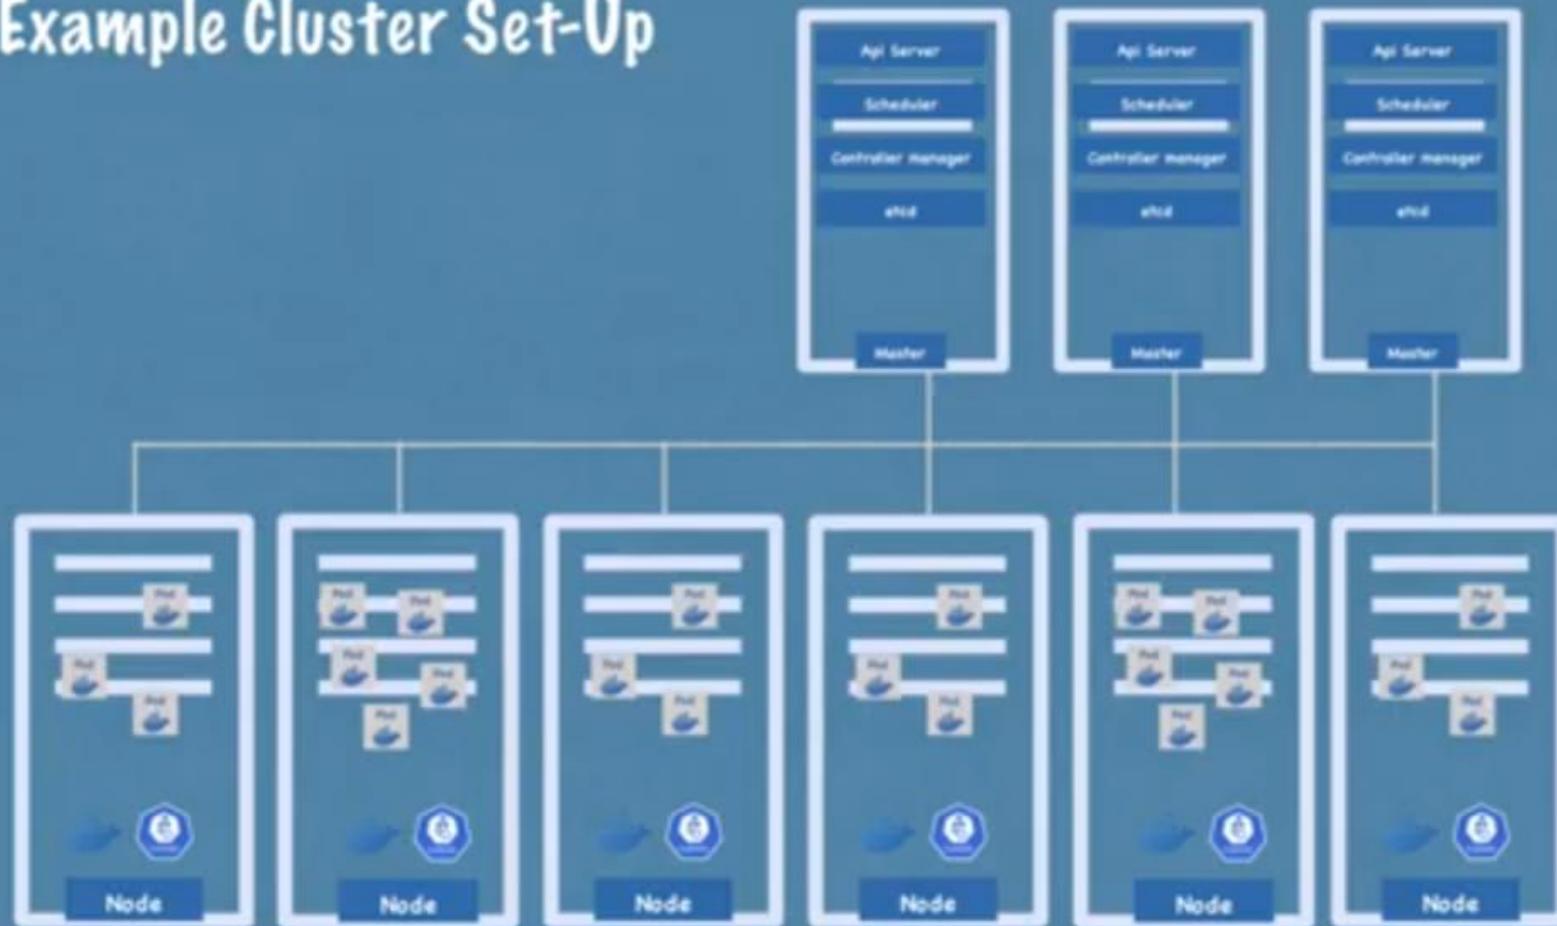
CPU | RAM | STORAGE



## Example Cluster Set-Up

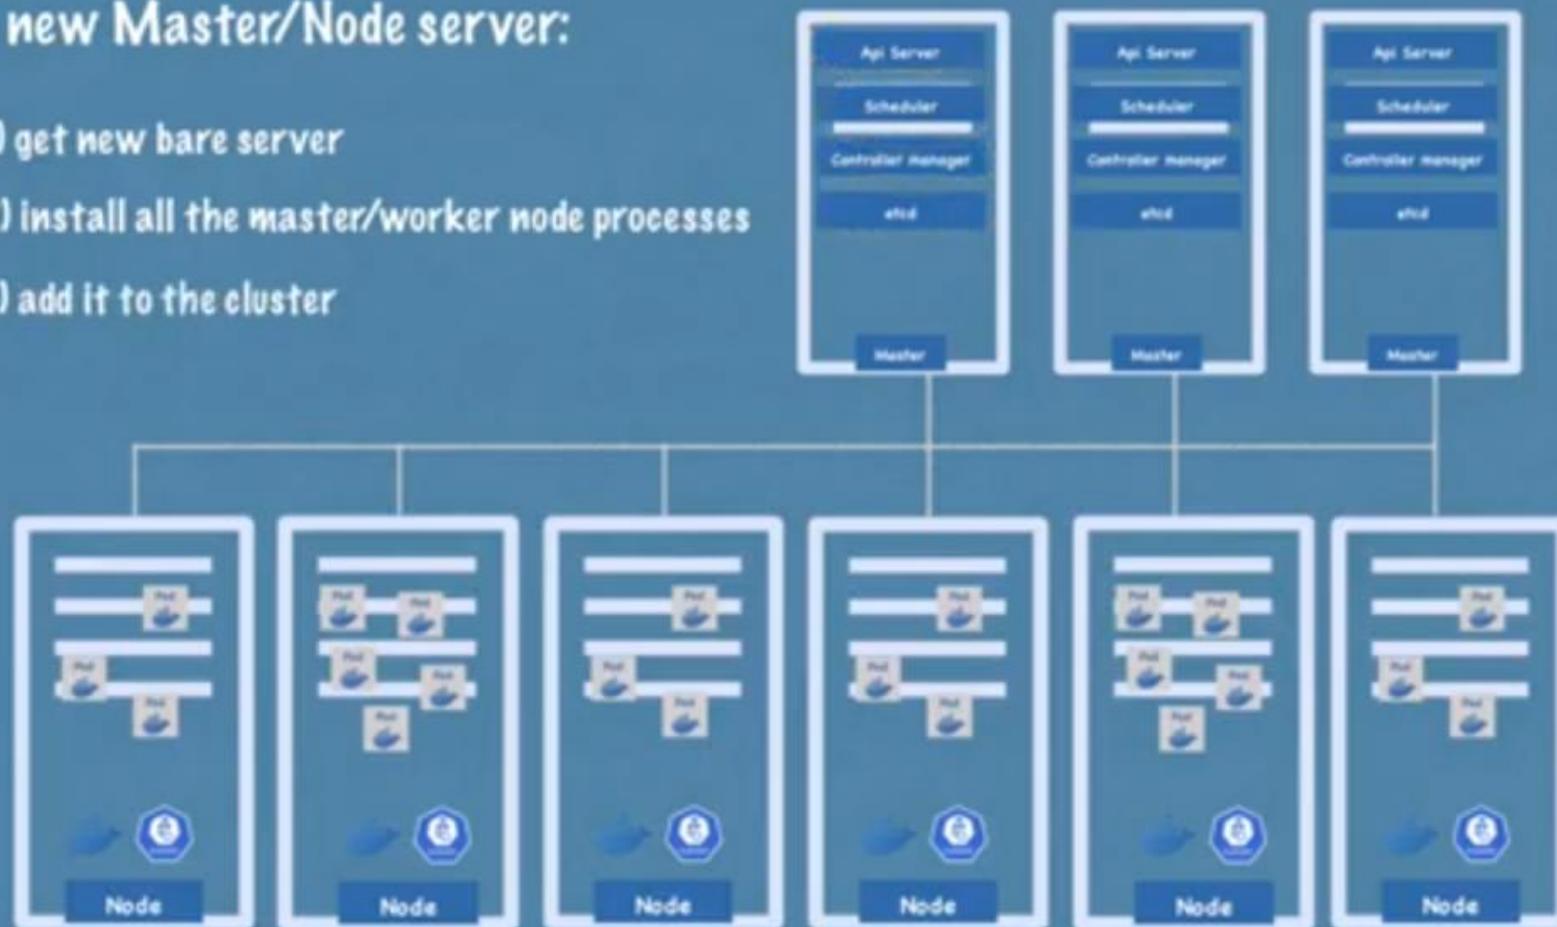


## Example Cluster Set-Up



## Add new Master/Node server:

- 1) get new bare server
- 2) install all the master/worker node processes
- 3) add it to the cluster





Kubernetes is an **orchestration tool**



# high availability





scalability

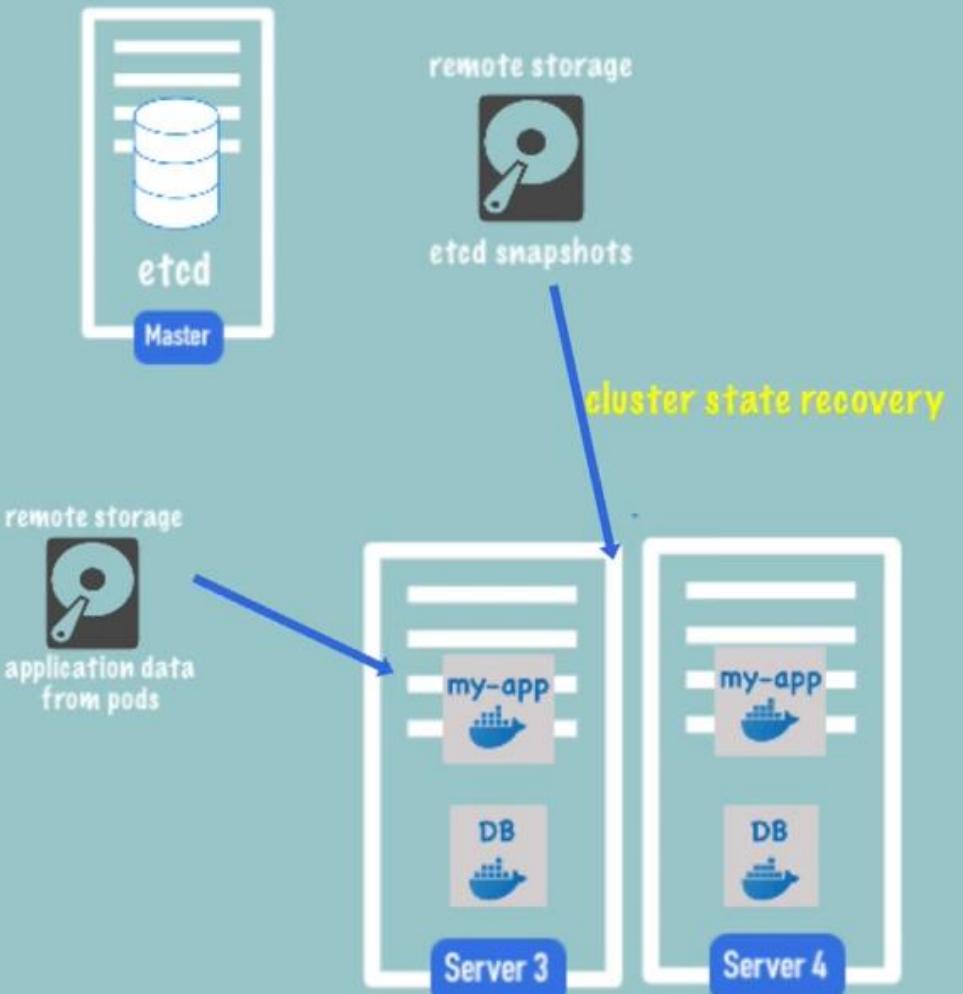
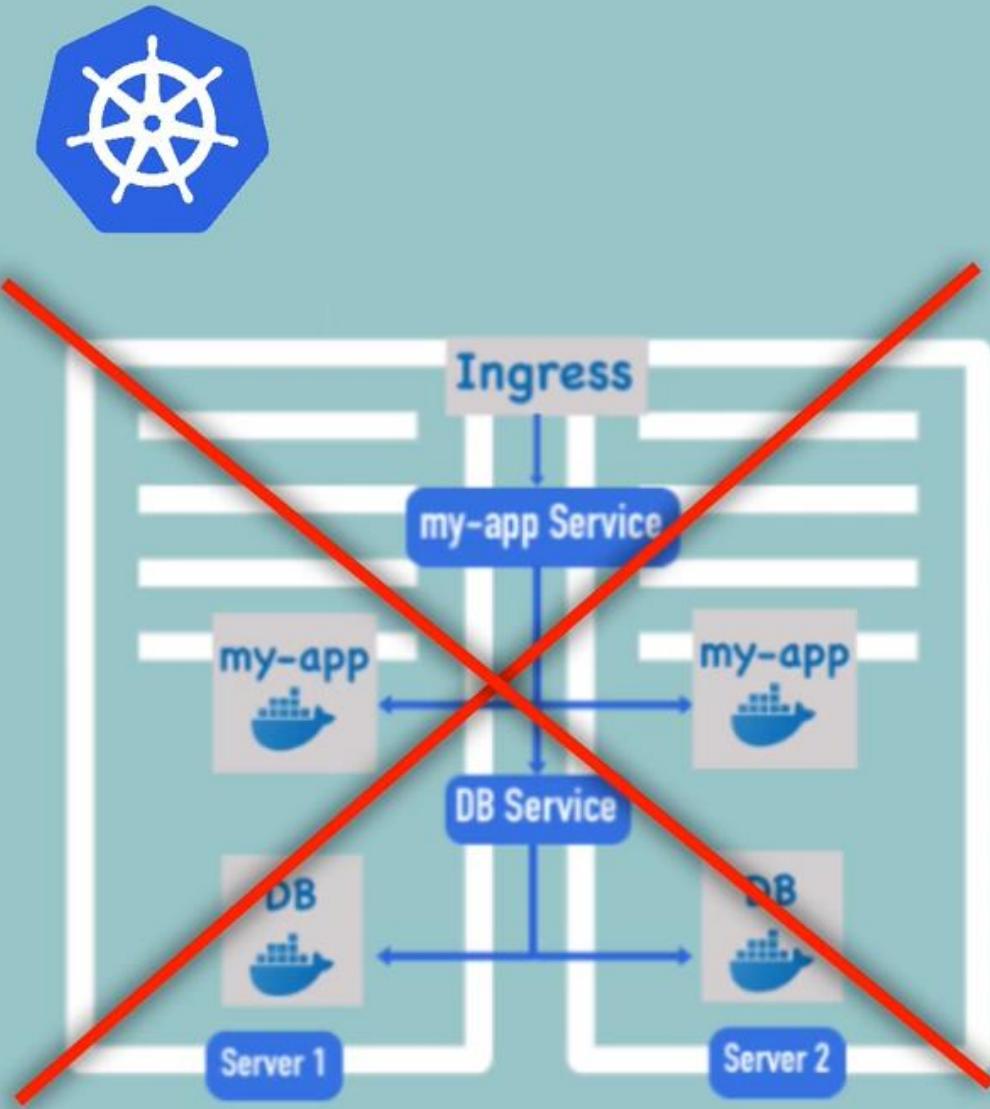
high availability

disaster recovery





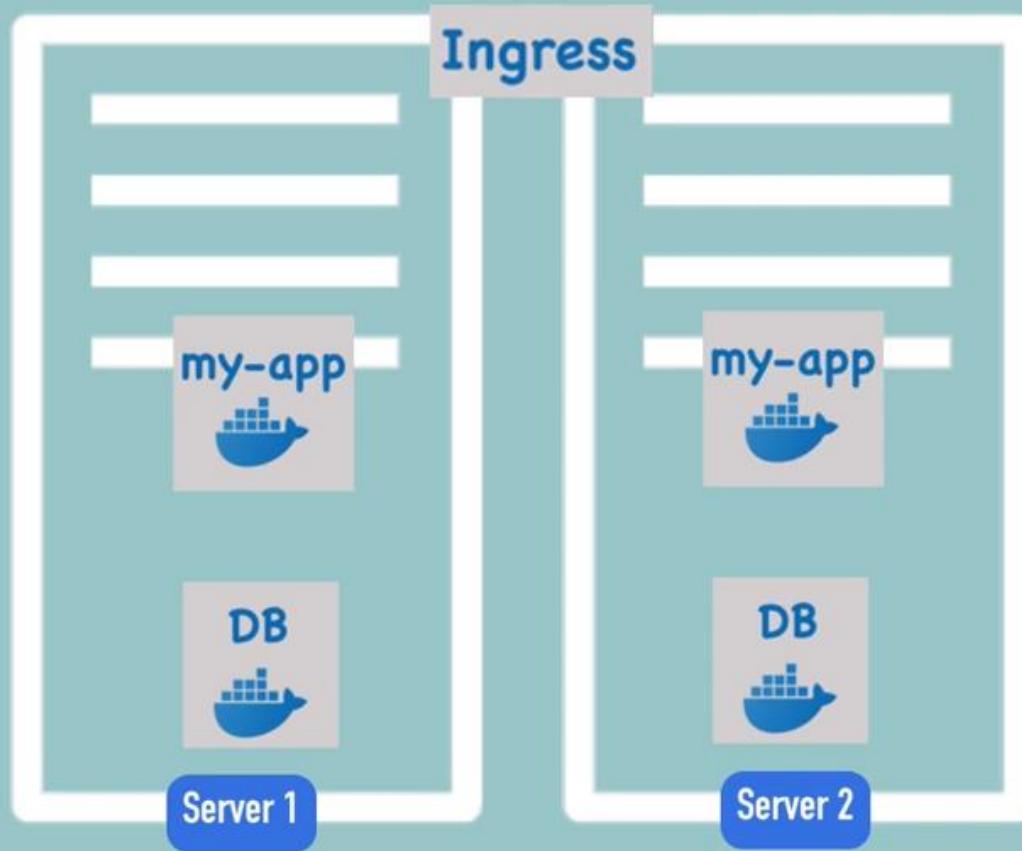
# How K8s actually makes - high availability - scalability - disaster recovery possible?



High availability  
and  
scalability



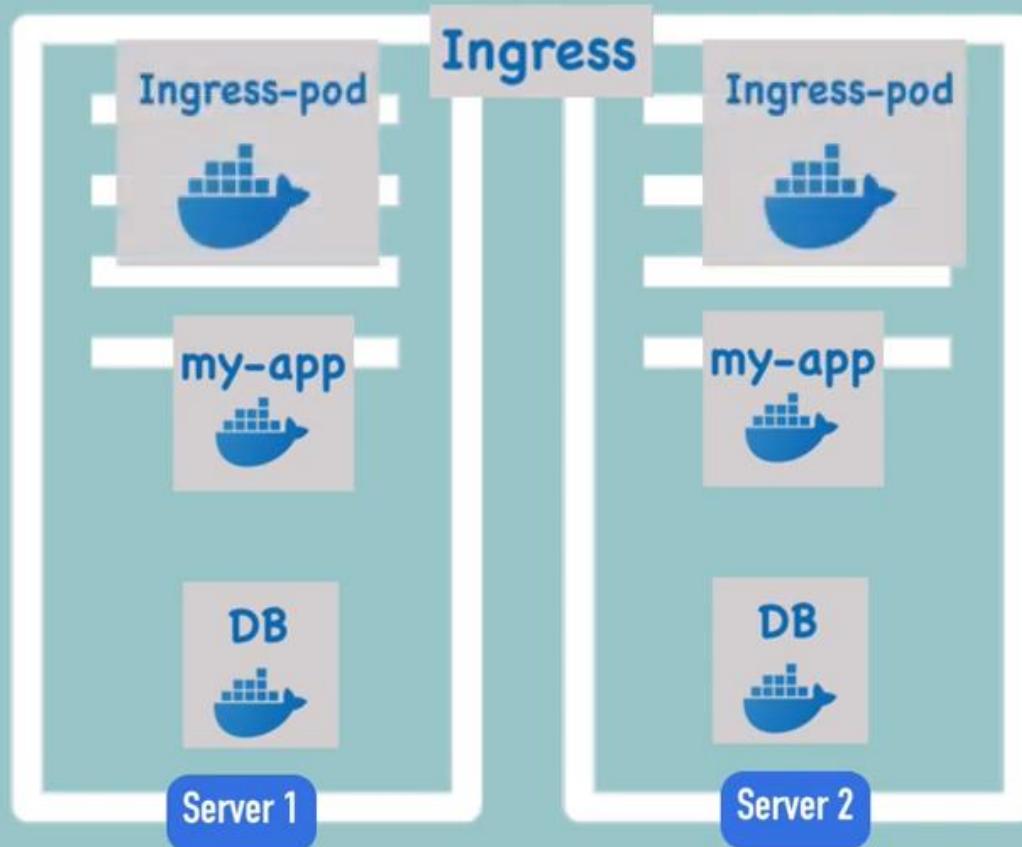
## High availability and scalability



Ingress handles every incoming request



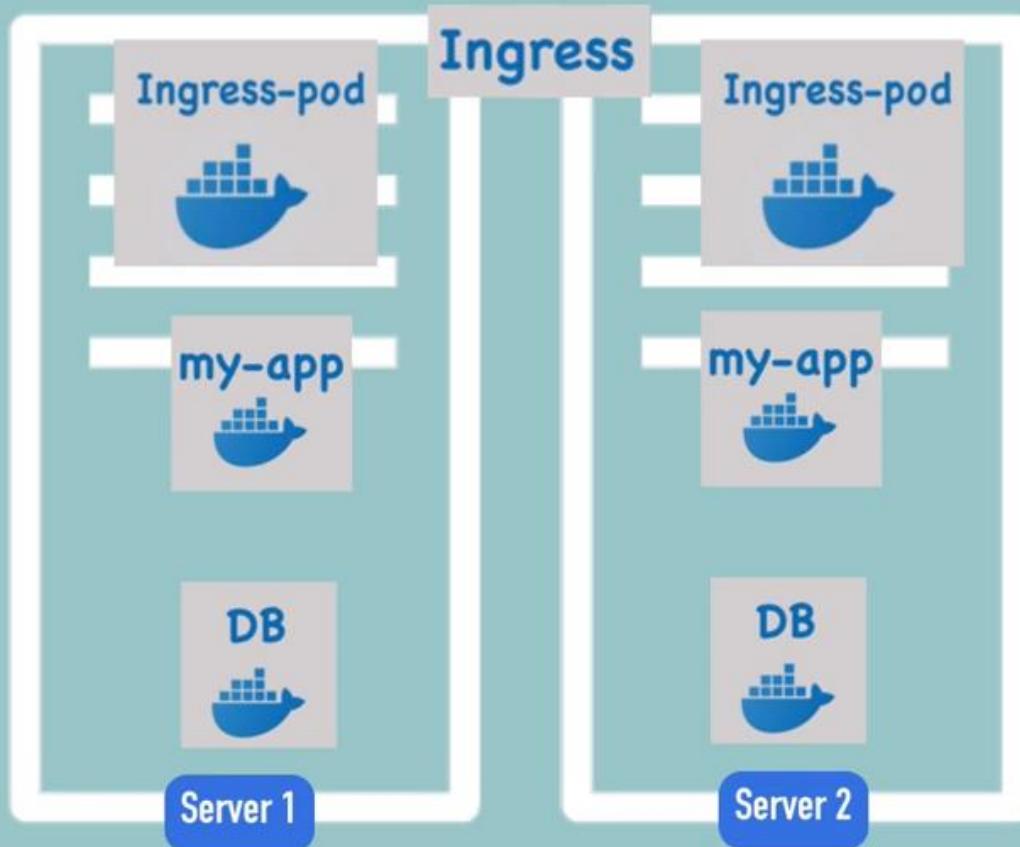
## High availability and scalability



Ingress handles every incoming request

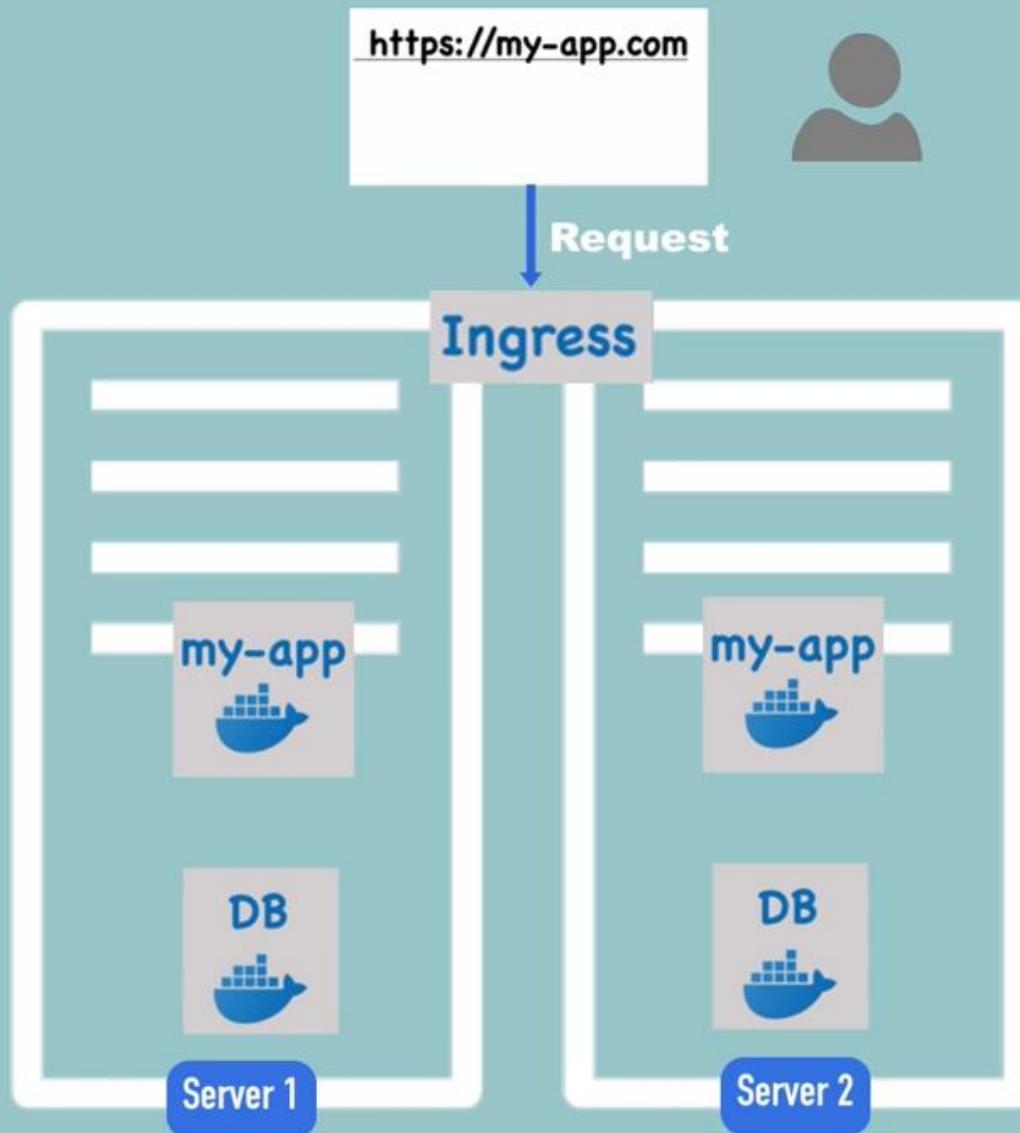


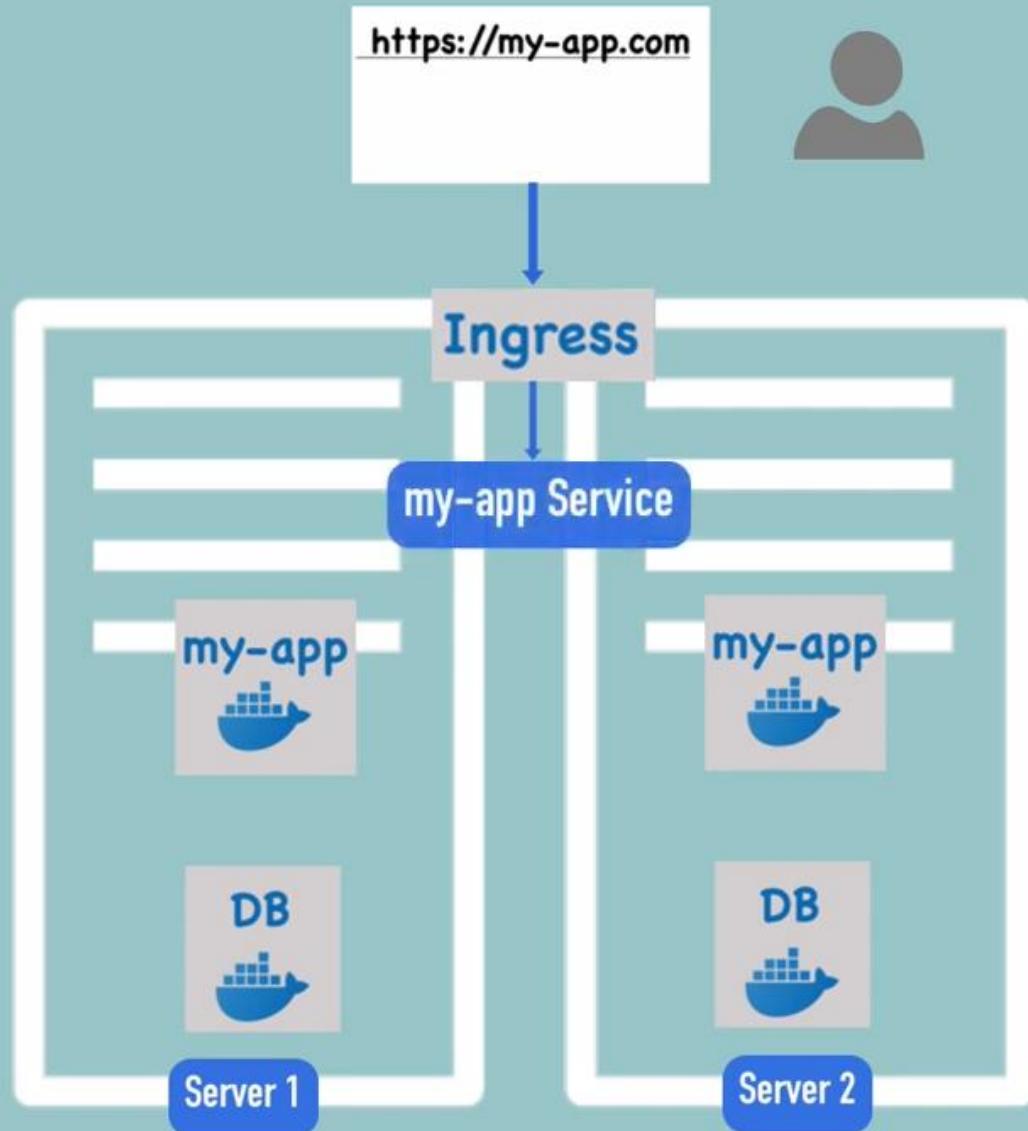
## High availability and scalability

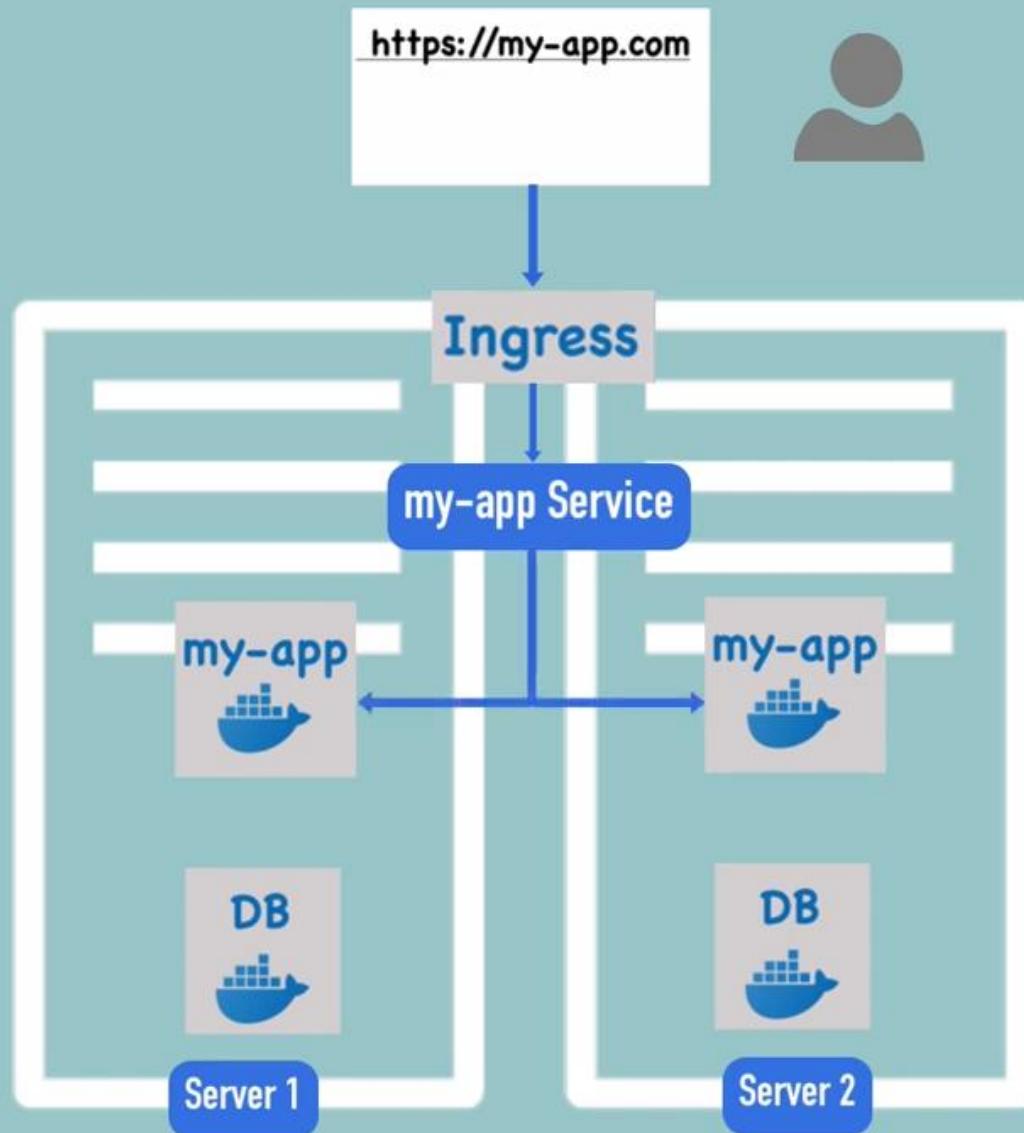


Ingress handles every incoming request

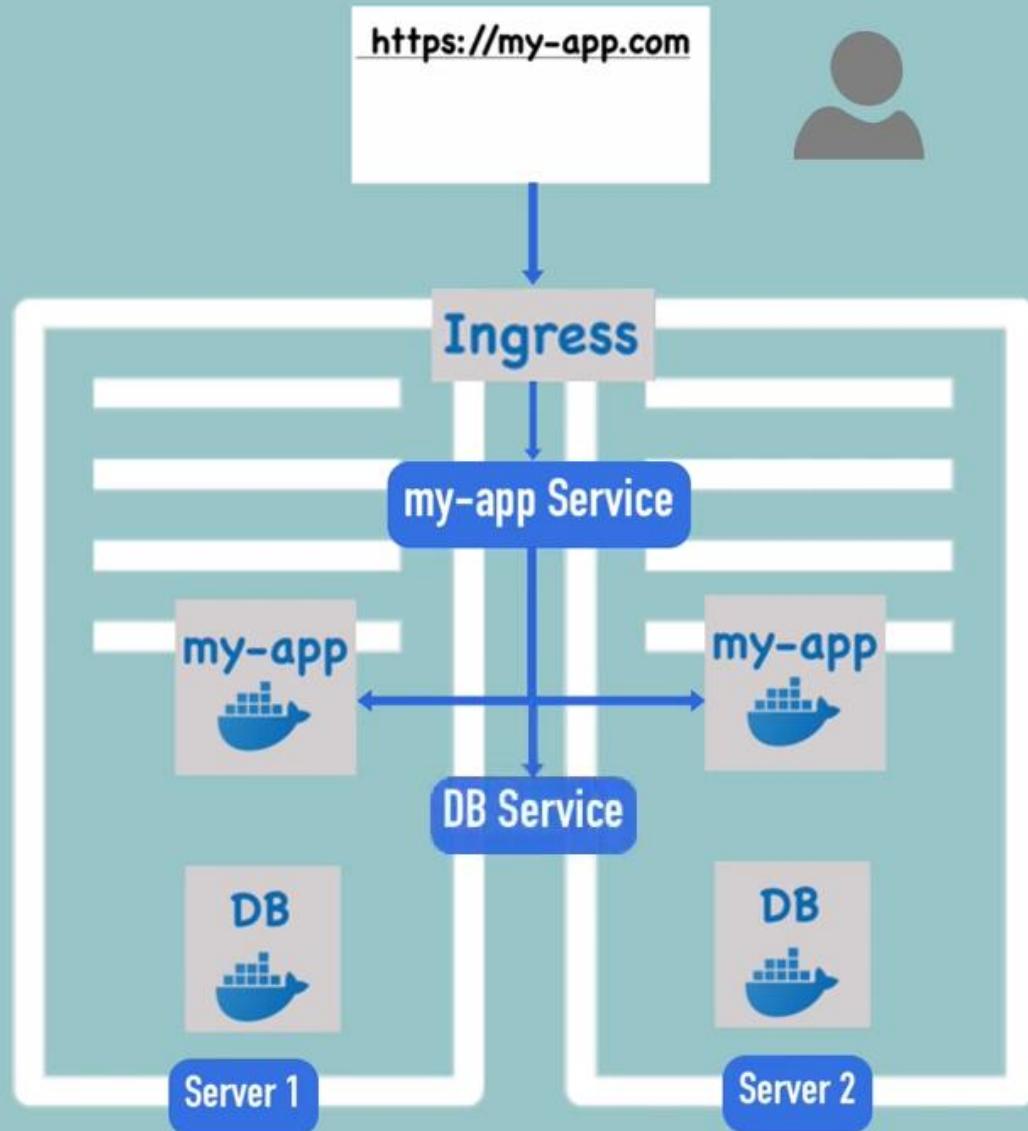
Ingress is also replicated.

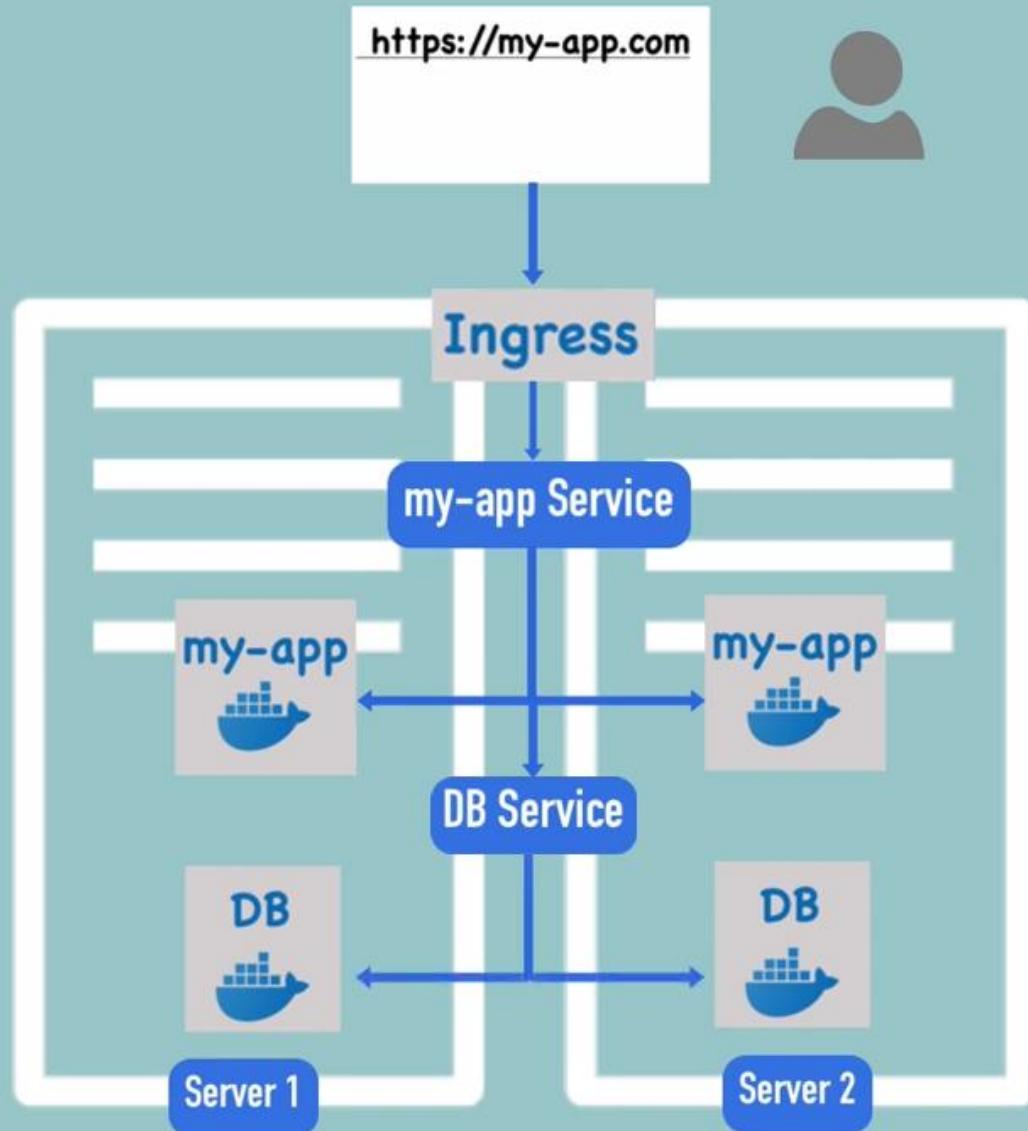


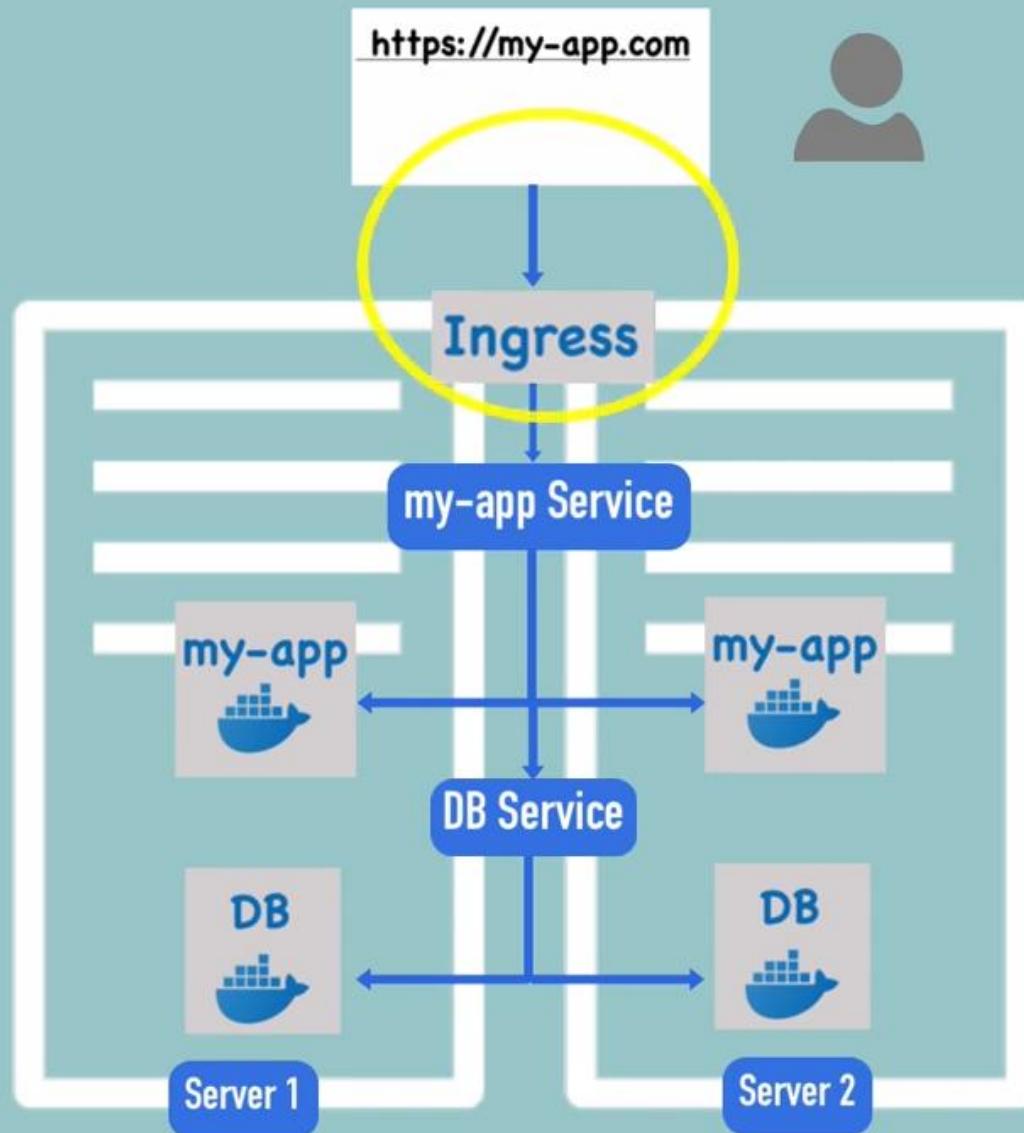


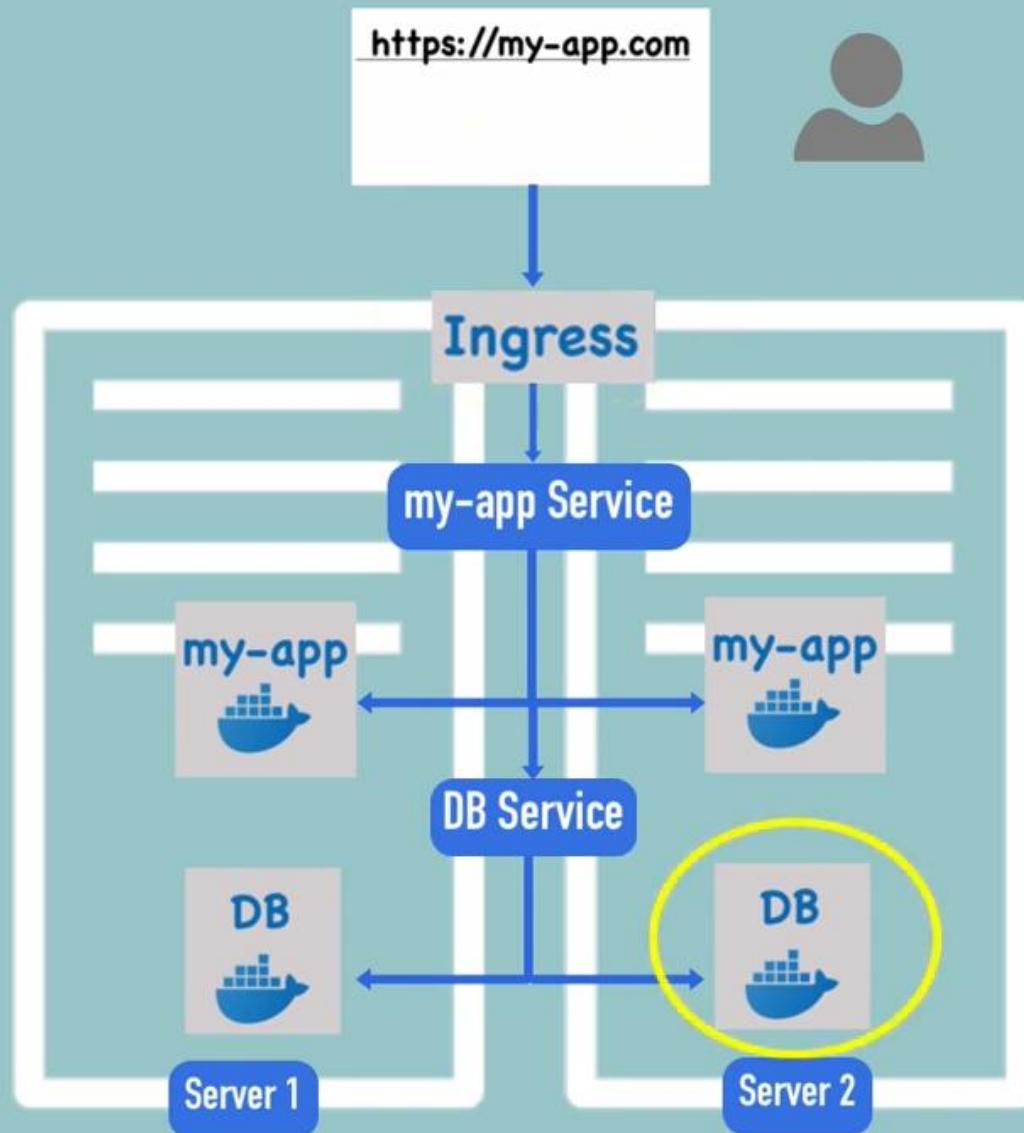


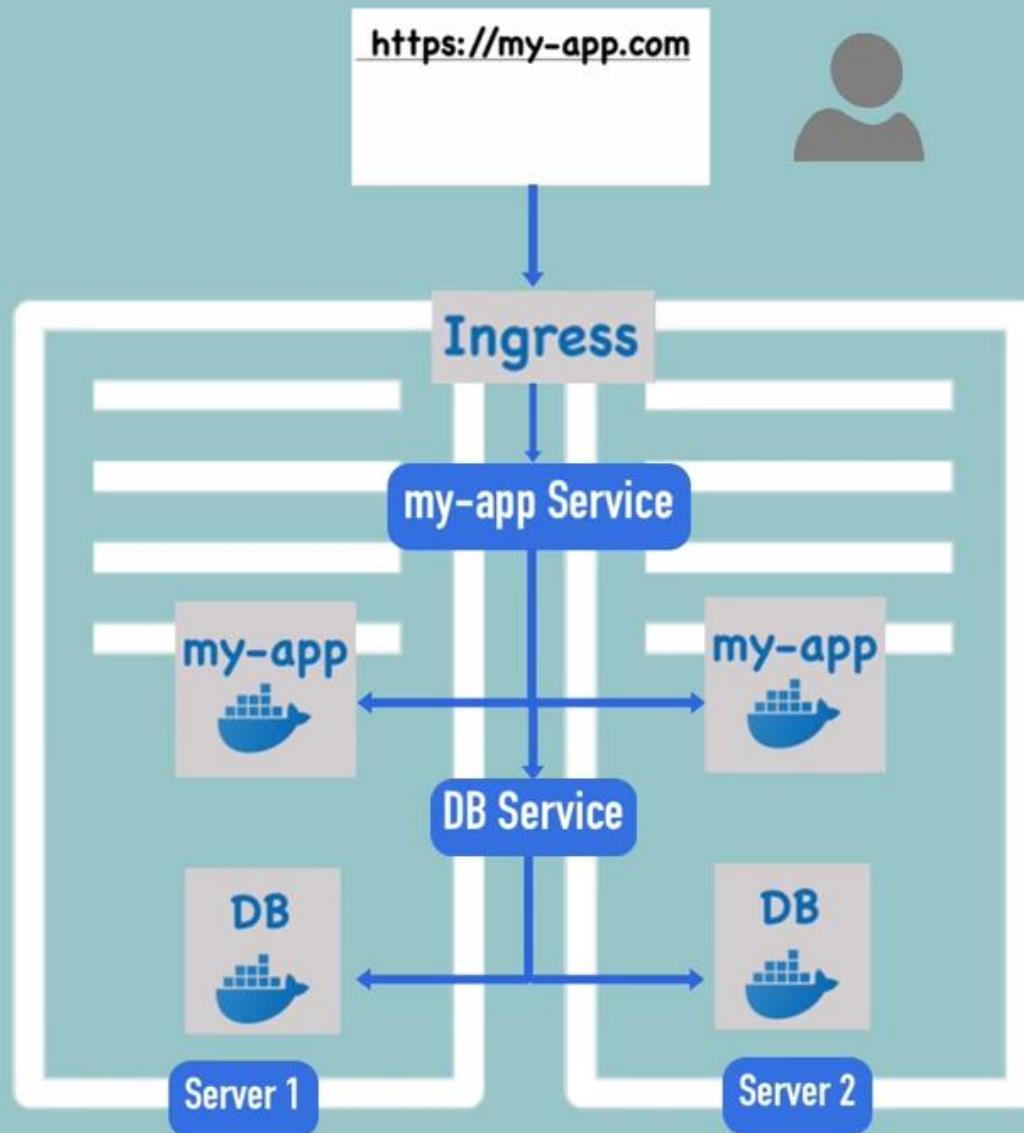
Service is a load balancer



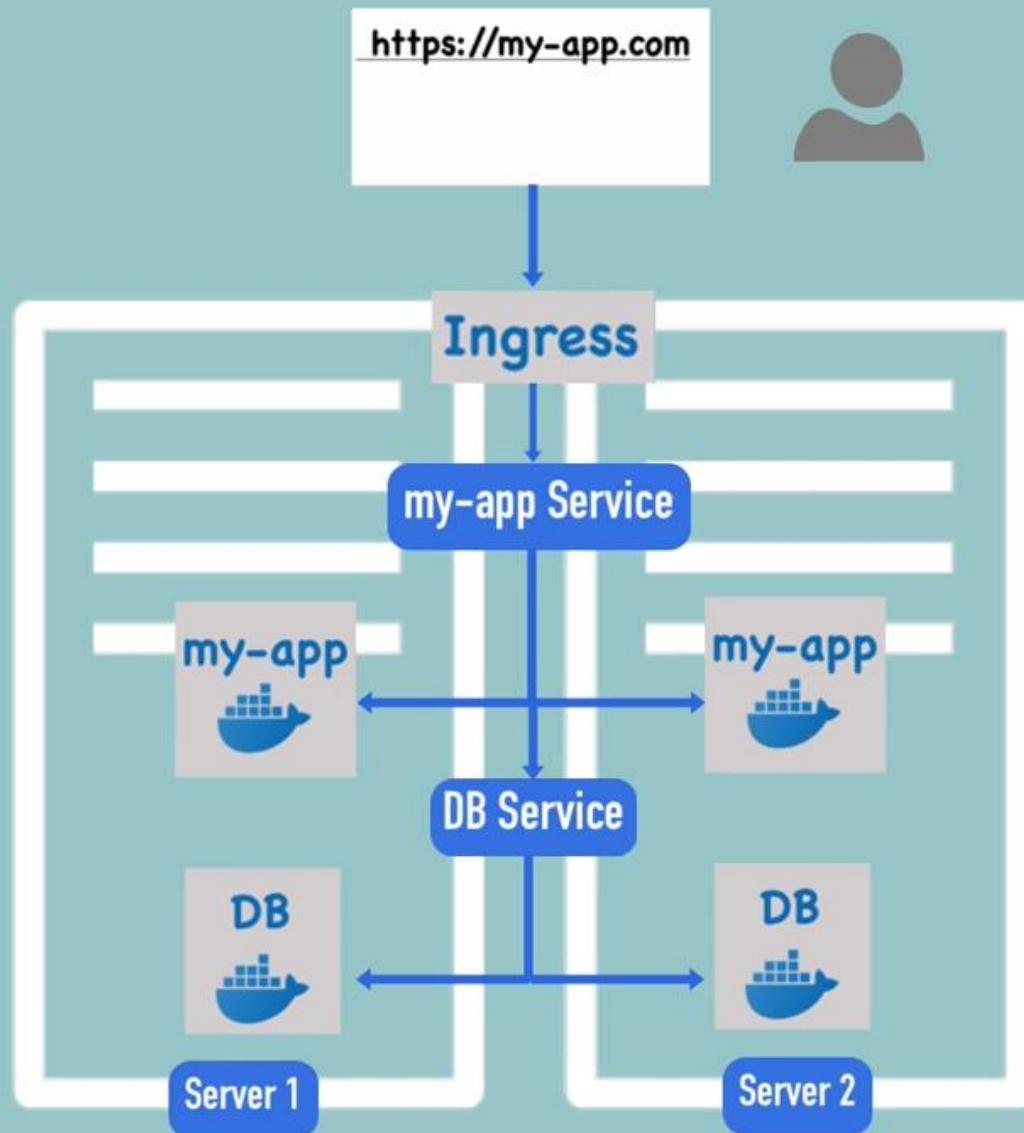




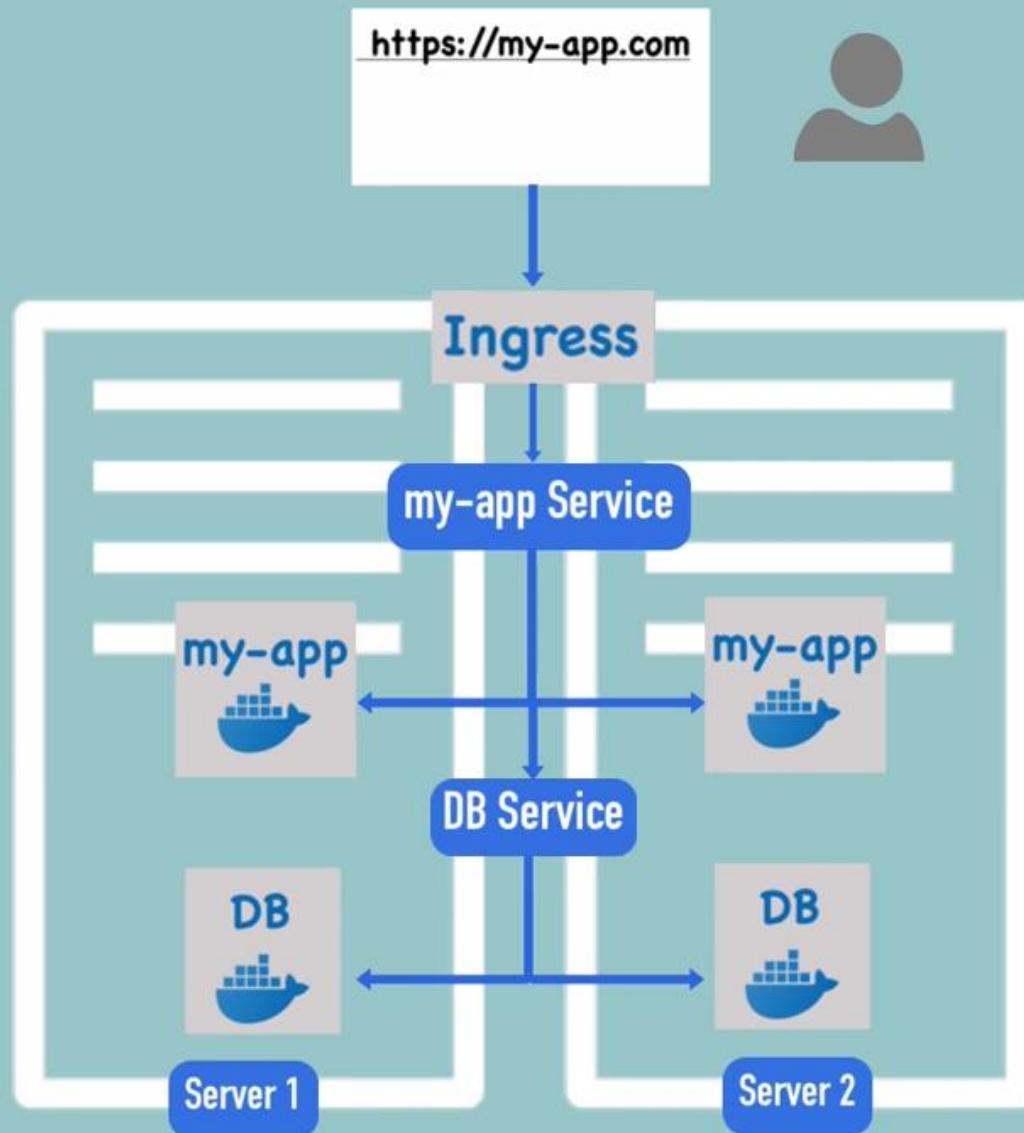




- **EVERY** component is replicated
- load balanced

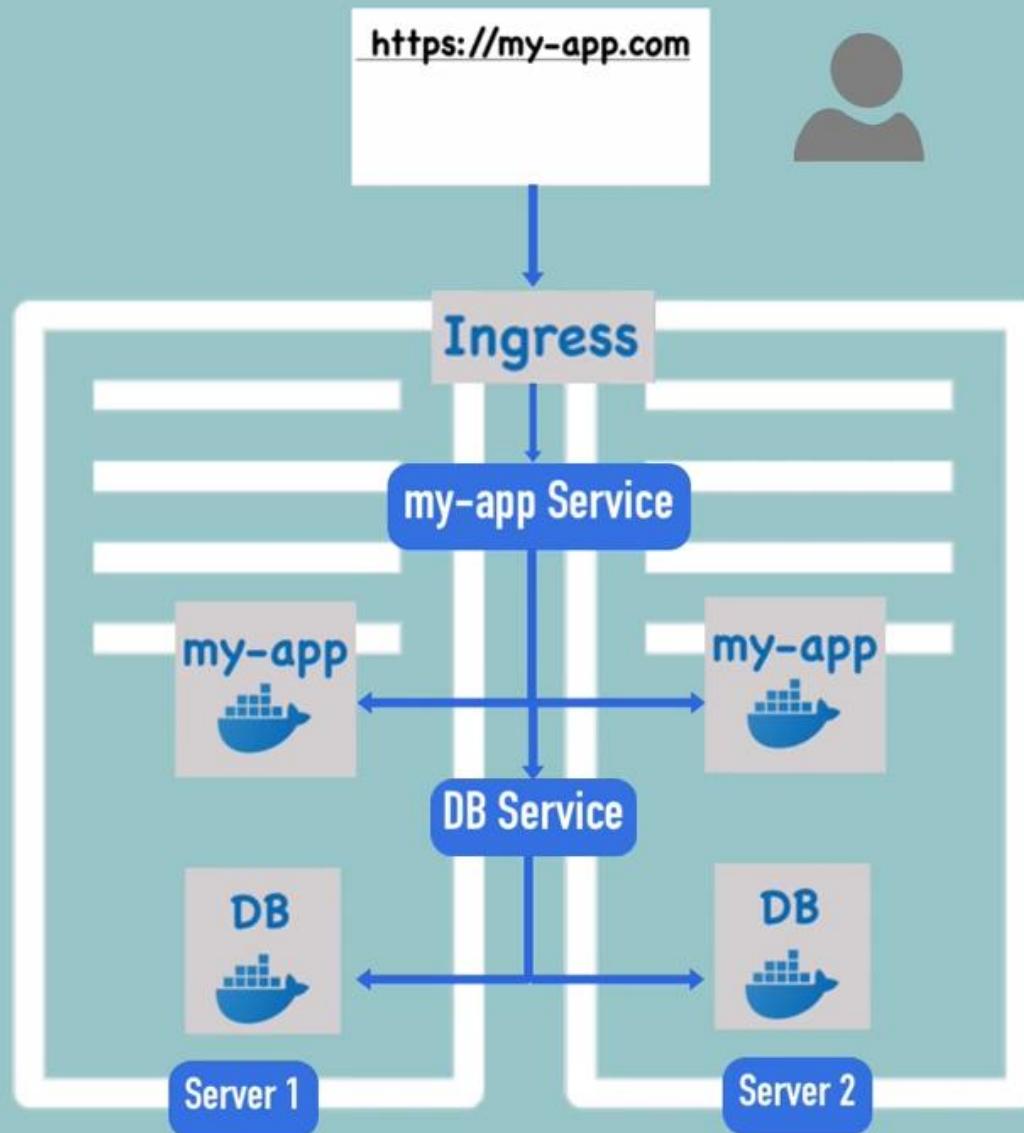


- **EVERY component is replicated**
- **load balanced**
- **no bottleneck that slows down the responses**

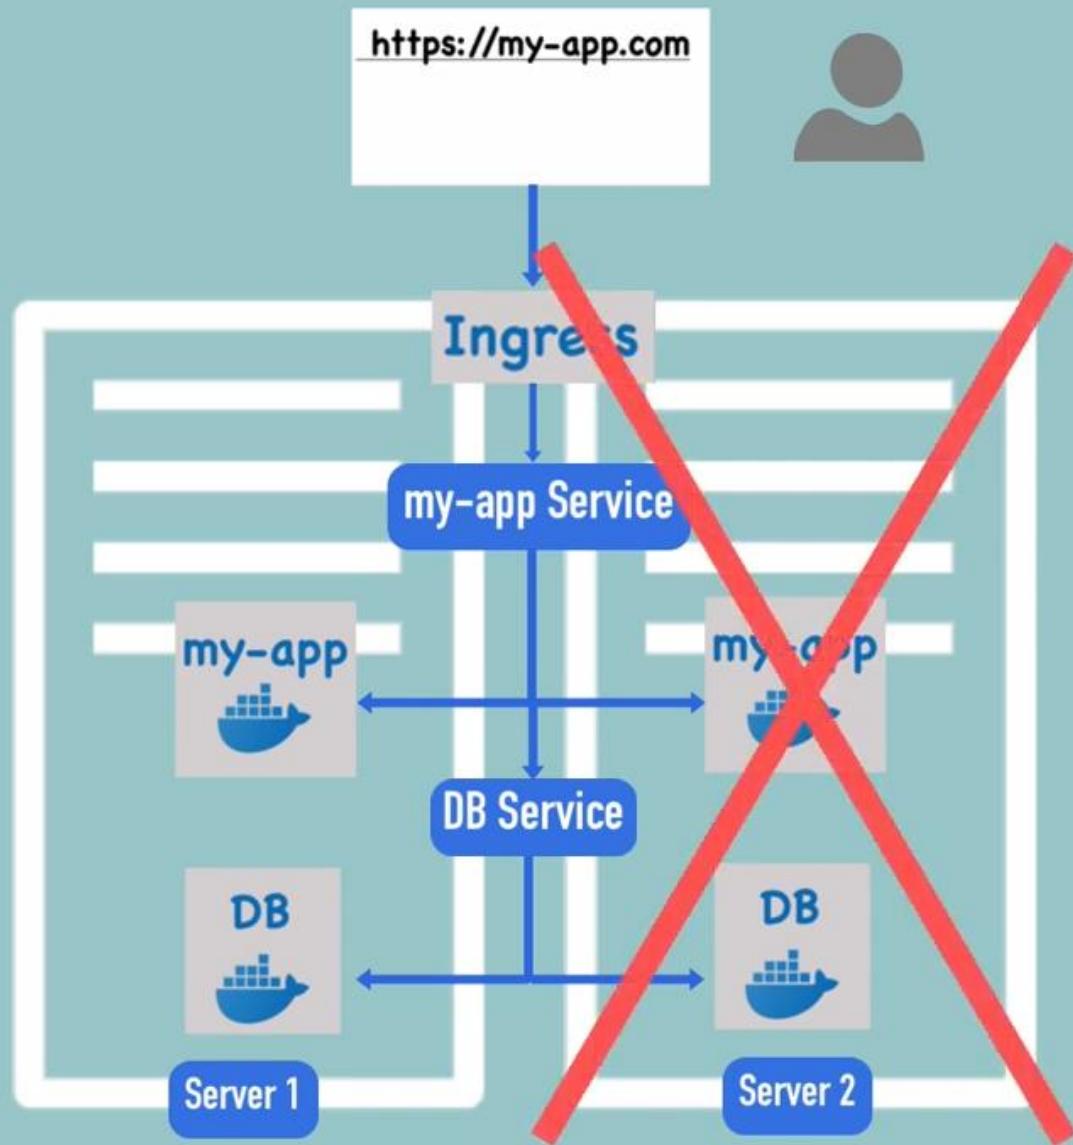


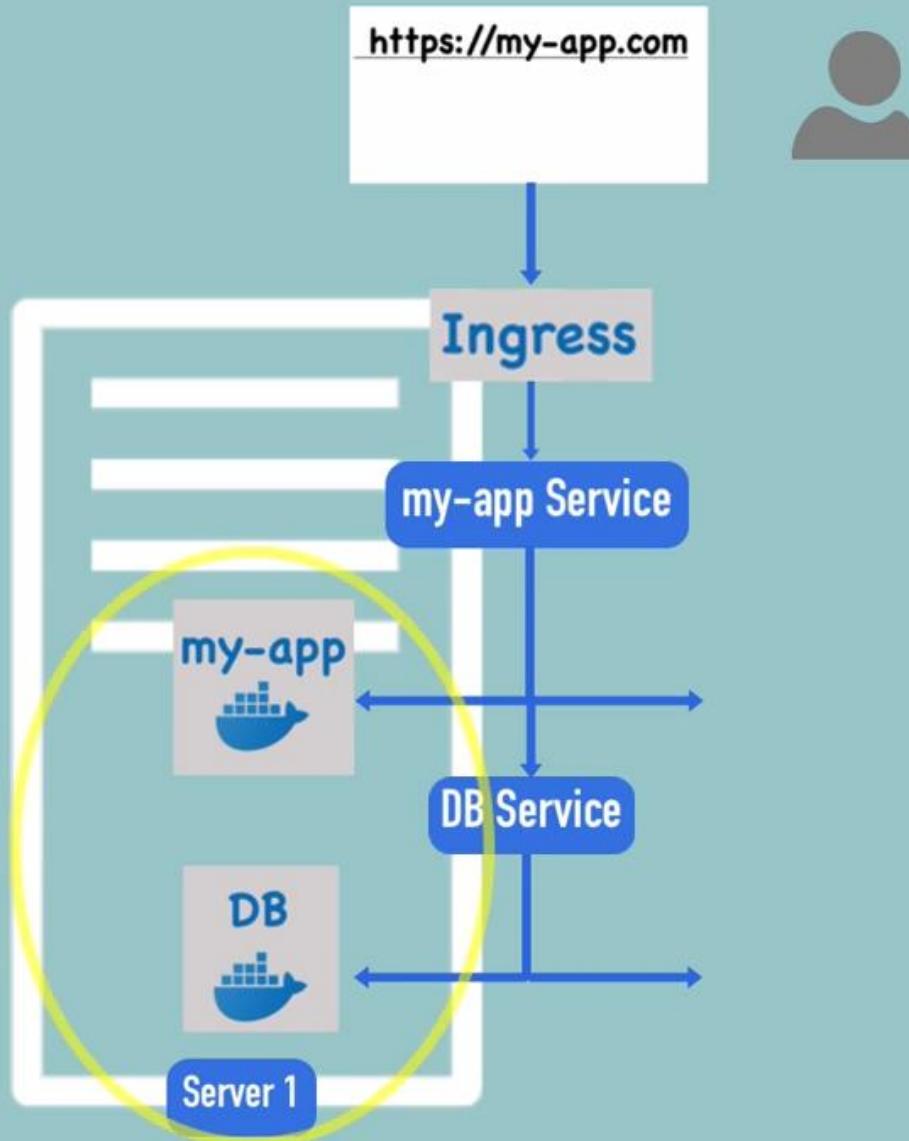
Scalability or  
High Performance  
achieved!

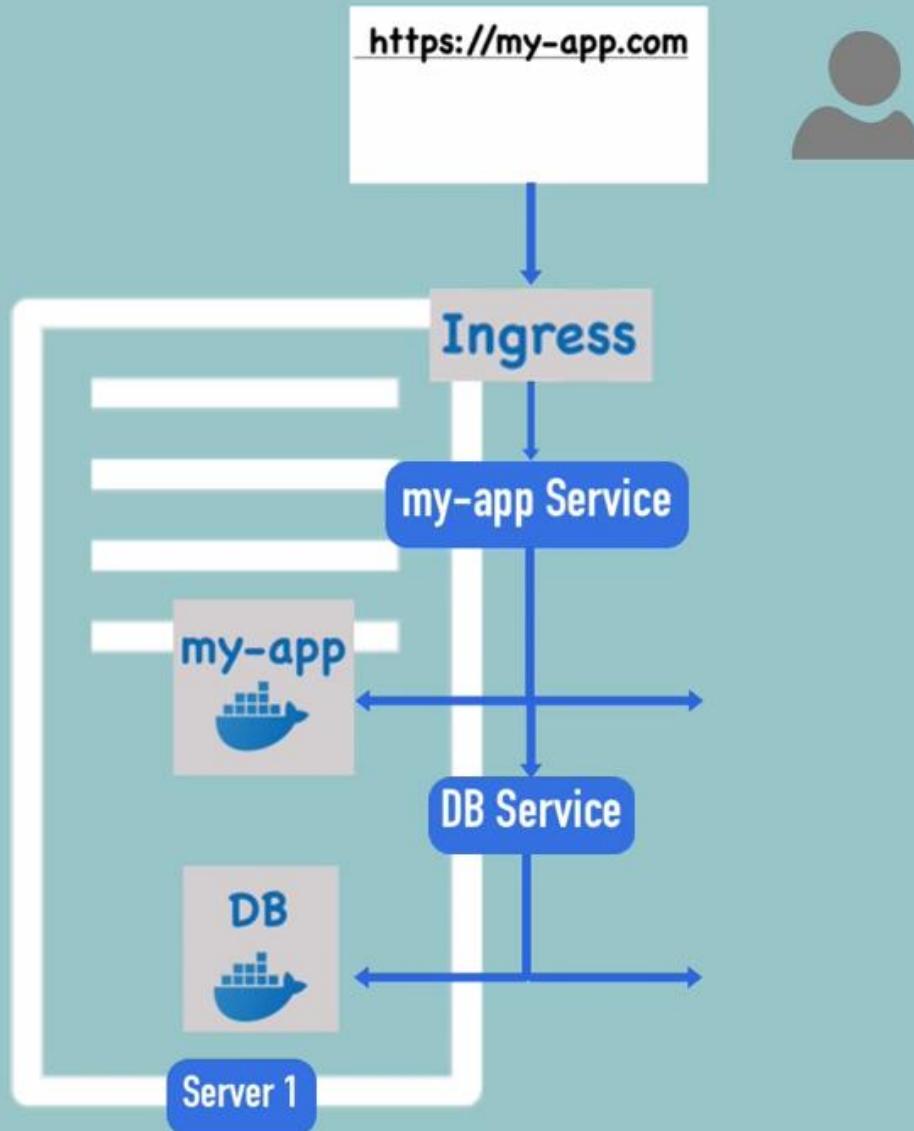




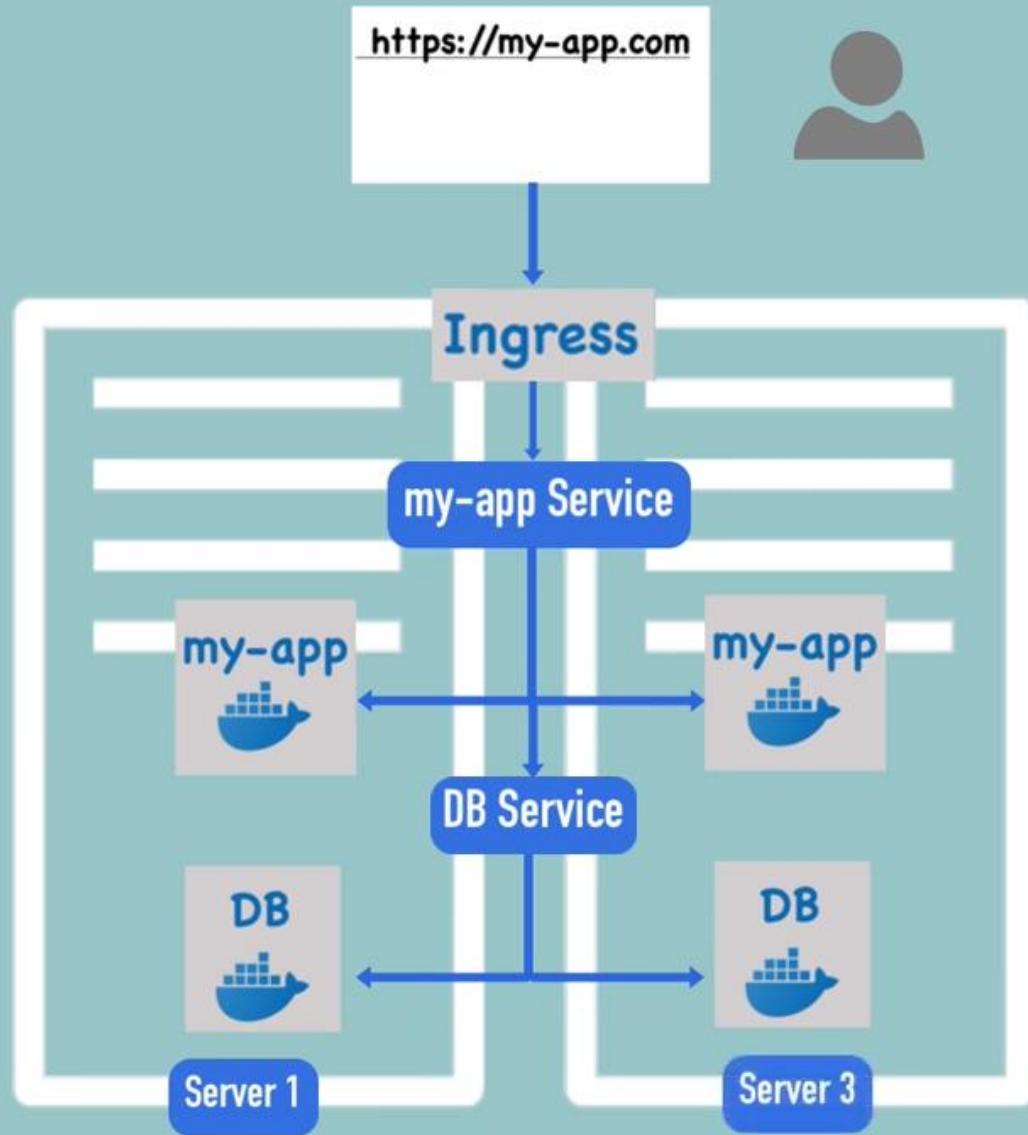
**Nonetheless, you need a  
properly designed application!**

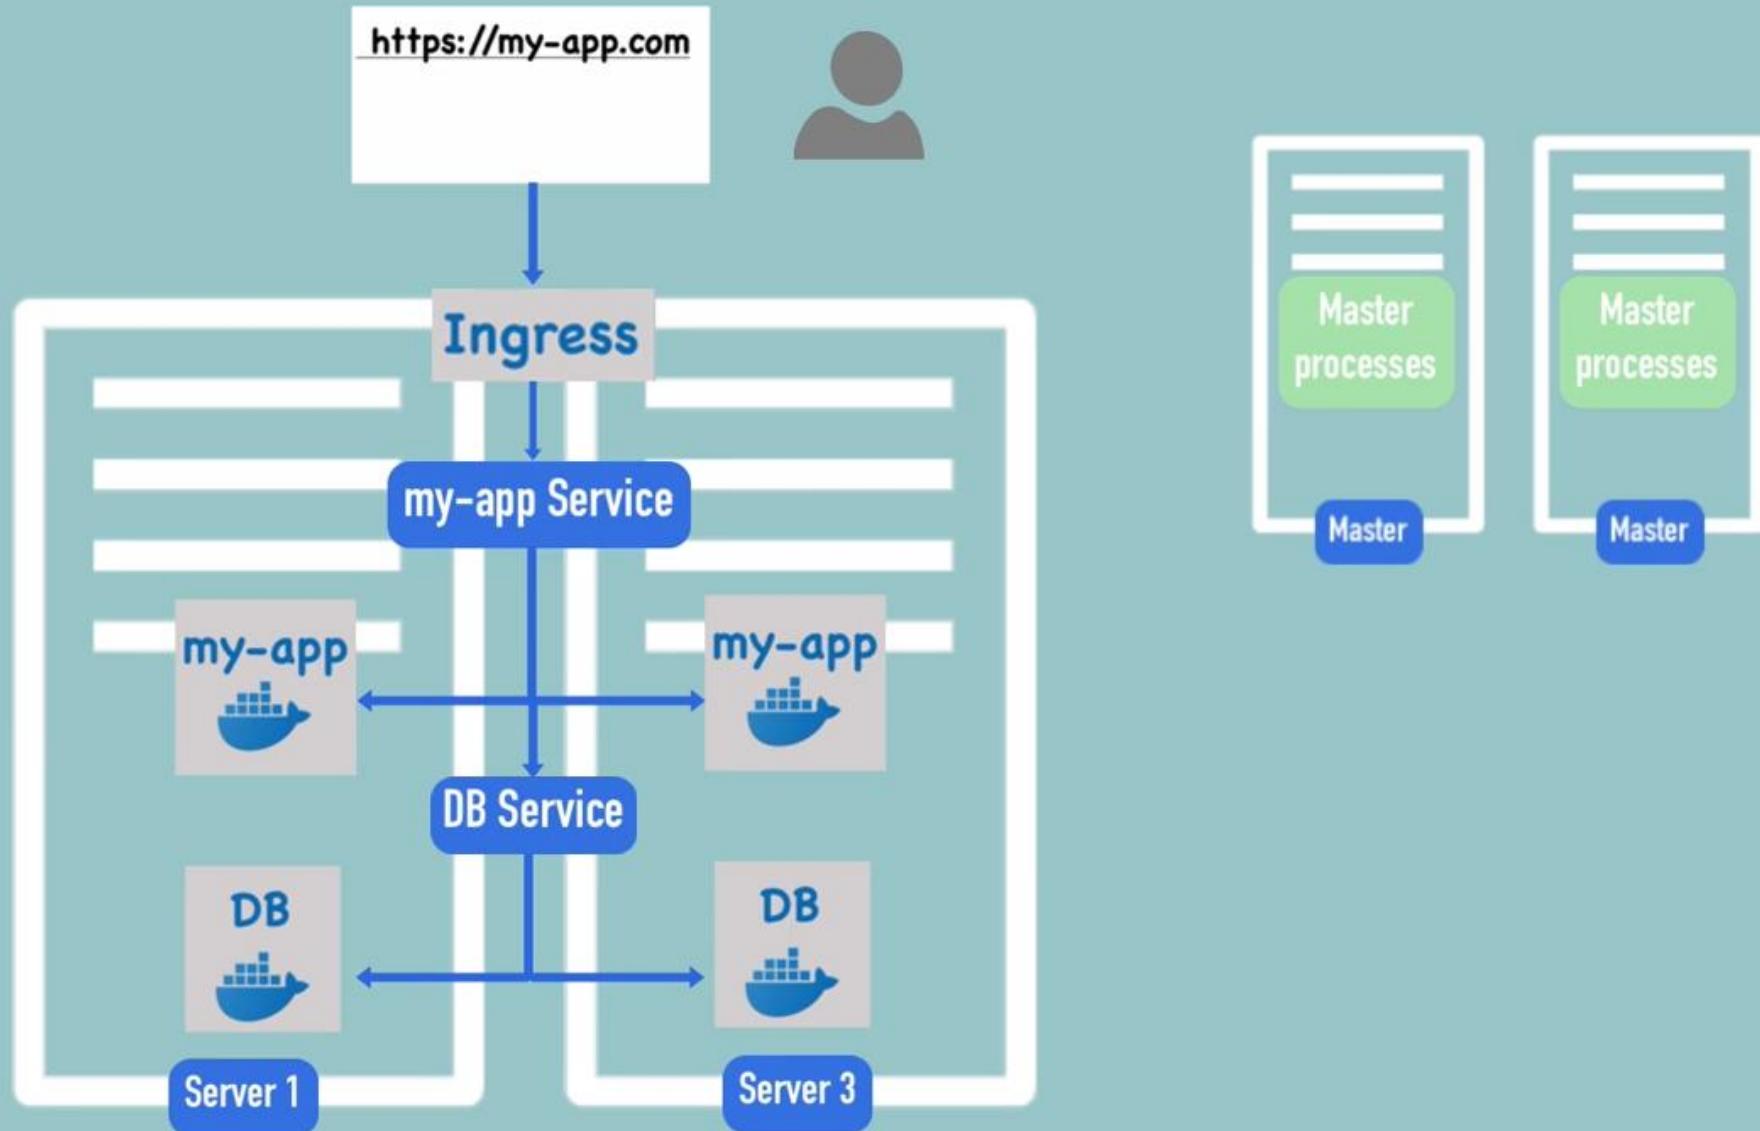


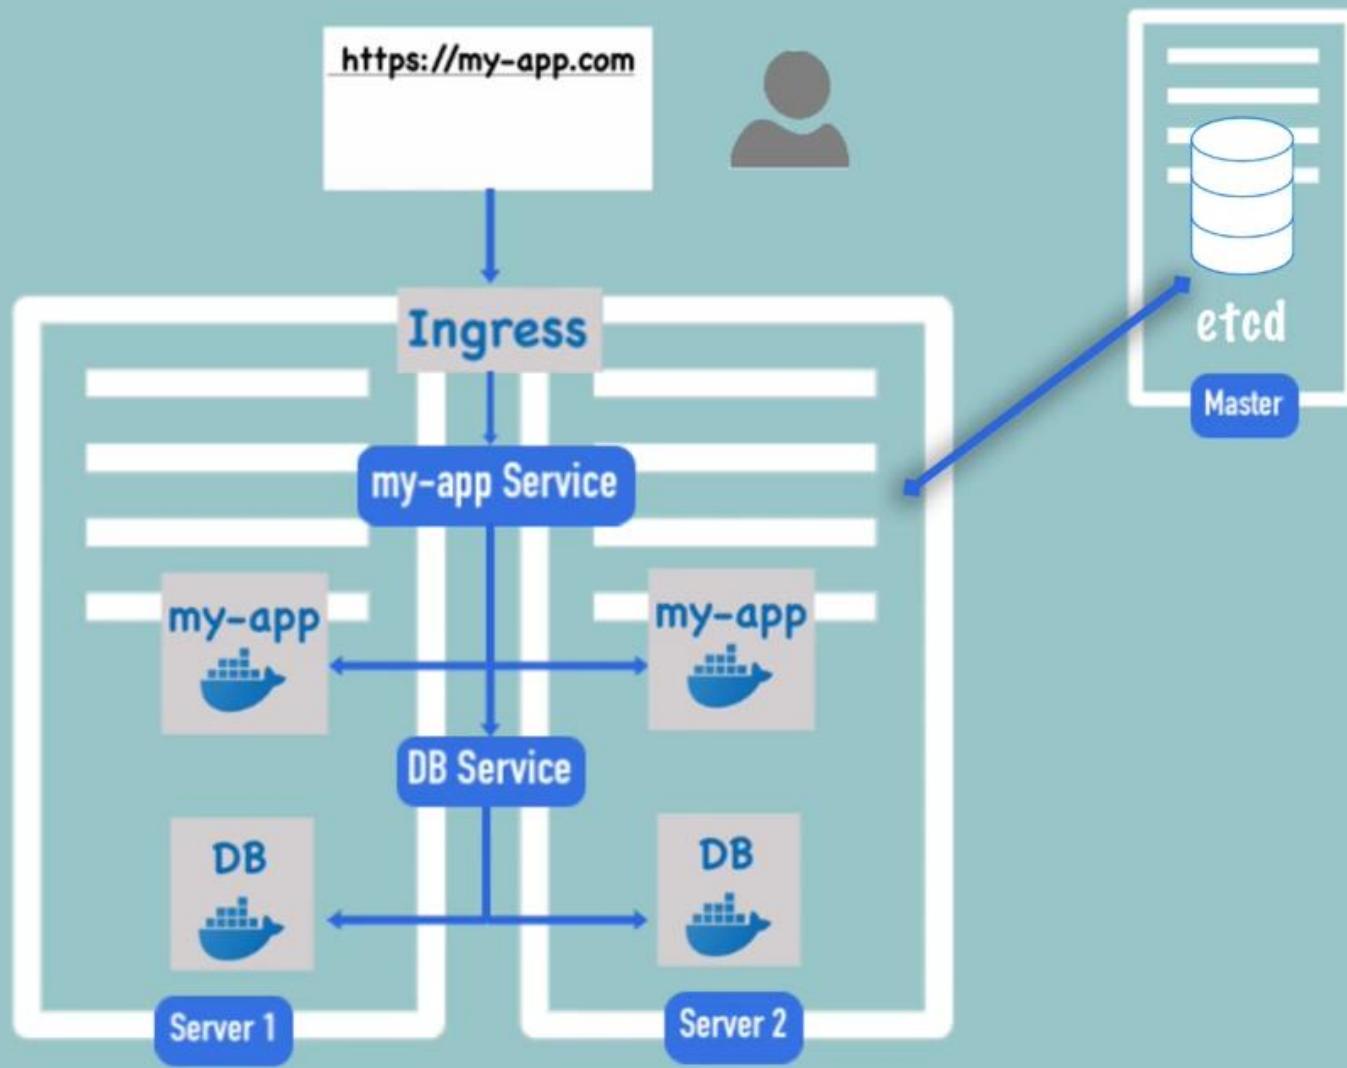


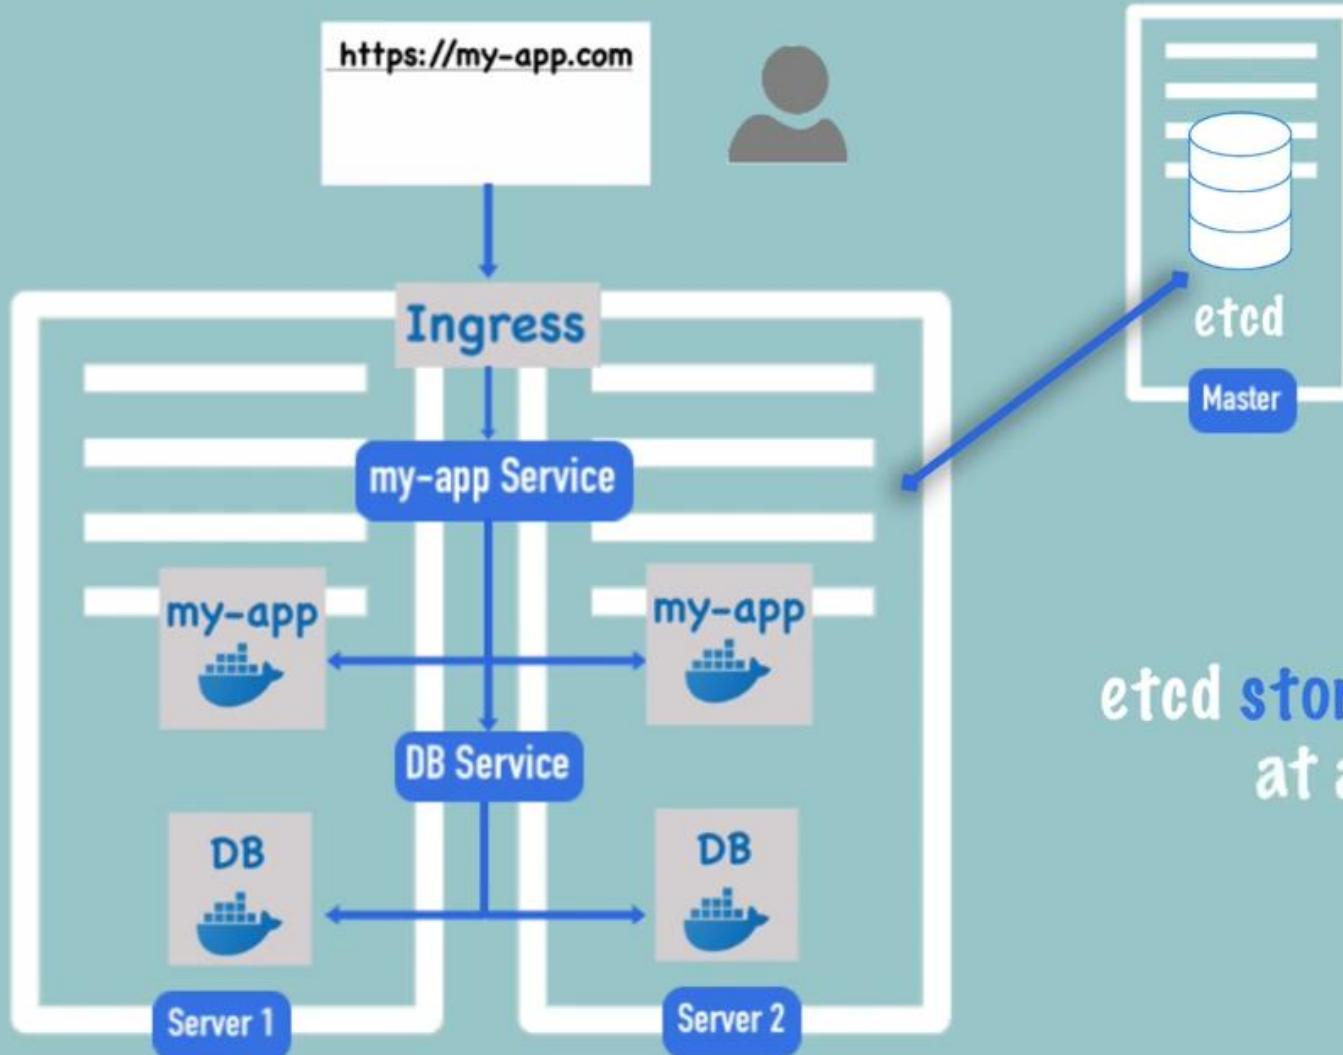


Controller Manager schedules  
new replicas

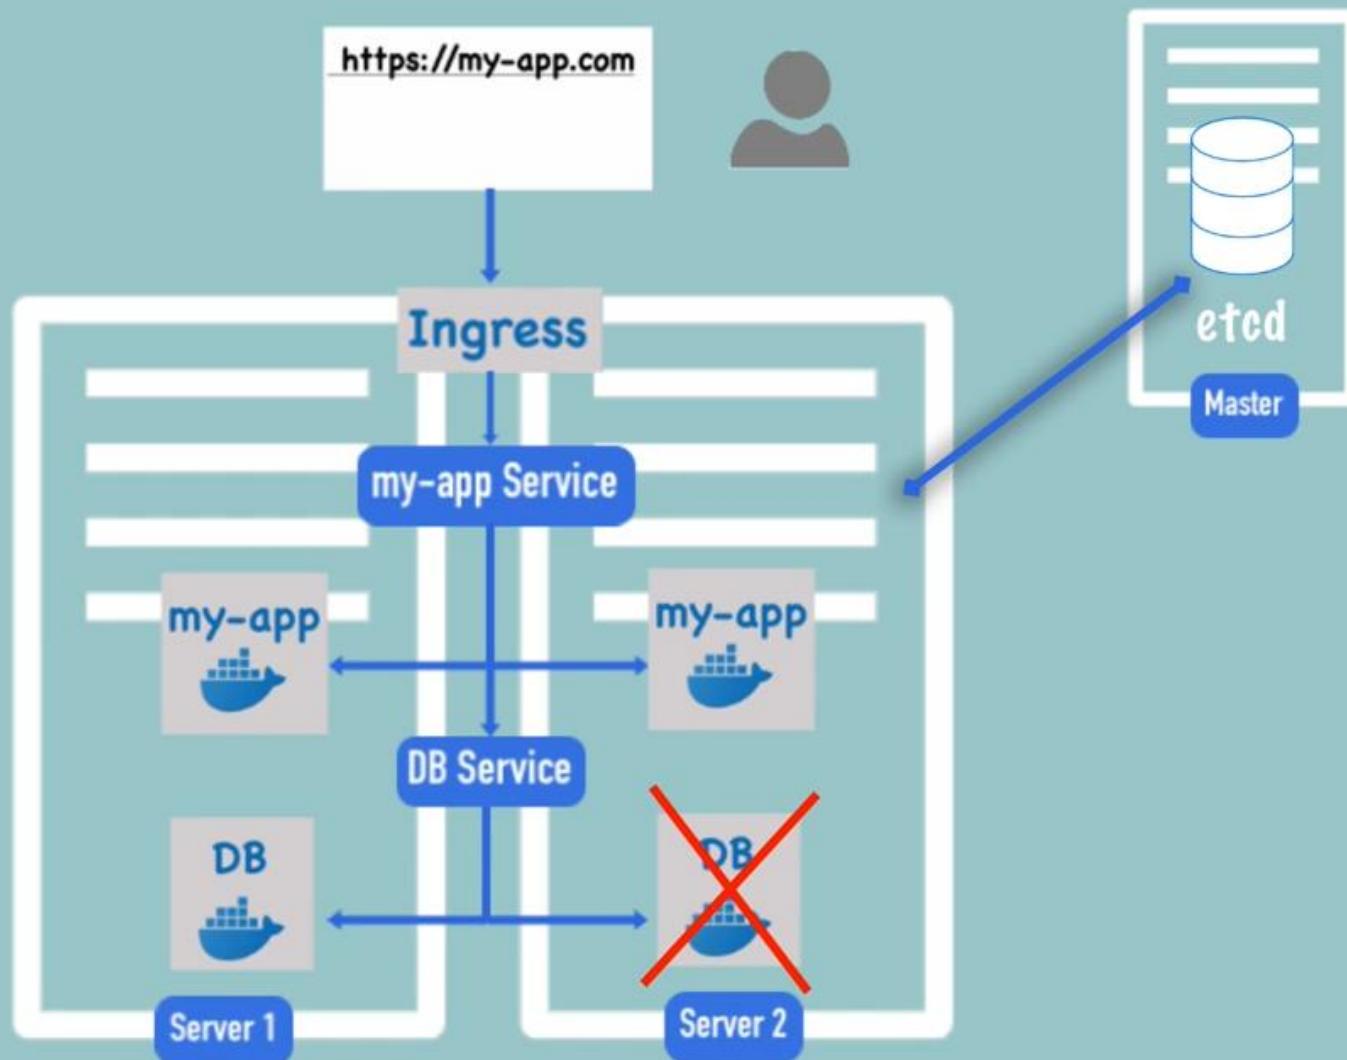


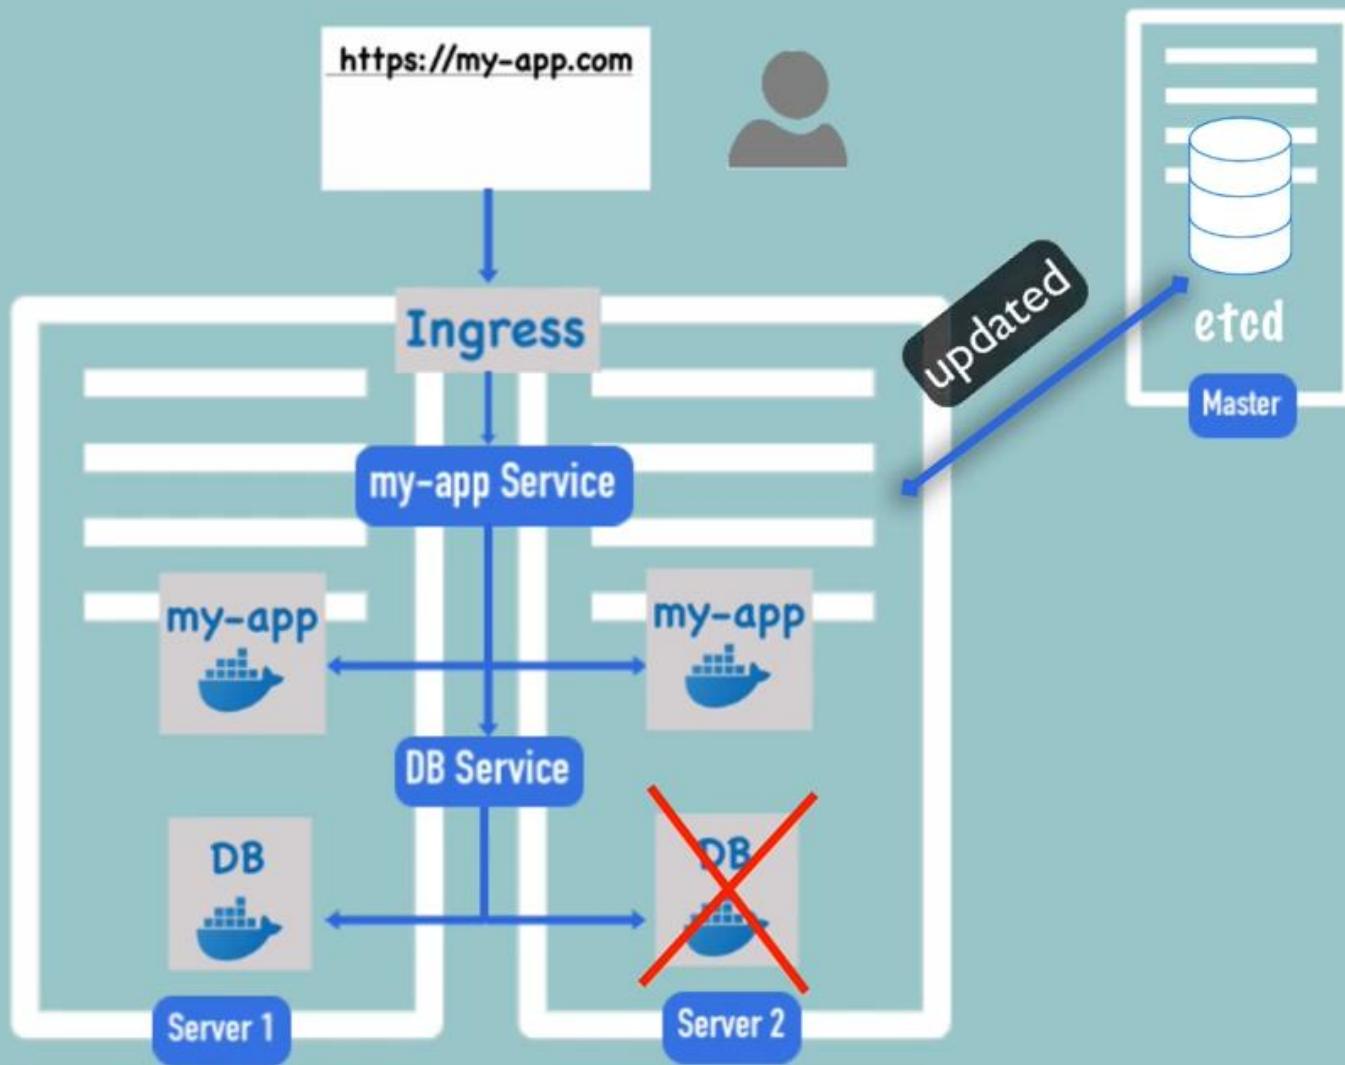


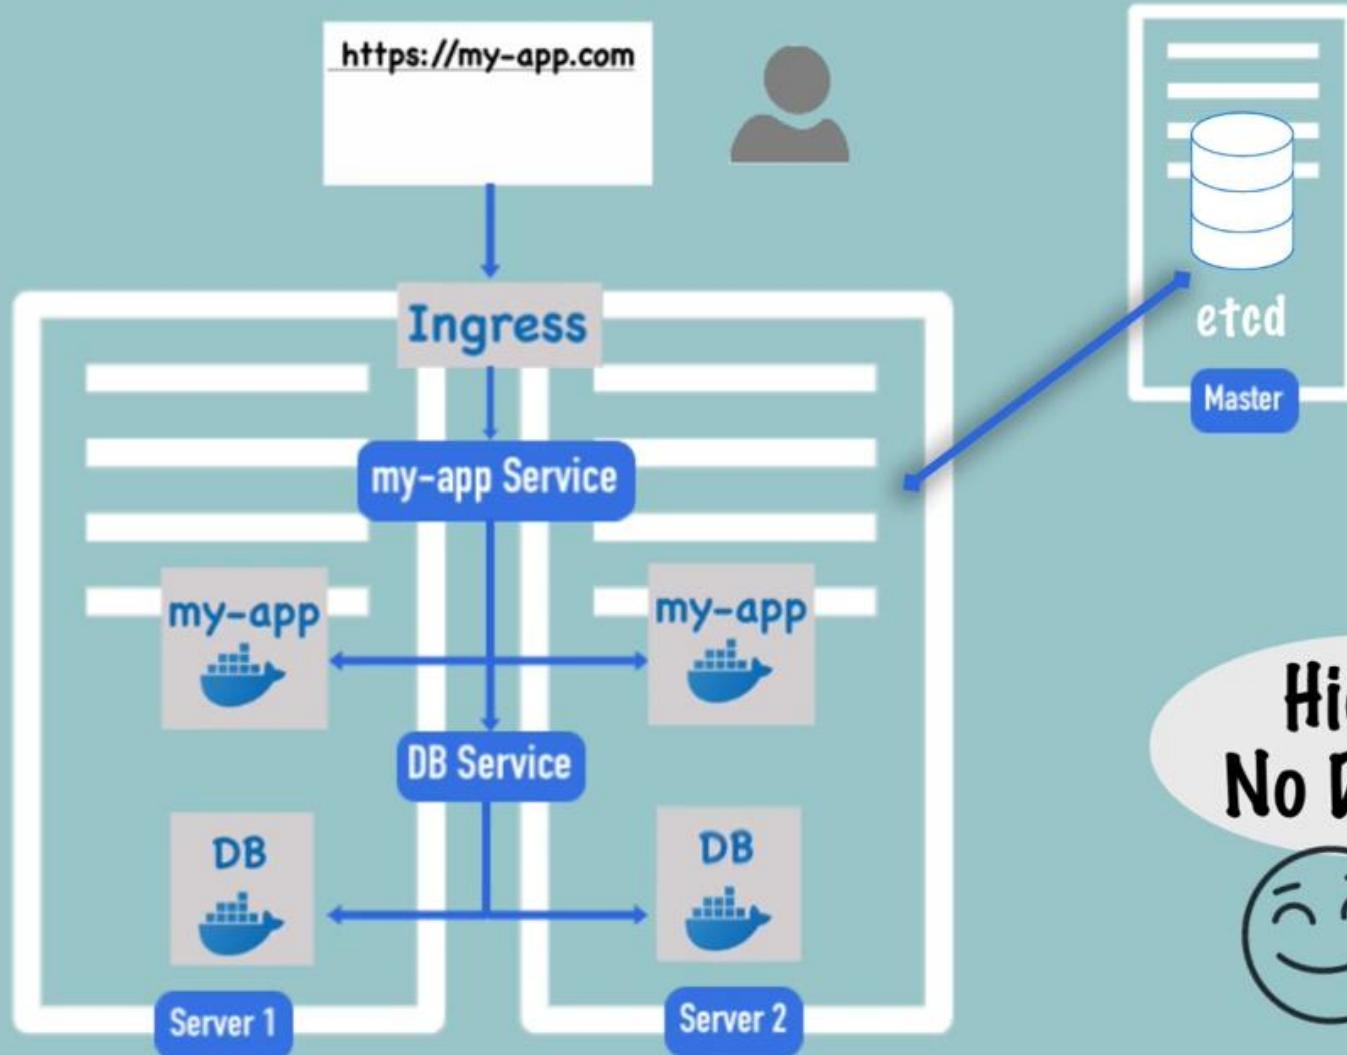




etcd stores the cluster state  
at any given time

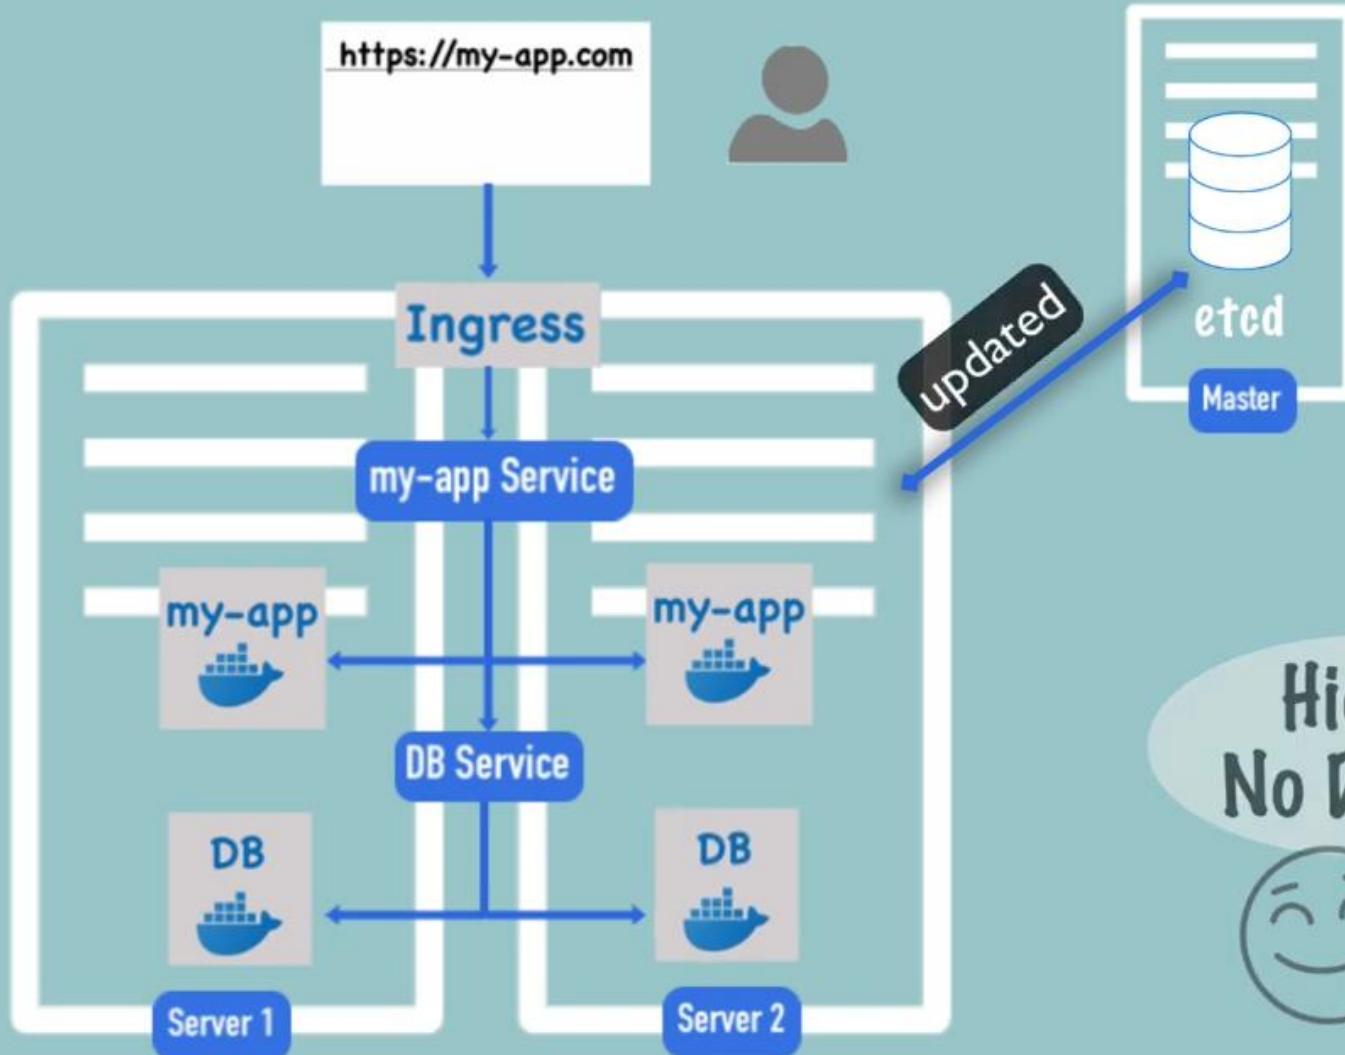






**High Availability or  
No Downtime achieved!**



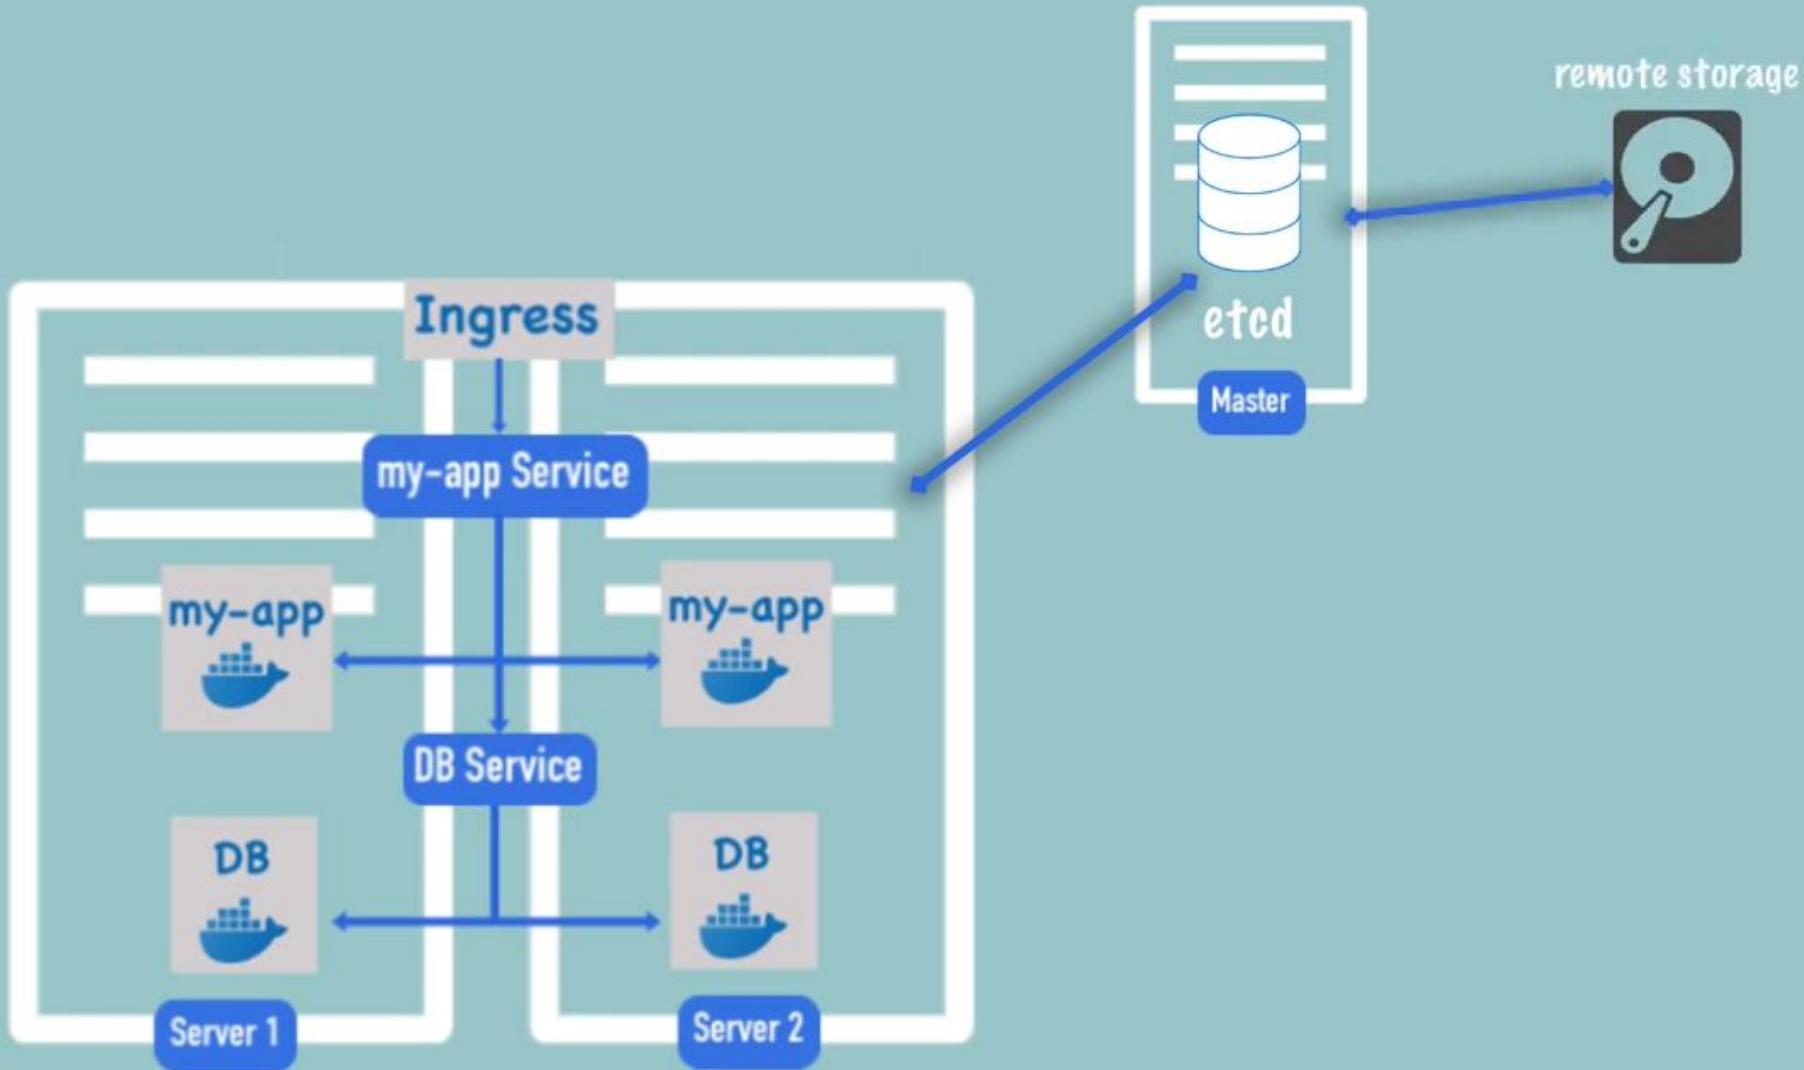


High Availability or  
No Downtime achieved!

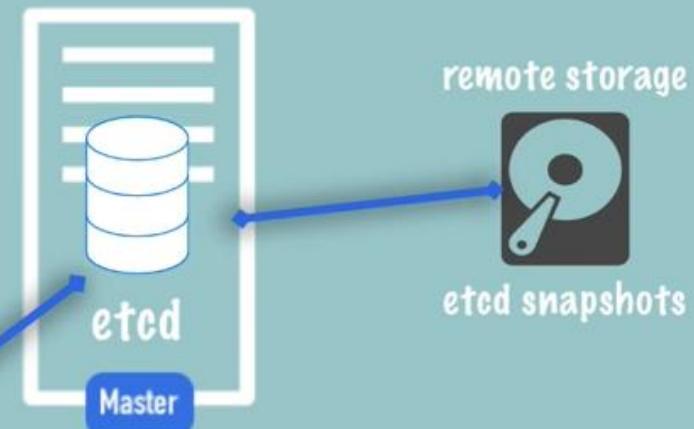
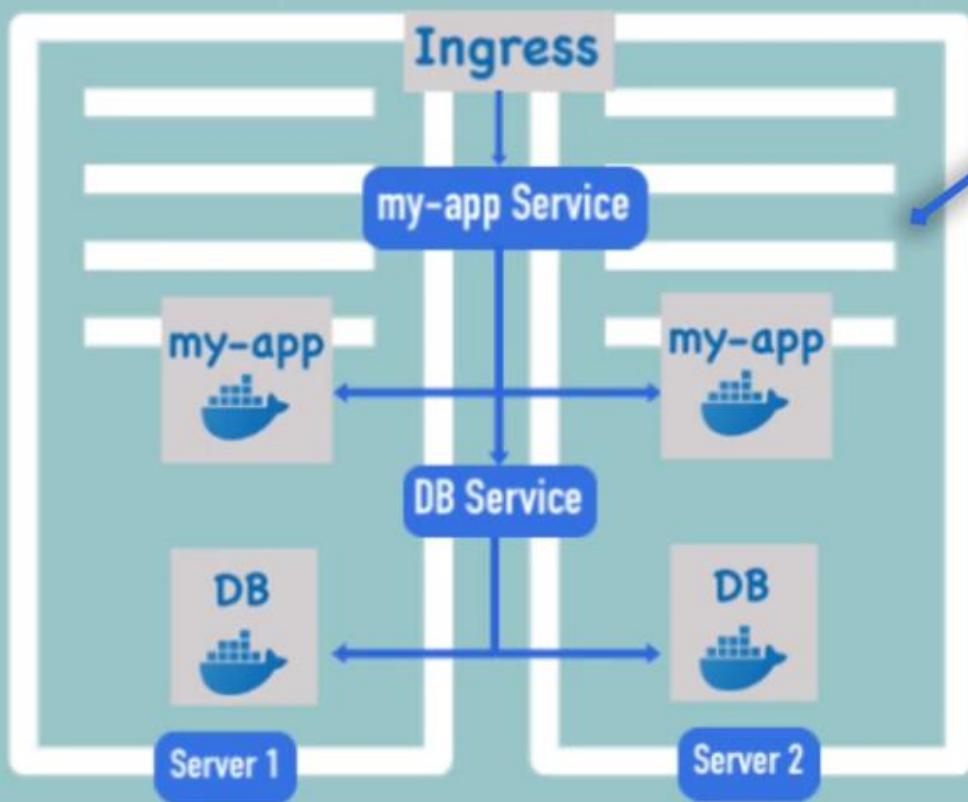


# Disaster recovery

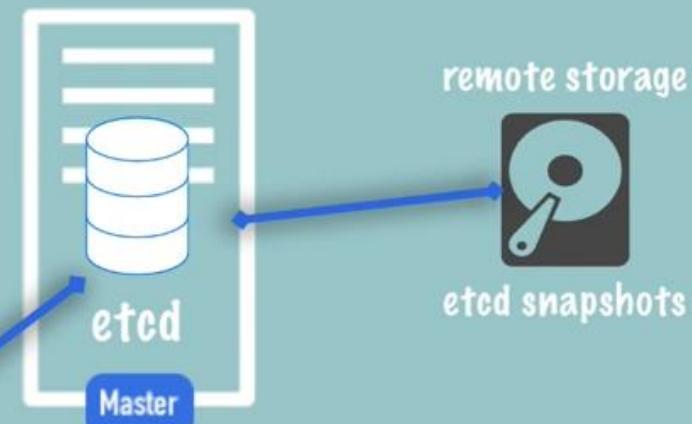
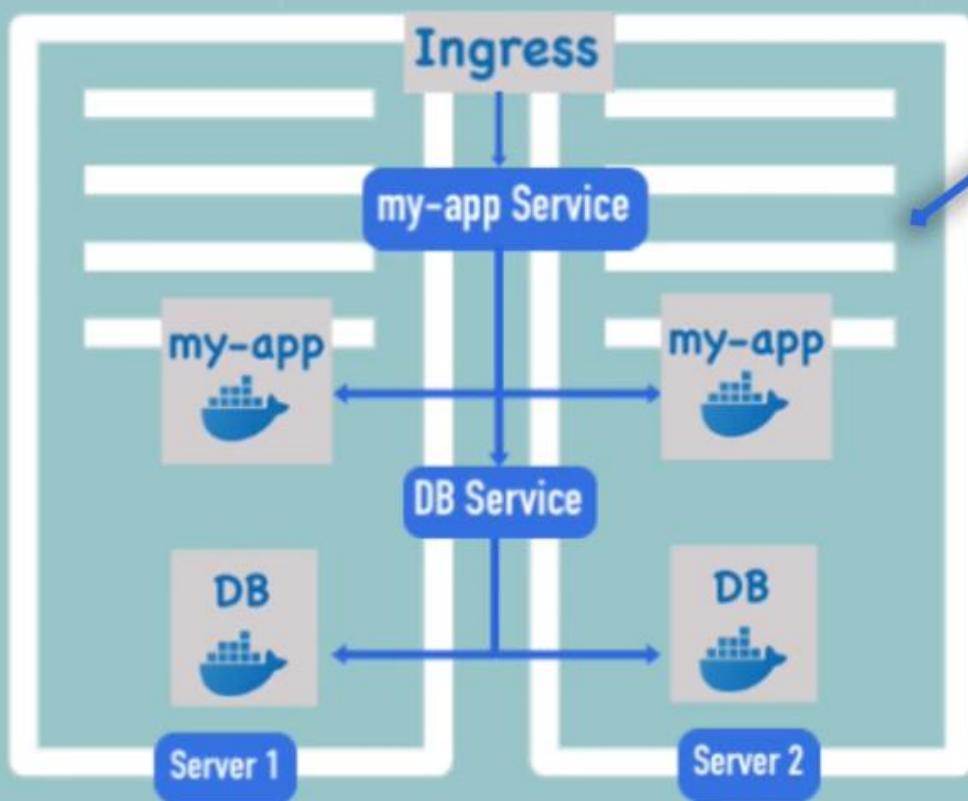




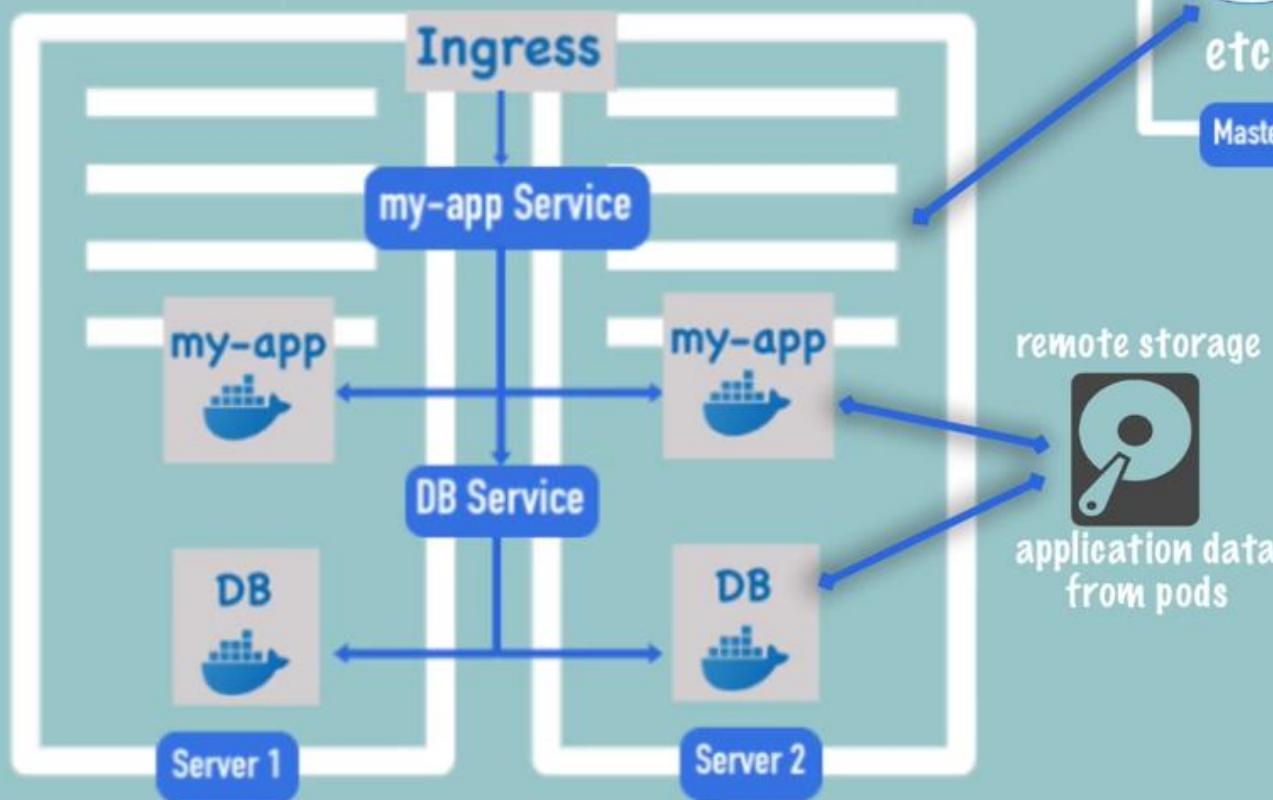
You need to take care  
of the backup!



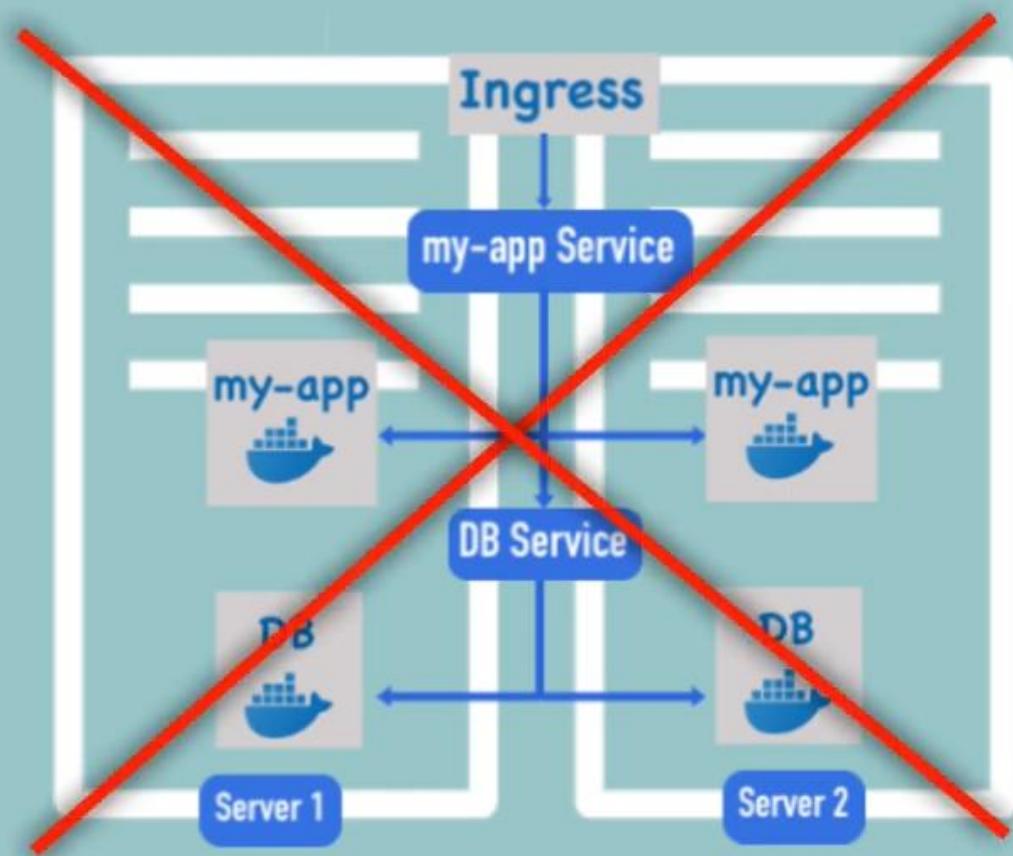
**NO database data is stored in etcd!**



**NO database data is stored in etcd!**



**Storage outside of the cluster!**



remote storage

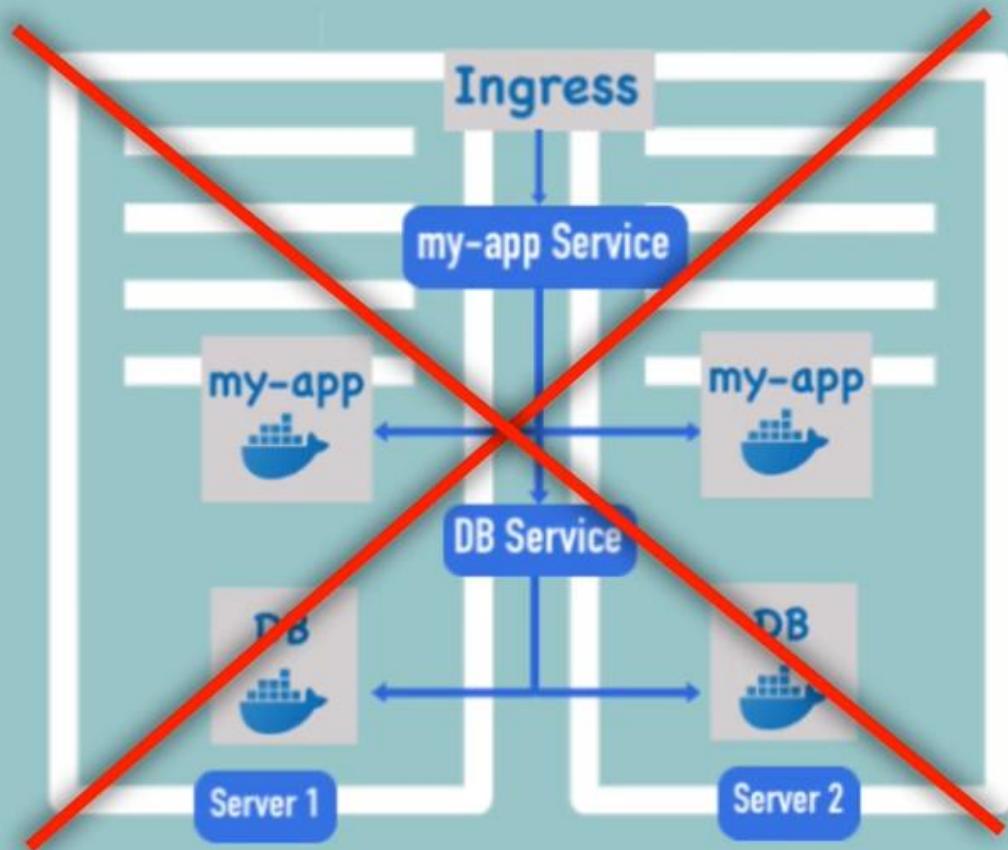


etcd snapshots

remote storage



application data  
from pods



remote storage

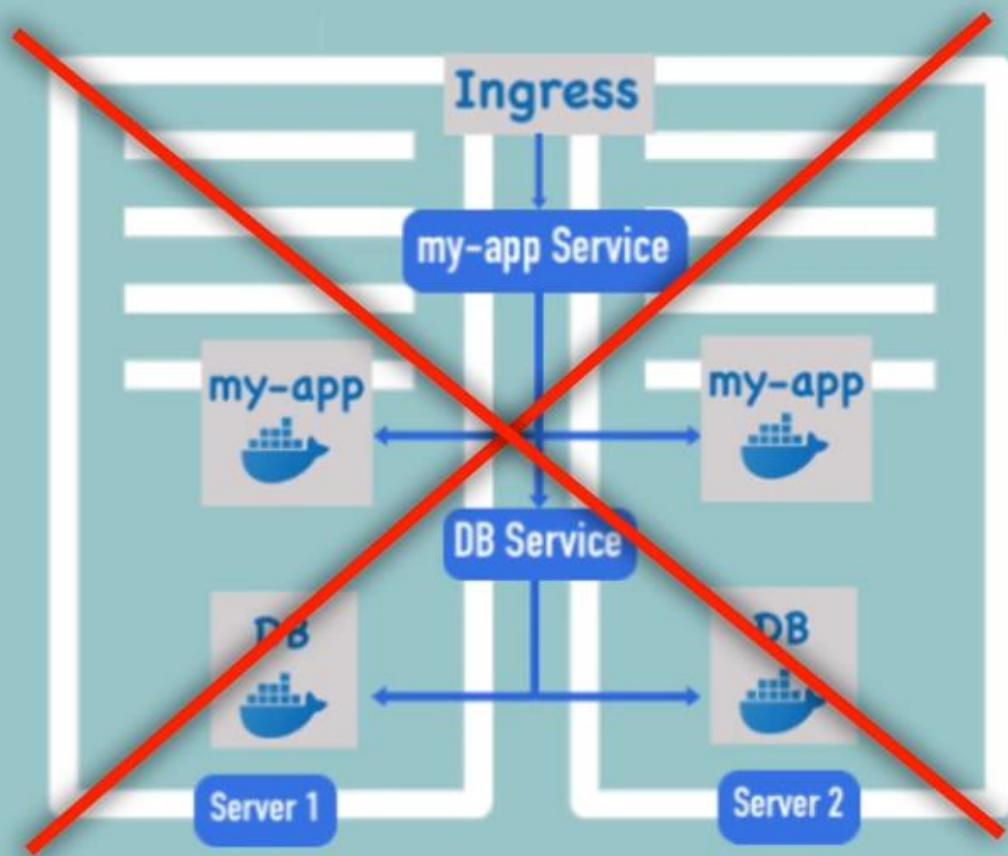


etcd snapshots

remote storage



application data  
from pods



remote storage



etcd snapshots

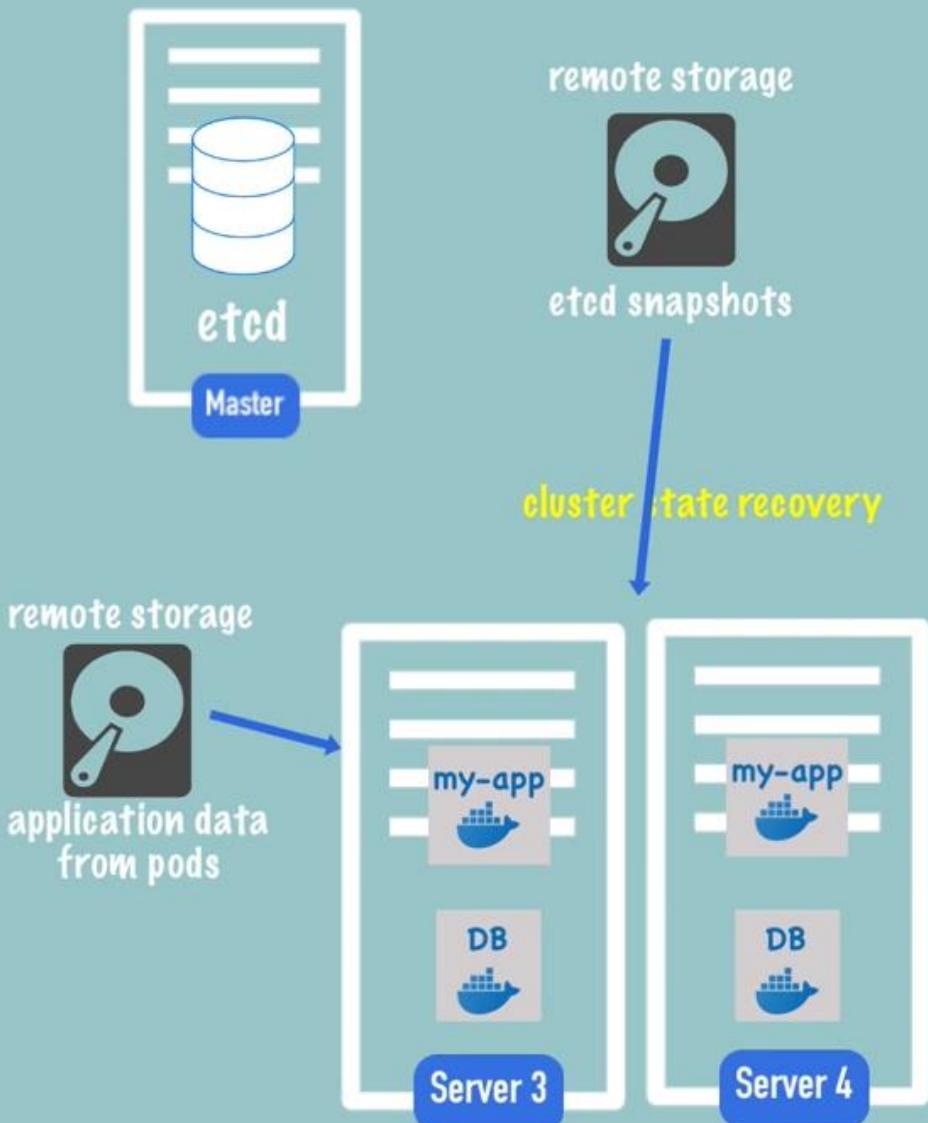
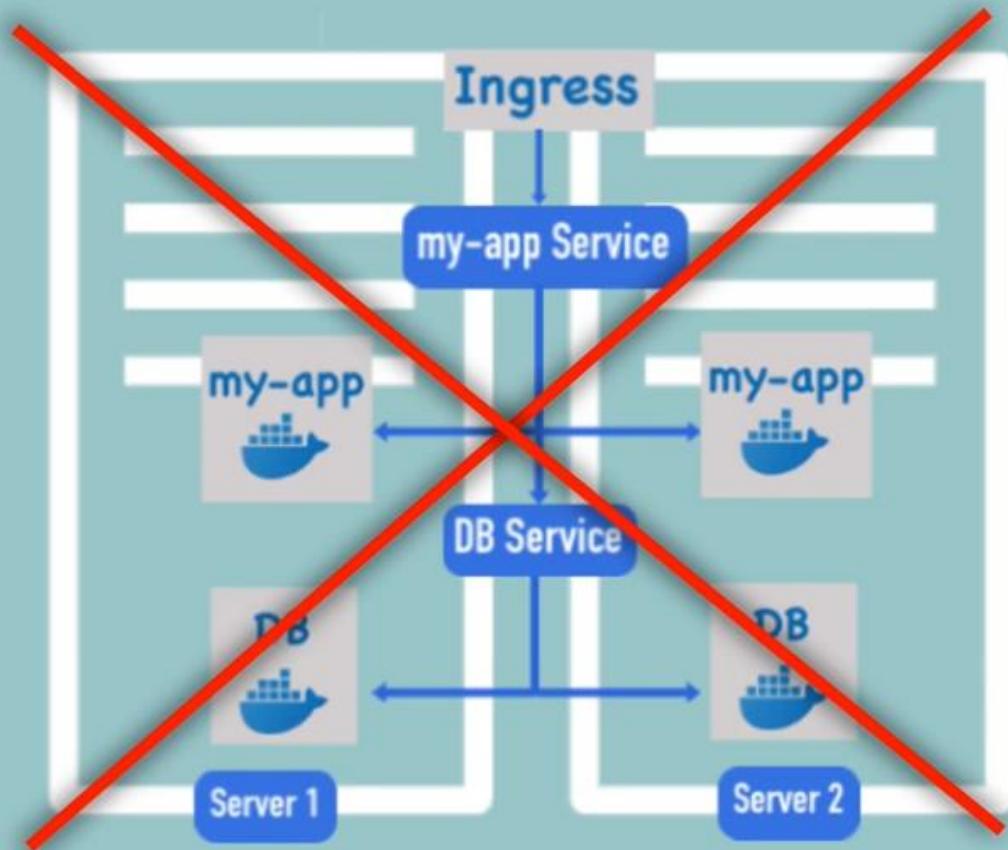
cluster state recovery

remote storage



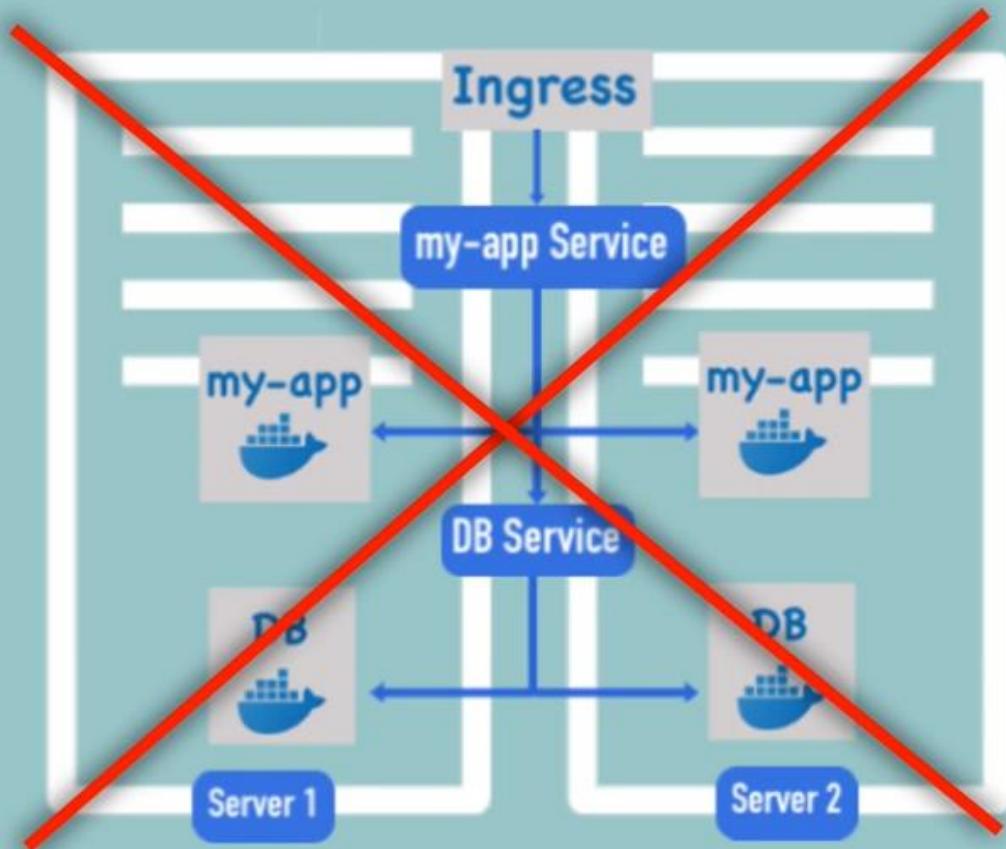
application data  
from pods







# Easy Disaster Recovery !



remote storage

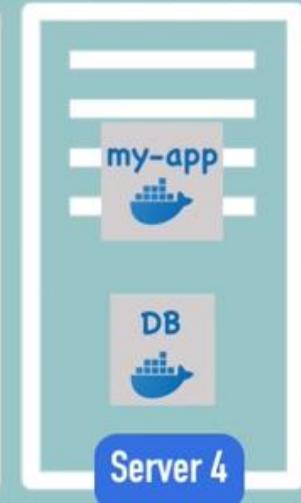
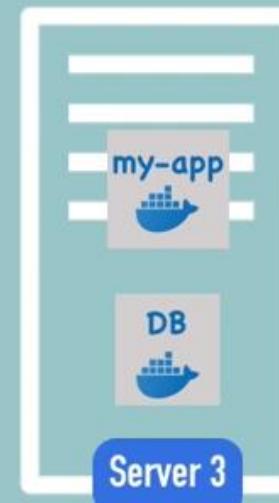


etcd snapshots

remote storage

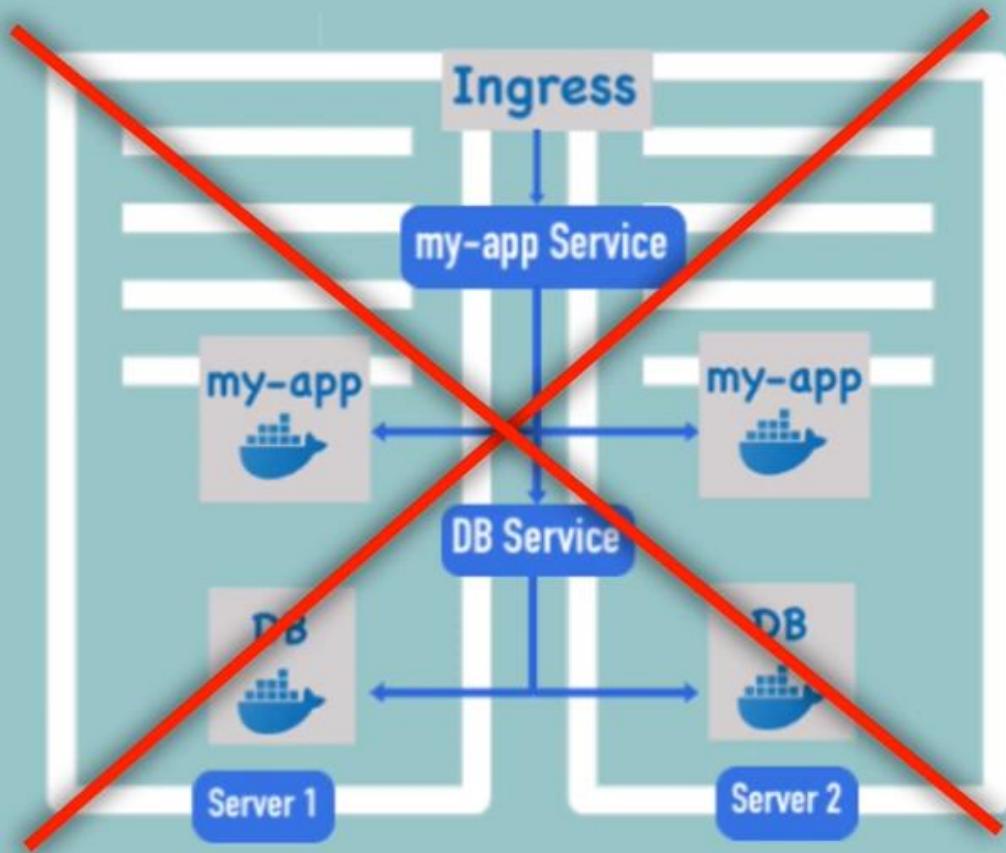


application data  
from pods





# Easy Disaster Recovery !



remote storage

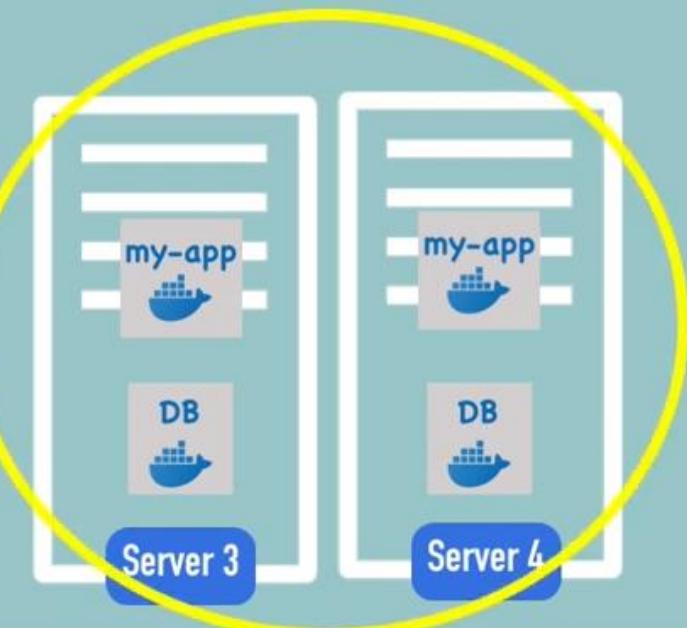


etcd snapshots

remote storage



application data  
from pods





# Kubernetes vs. AWS





# Advantages of Kubernetes

- replication is much easier
- self-healing of K8s



or  
other platforms

or  
own setup



# Advantages of Kubernetes

- replication is much easier
- self-healing of K8s



or  
other platforms

or  
own setup



# Advantages of Kubernetes

- replication is much easier
- self-healing of K8s



or  
other platforms

or  
own setup



# Advantages of Kubernetes

- replication is much easier
- self-healing of K8s
- smart scheduling



or  
other platforms

or  
own setup



# Advantages of Kubernetes

- replication is much easier
- self-healing of K8s
- smart scheduling



or  
other platforms

or  
own setup